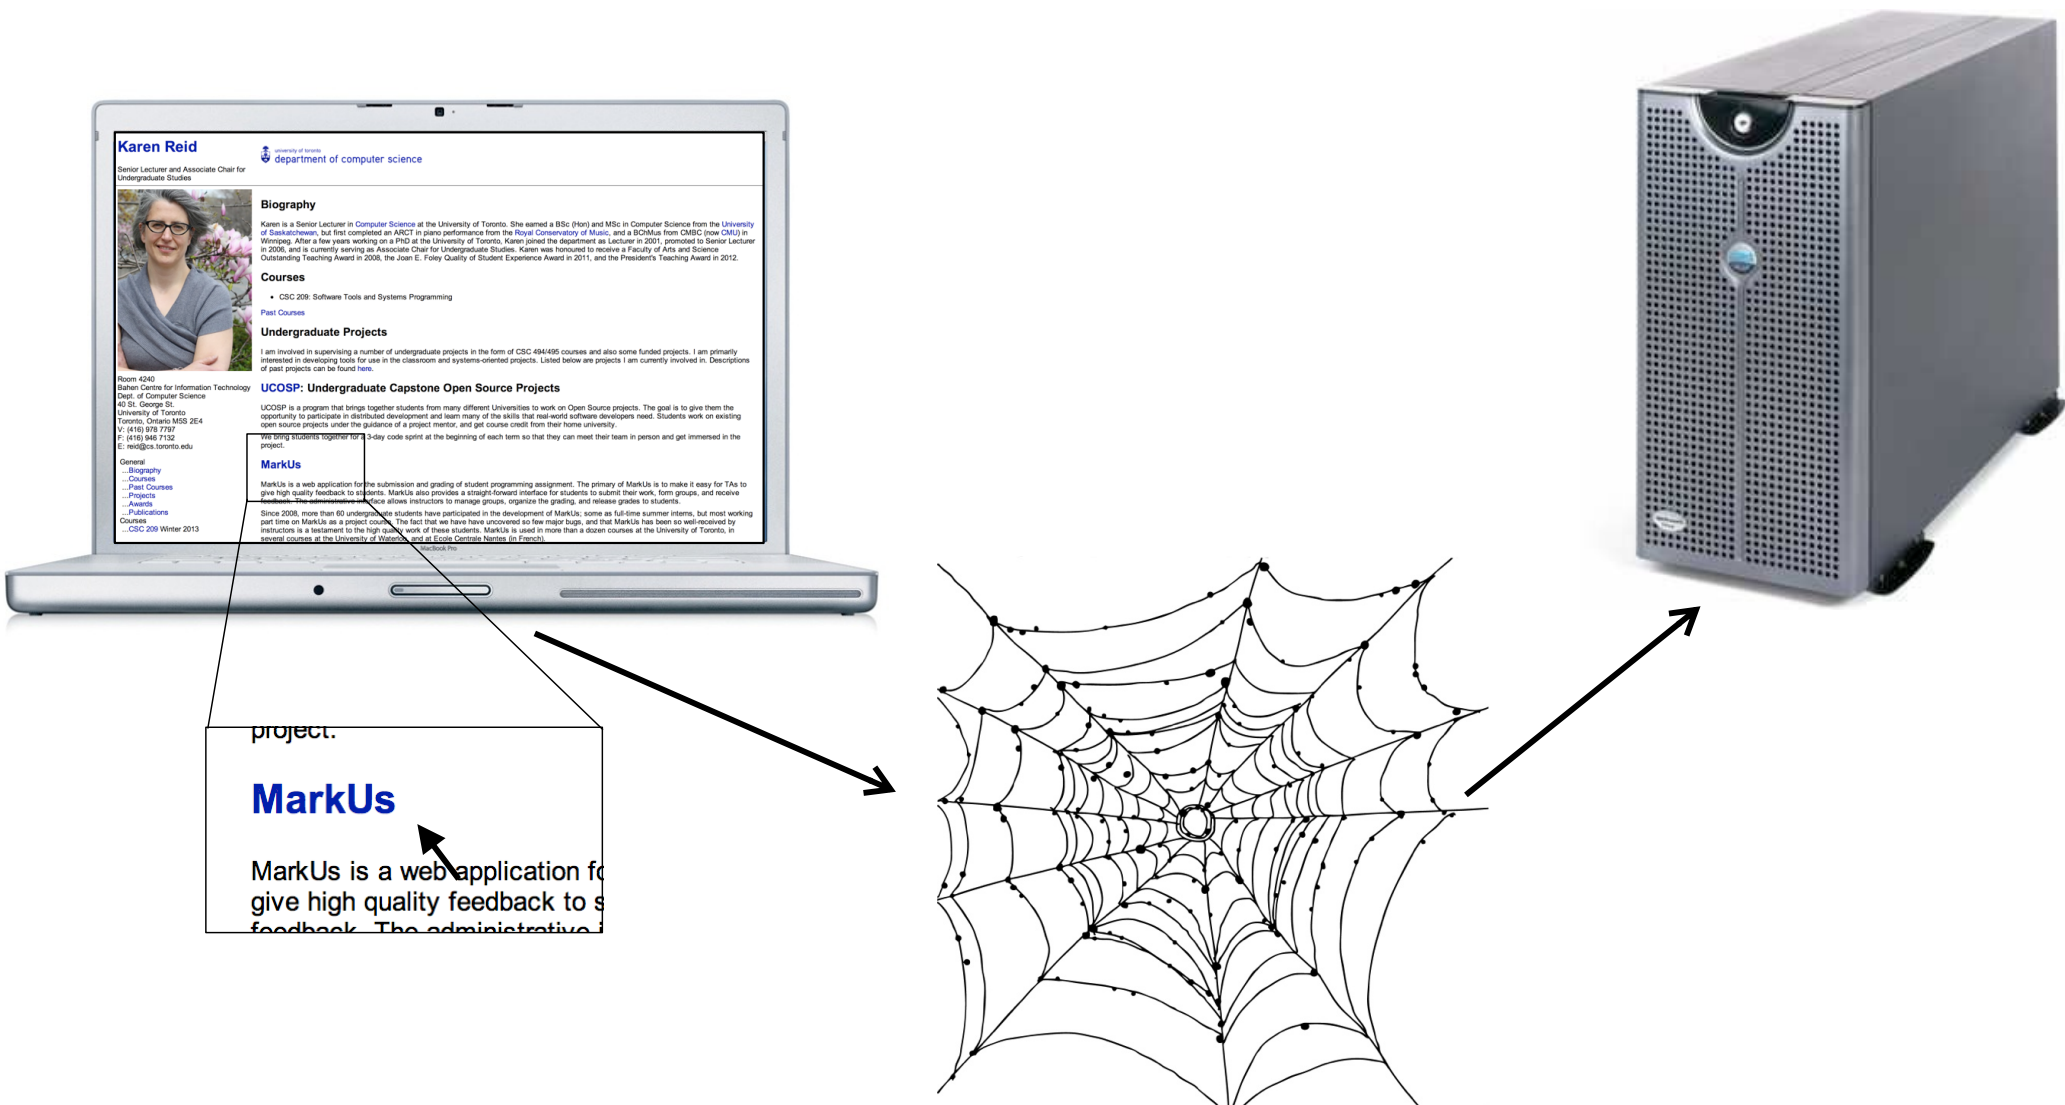





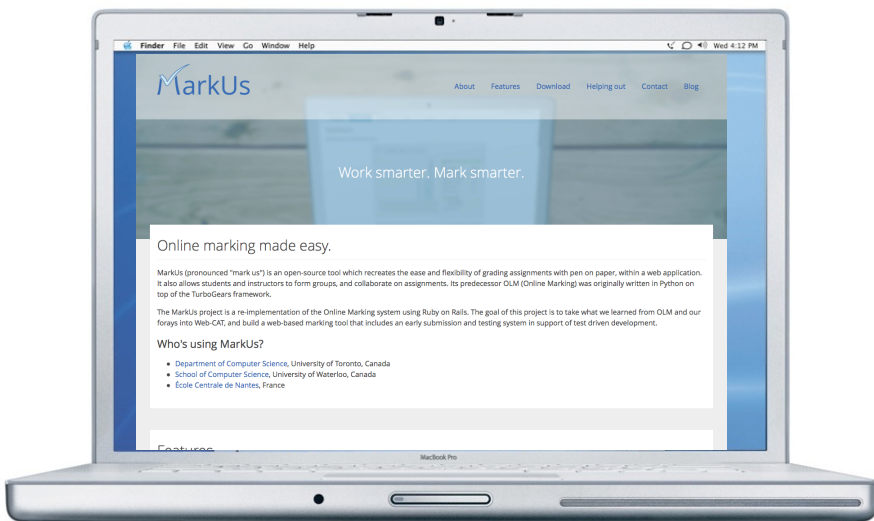
What makes the world wide web work?

Karen Reid

Simple Web Request



Response



The Request

- How do we tell the web server what we want?
- How do we even find the web server?
- How do the web server and browser talk to each other?

HTTP Request

request

```
GET / HTTP/1.1  
Host: markusproject.org  
...
```



reply

```
HTTP/1.1 200 OK  
Date: Tue, 13 Mar 2017  
Server: Apache/2.2.22(Debian)  
Content-Type: text/html
```

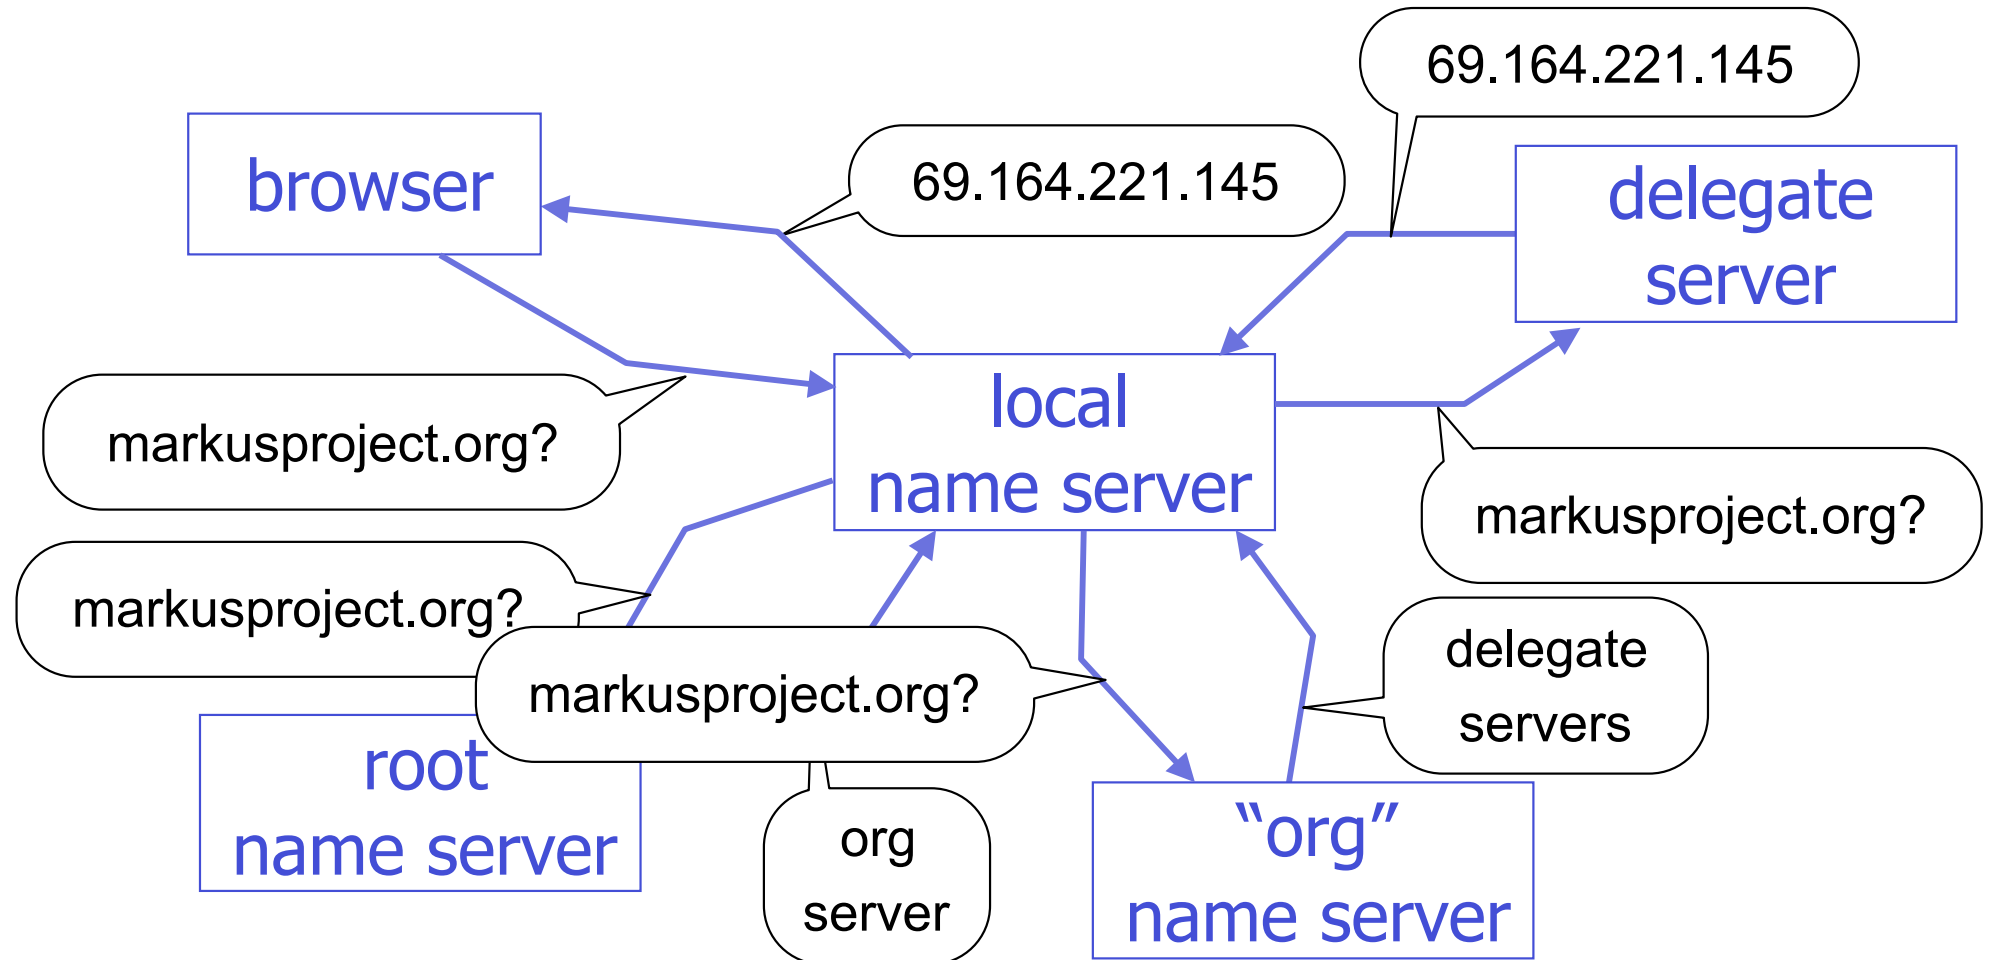


How do we find the server?


- Every computer on the Internet has an Internet address.
- Called an IP address (Internet Protocol)
- An IP address is 4 numbers separated by dots.

markusproject.org = 69.164.221.145

Domain Name Servers



This is getting complicated!



Number of
messages?

8 (or so)

Now what?

- Okay, we have the address.
- What do we do with it?
- Let's look at how two computers communicate.
- HTTP is a high-level protocol
- HTTP is specific to the web.
- Computers communicate for many reasons.

Protocols

- Computers use several layers of general protocols to communicate.
- To understand why these layers are important, think about how a company sends you an invoice for a purchase.

Protocols

Invoice:

Customer: Karen Reid
Order No: 5379

Qty:		Unit Price	Total
1	Athalon	219.00	219.00
2	128 MB	149.95	299.90
	Subtotal		518.90
	Tax		77.84
	TOTAL		596.74

Karen Reid

Feb 18, 2001


Payable to: CPUS are us \$596.74
Five hundred ninety six 74/100

CPUS are us

Karen Reid
Dept. of Computer Science
University of Toronto

Karen Reid

CPUS are us
0 College Street
Toronto Ontario M5S 3G4

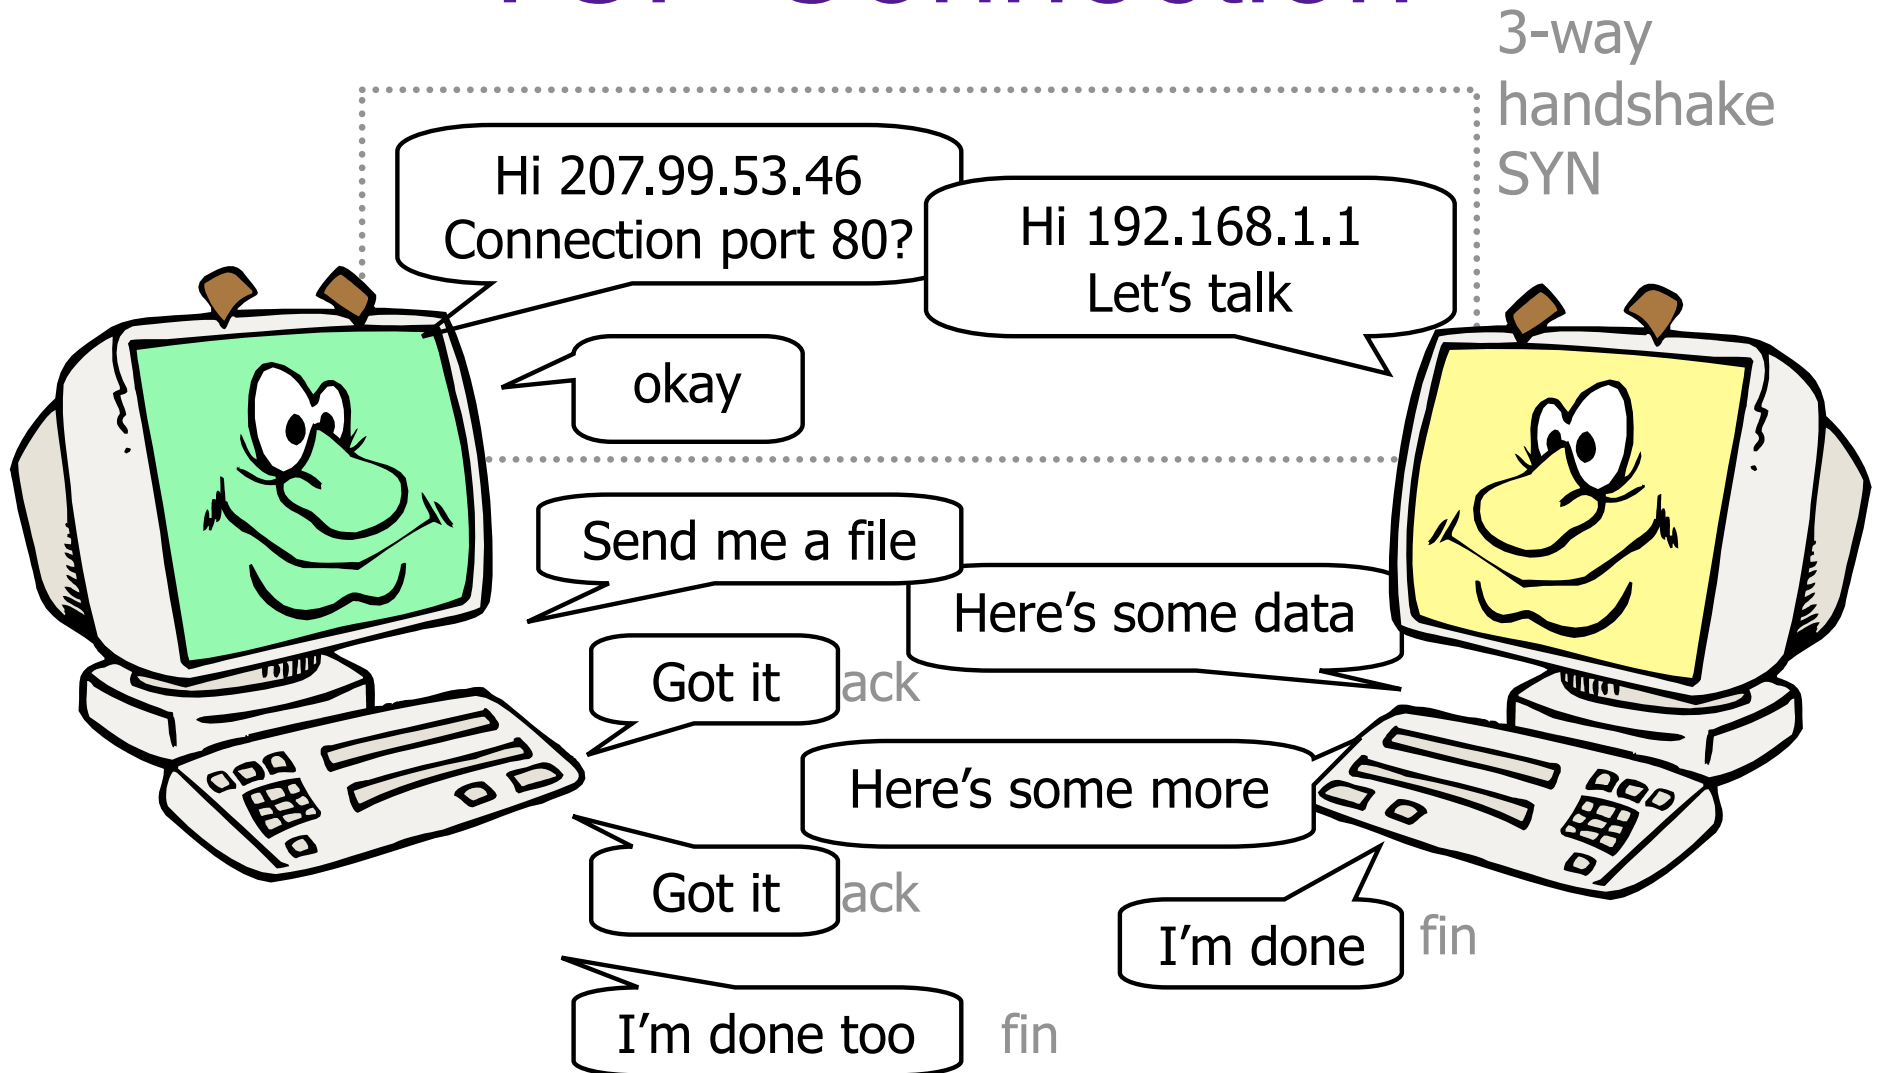


TCP/IP

- Transmission Control Protocol.
- Tells us how to package up the data.

source address		dest. address
bytes	ack	port
data		


TCP Connection



Packaging up the data

- make packets

01100111001001
00100010001111
10100010111







- Each TCP packet is given a header
 - sequence number
 - checksum

101010001
111010101
100110010
110101111
001011011

101010001
111010101
100110010
110101111
001011011

101010001
111010101
100110010
110101111
001011011

101010001
111010101
100110010
110101111
001011011



- put in an IP envelope with another header


To
207.99.53.46

To
207.99.53.46

To
207.99.53.46

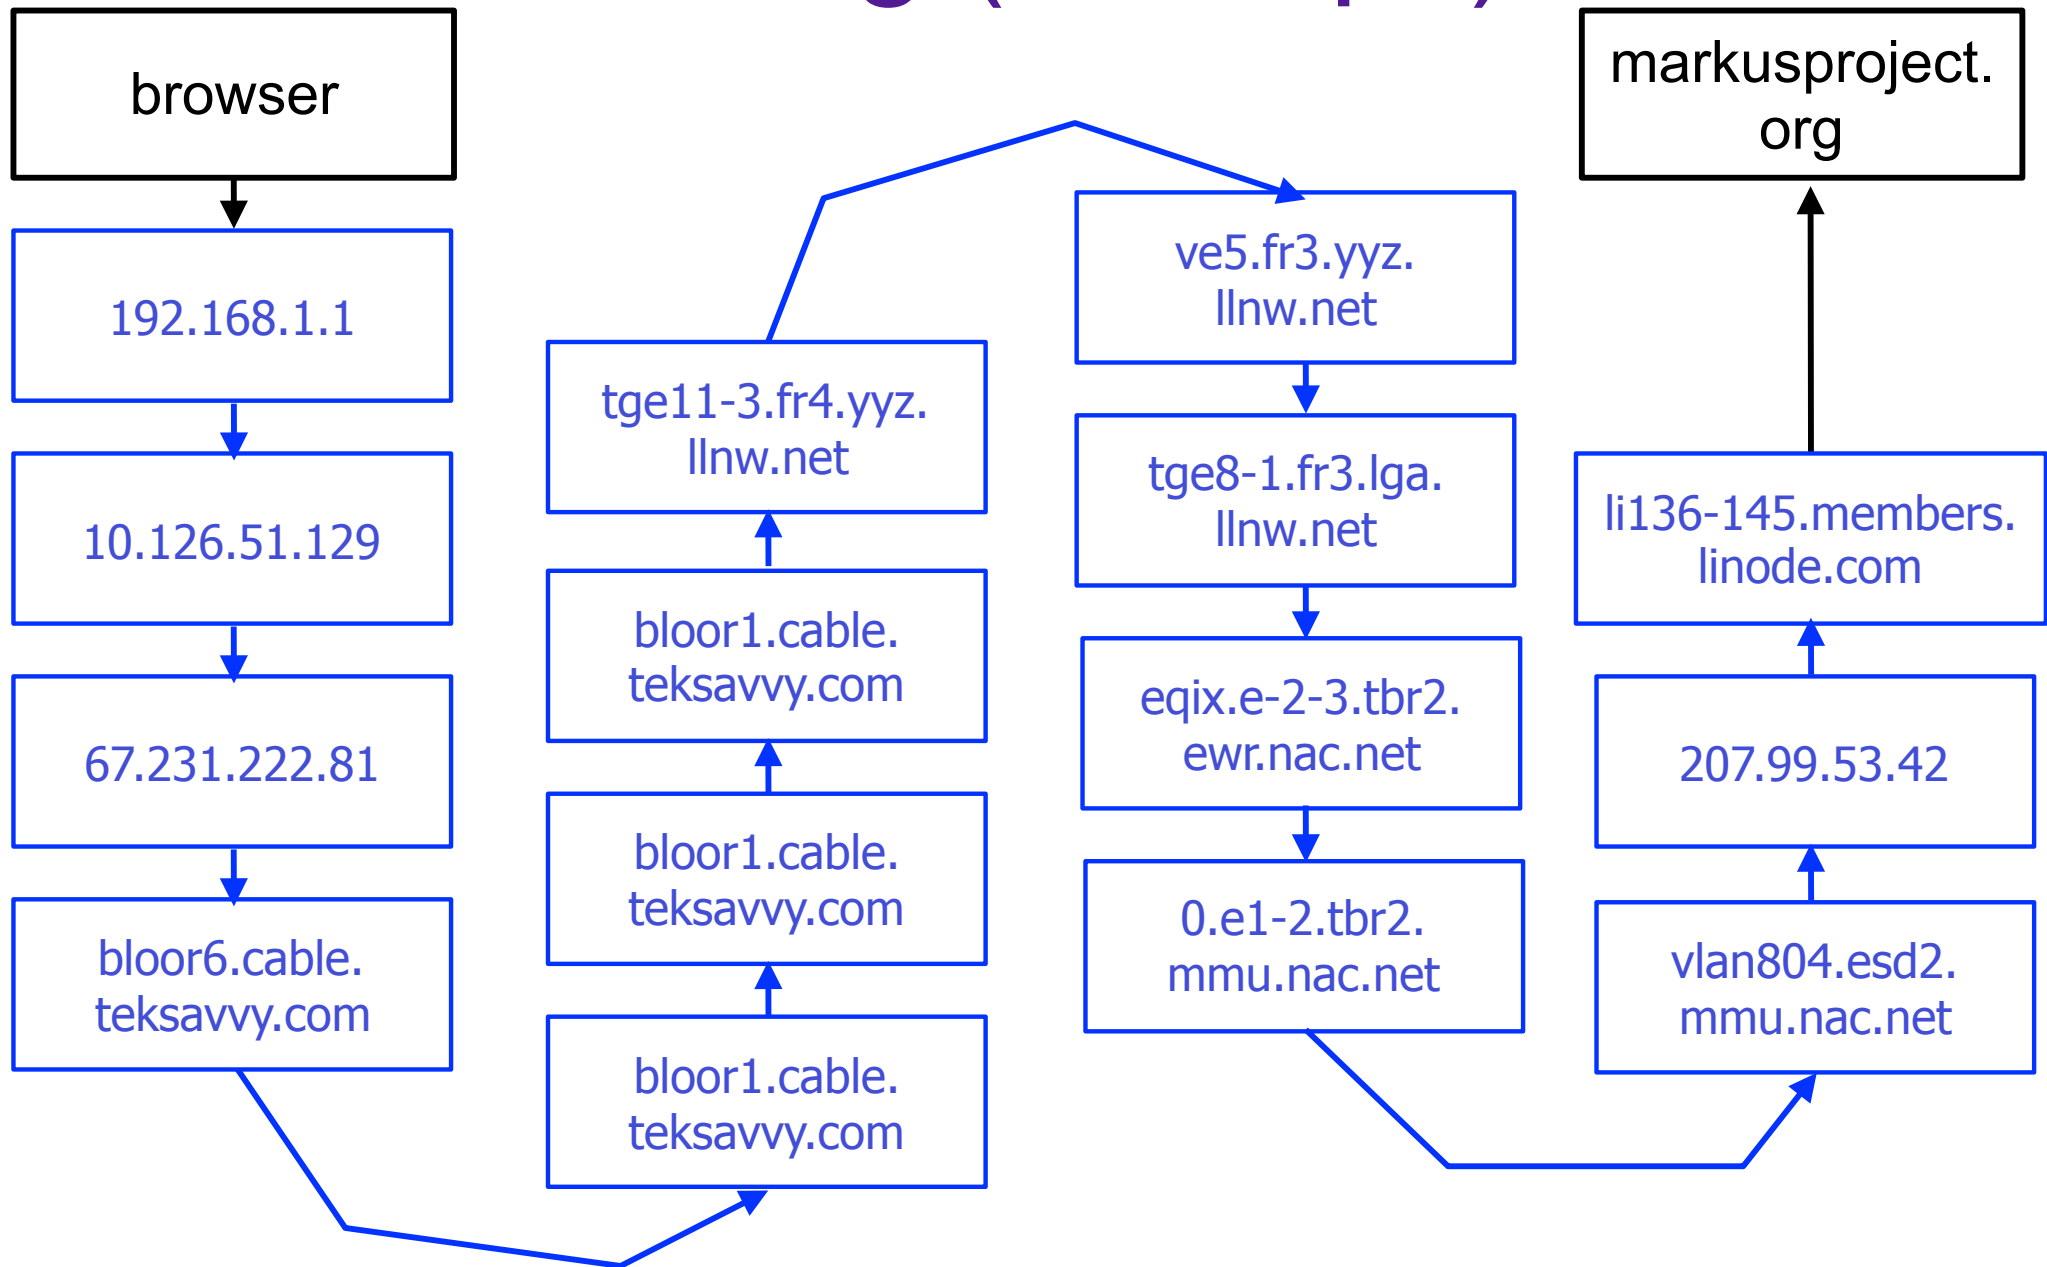
To
207.99.53.46

The Big Picture

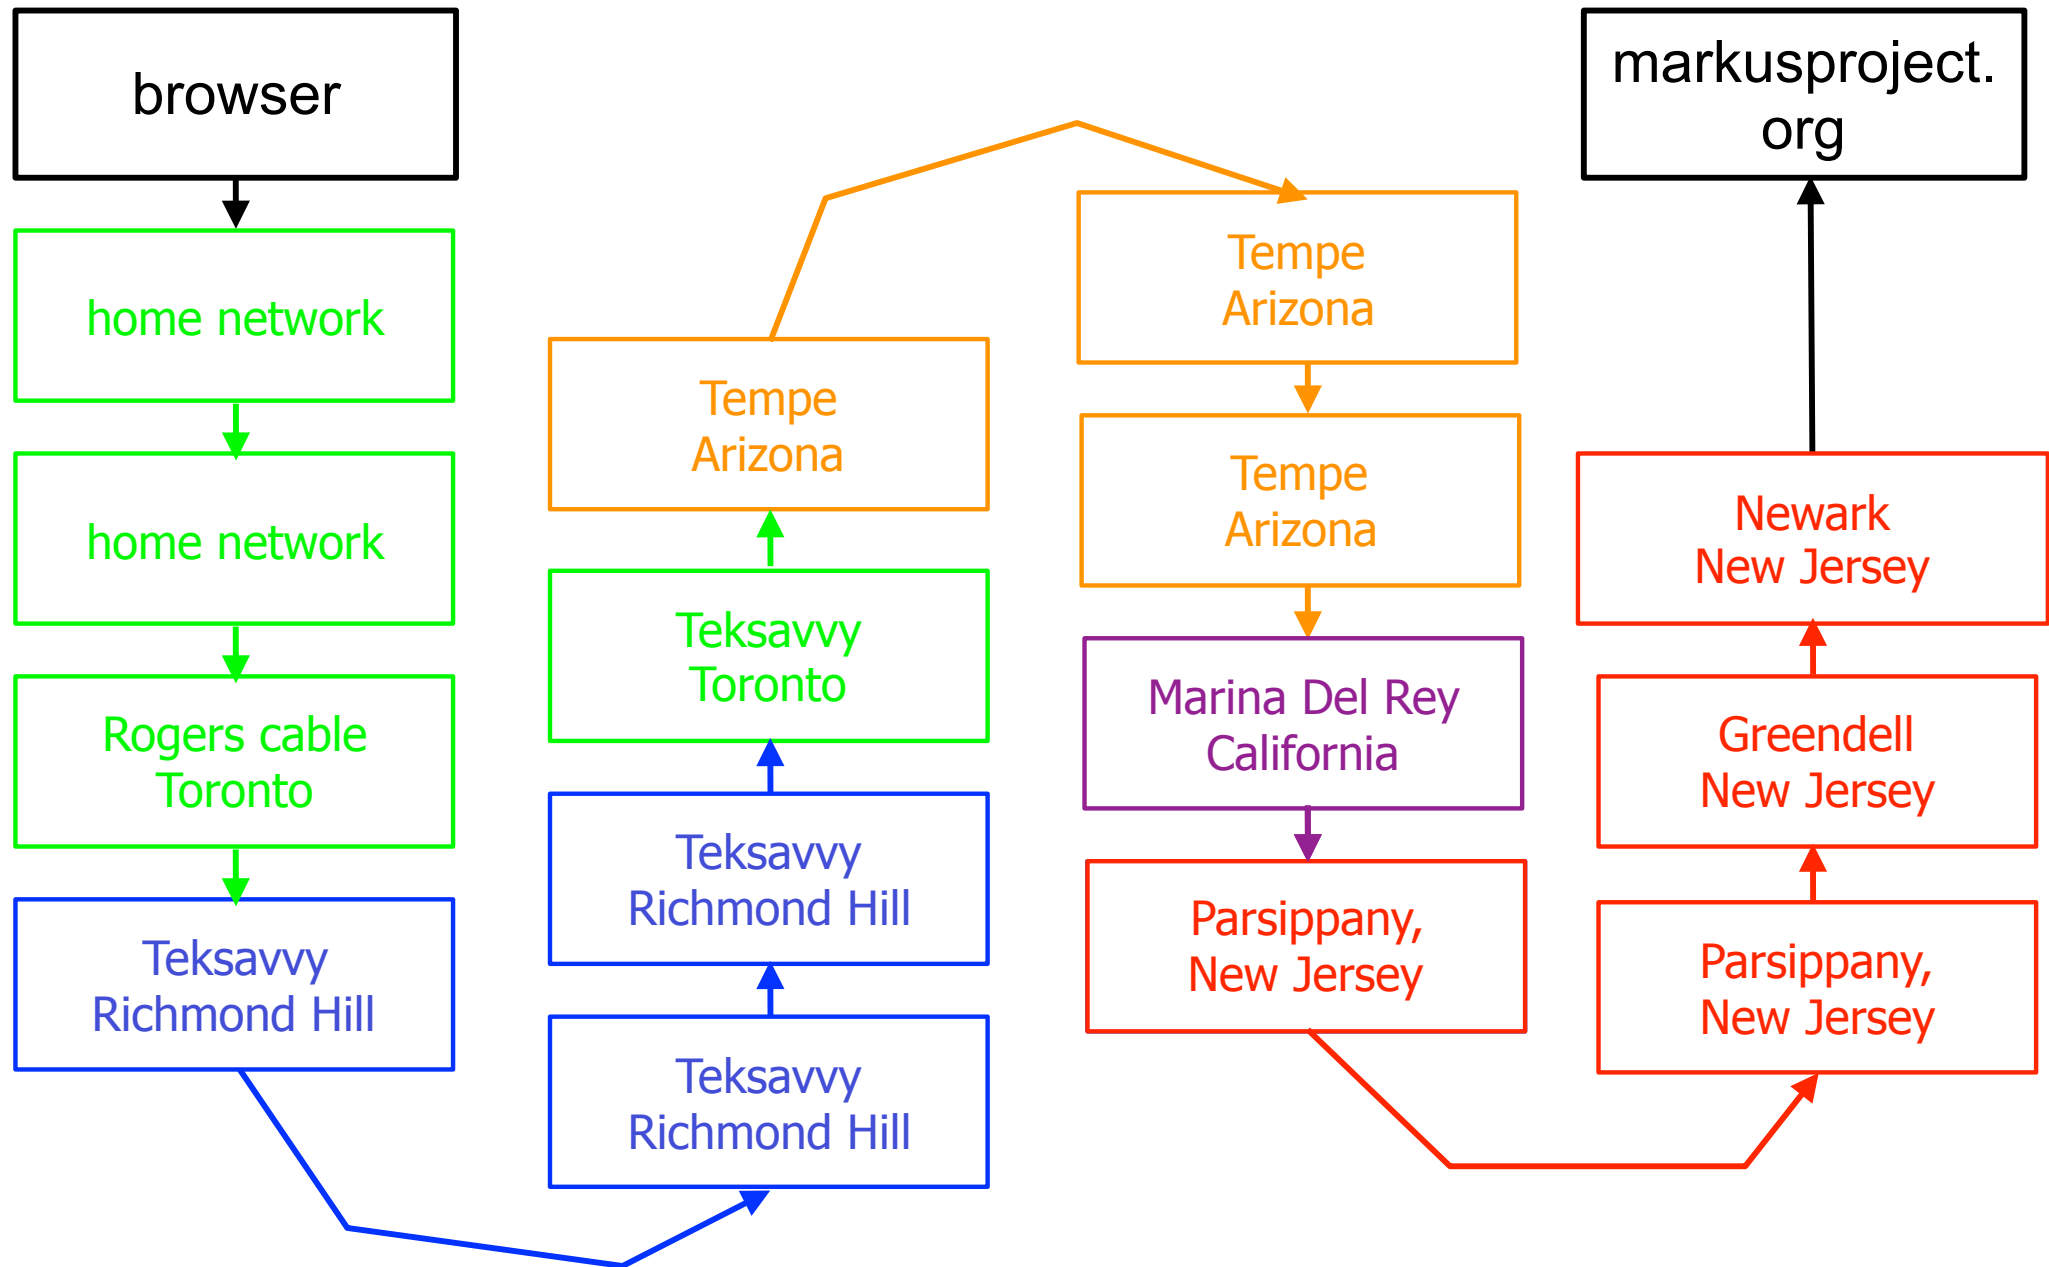


- **Client-Server model:** a client process wants to talk to a server process
- Client must find server - **DNS lookup**
- Client must find process on server - **ports**
- Finally **establish a connection** so two processes can talk

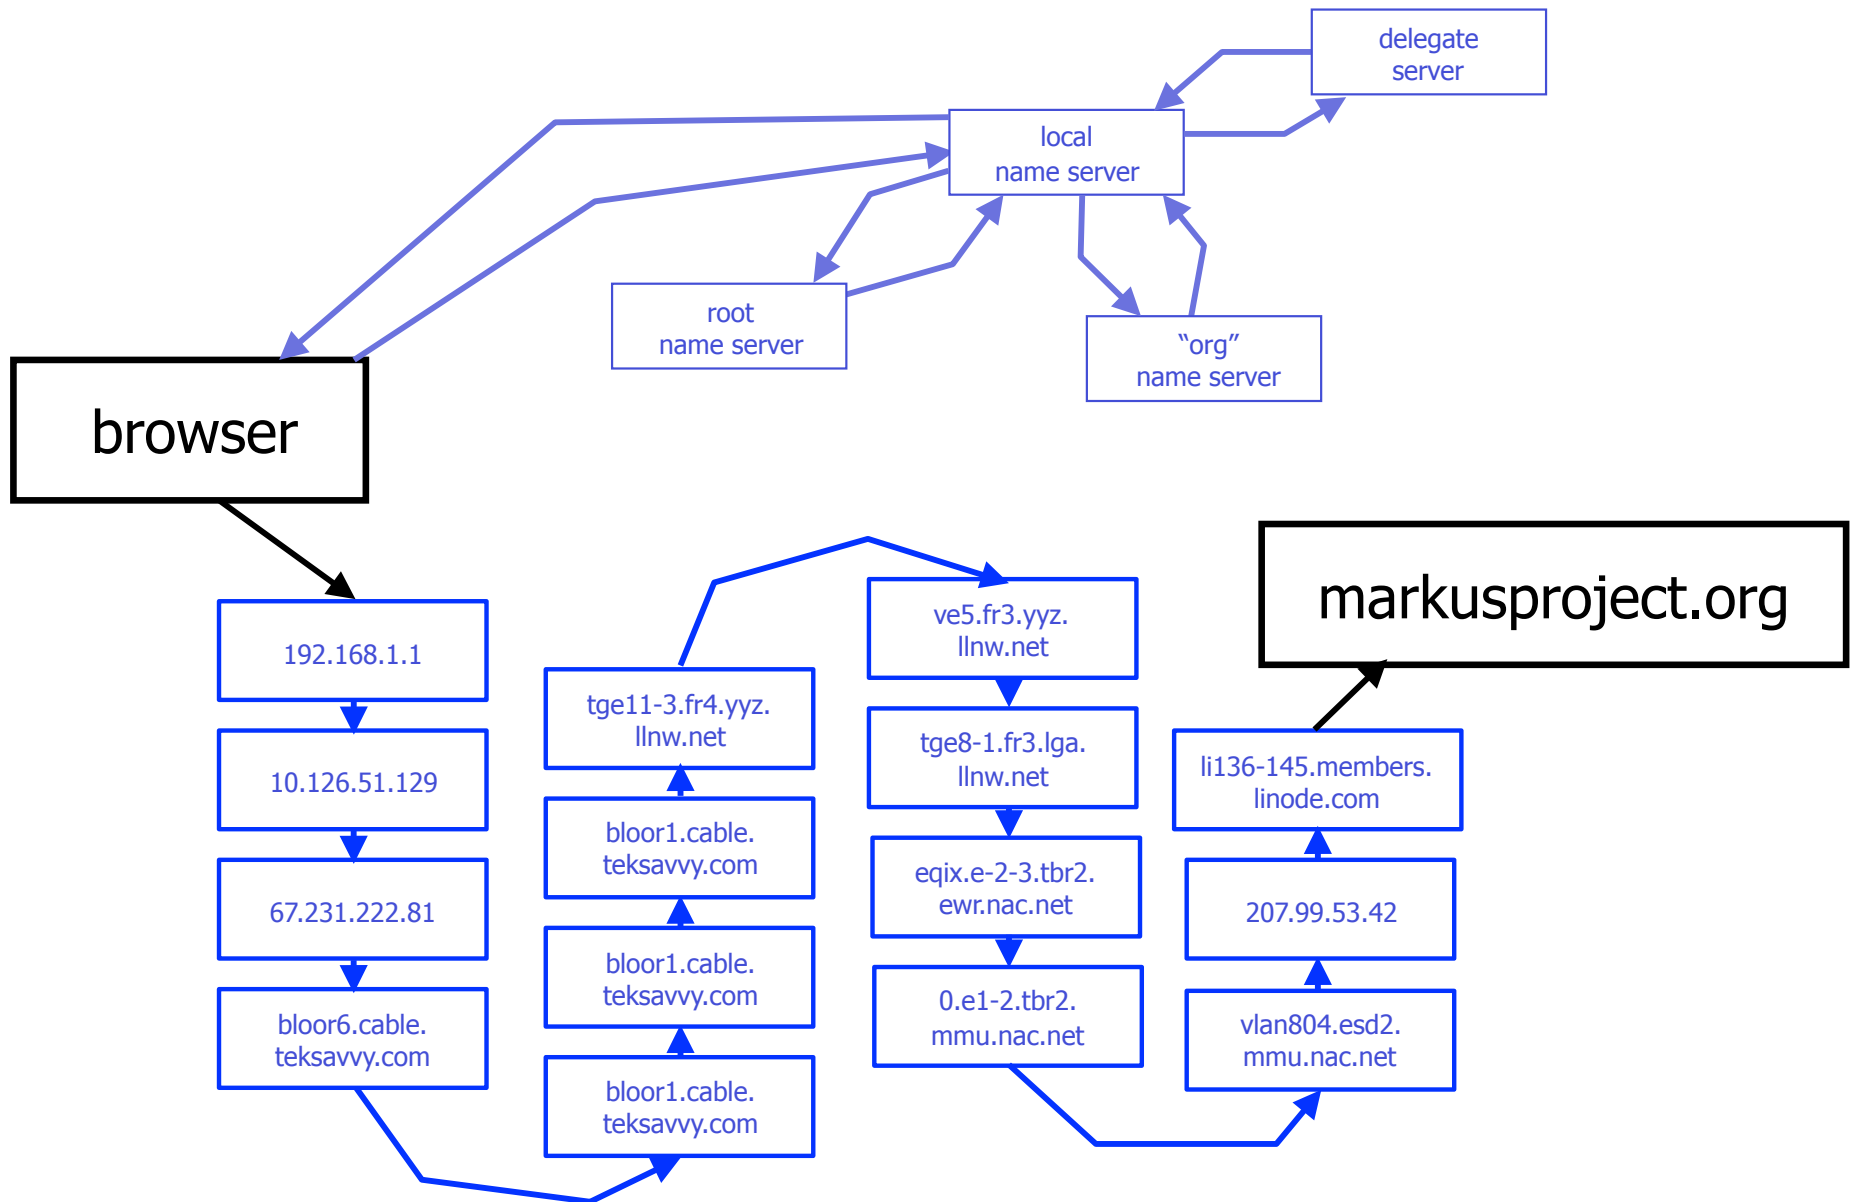
Routing (15 hops)



7 cities, 5 states/prov, 2 countries



Putting it together



How many messages?

- It depends on the size of the web page
- The web page that appears for markusproject.org is less than 30 Kbytes
- If the web page is 30 Kbytes (small!) it will likely be broken up into ~20 IP packets.


$$\begin{aligned} &8 \text{ (DNS)} + 20 * 15 \text{ hops} \\ &= 308 \text{ messages} \end{aligned}$$

When something goes wrong


- A packet might not arrive
 - traffic overload
 - bit corruption
- Receiver asks for missing packets to be resent.
- Want to send data as fast as possible.
- But sending too fast wastes resources.

TCP Congestion Control

- Window-based:
 - some number of packets allowed to be sent and not ack'd
 - as successful ack's arrive, grow window
 - if packet loss is detected, cut window size



TCP Congestion Control



All we did was click on a link...

Take aways

- The web today is made up of complex layers of software
- No one person, organization, or company could have created it in isolation
- We can understand it because we can study one layer at a time
- We can create new things by building on top of existing layers