


Status

- A4 will be out later today
 - No extensions are possible
 - You may use late penalties
- Exercises 10 and 11 are posted
 - We will take the best 9 out of 10 marks for labs.
 - Exercise 7 was not for credit.
 - This sort of makes exercise 11 optional

I/O Multiplexing

Kerrisk Ch 63

The problem



c2 send data to server.. read is blocked on c1, so c2 data is stuck in pipe / connection

When reading from multiple sources, blocking on one of the sources could be bad.

An example of denial of service.

One solution: one process for every client. What are the pros and cons of this solution?

1. accept call generate socket ,
2. incoming connection is being forked and accept data in a separate process
3. may be problematic
 - a. too many child process
 - b. shared data between child processes

Another way to look at the problem

Server

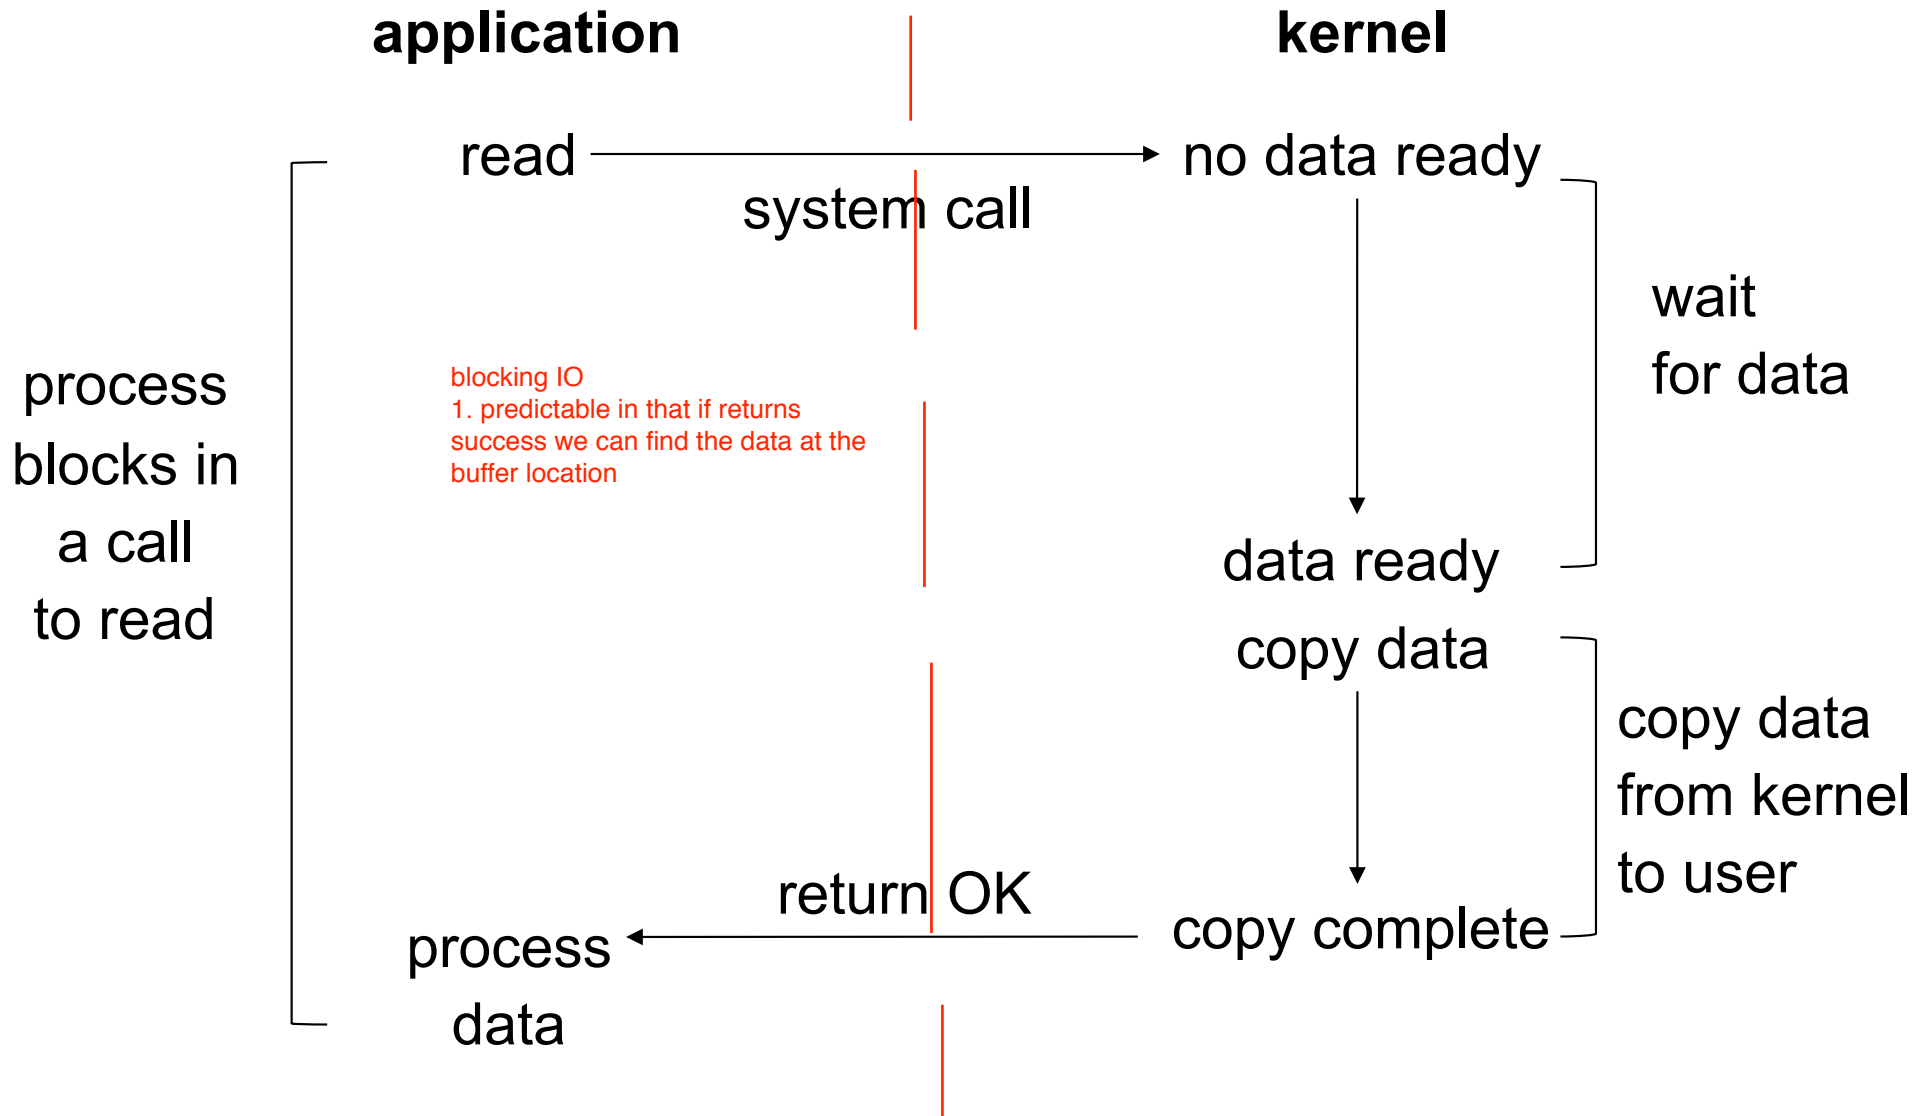
```
while(1)
    accept a new connection
    for each existing connection
        read
        write
```

Which of the system calls might block indefinitely?
read and accept

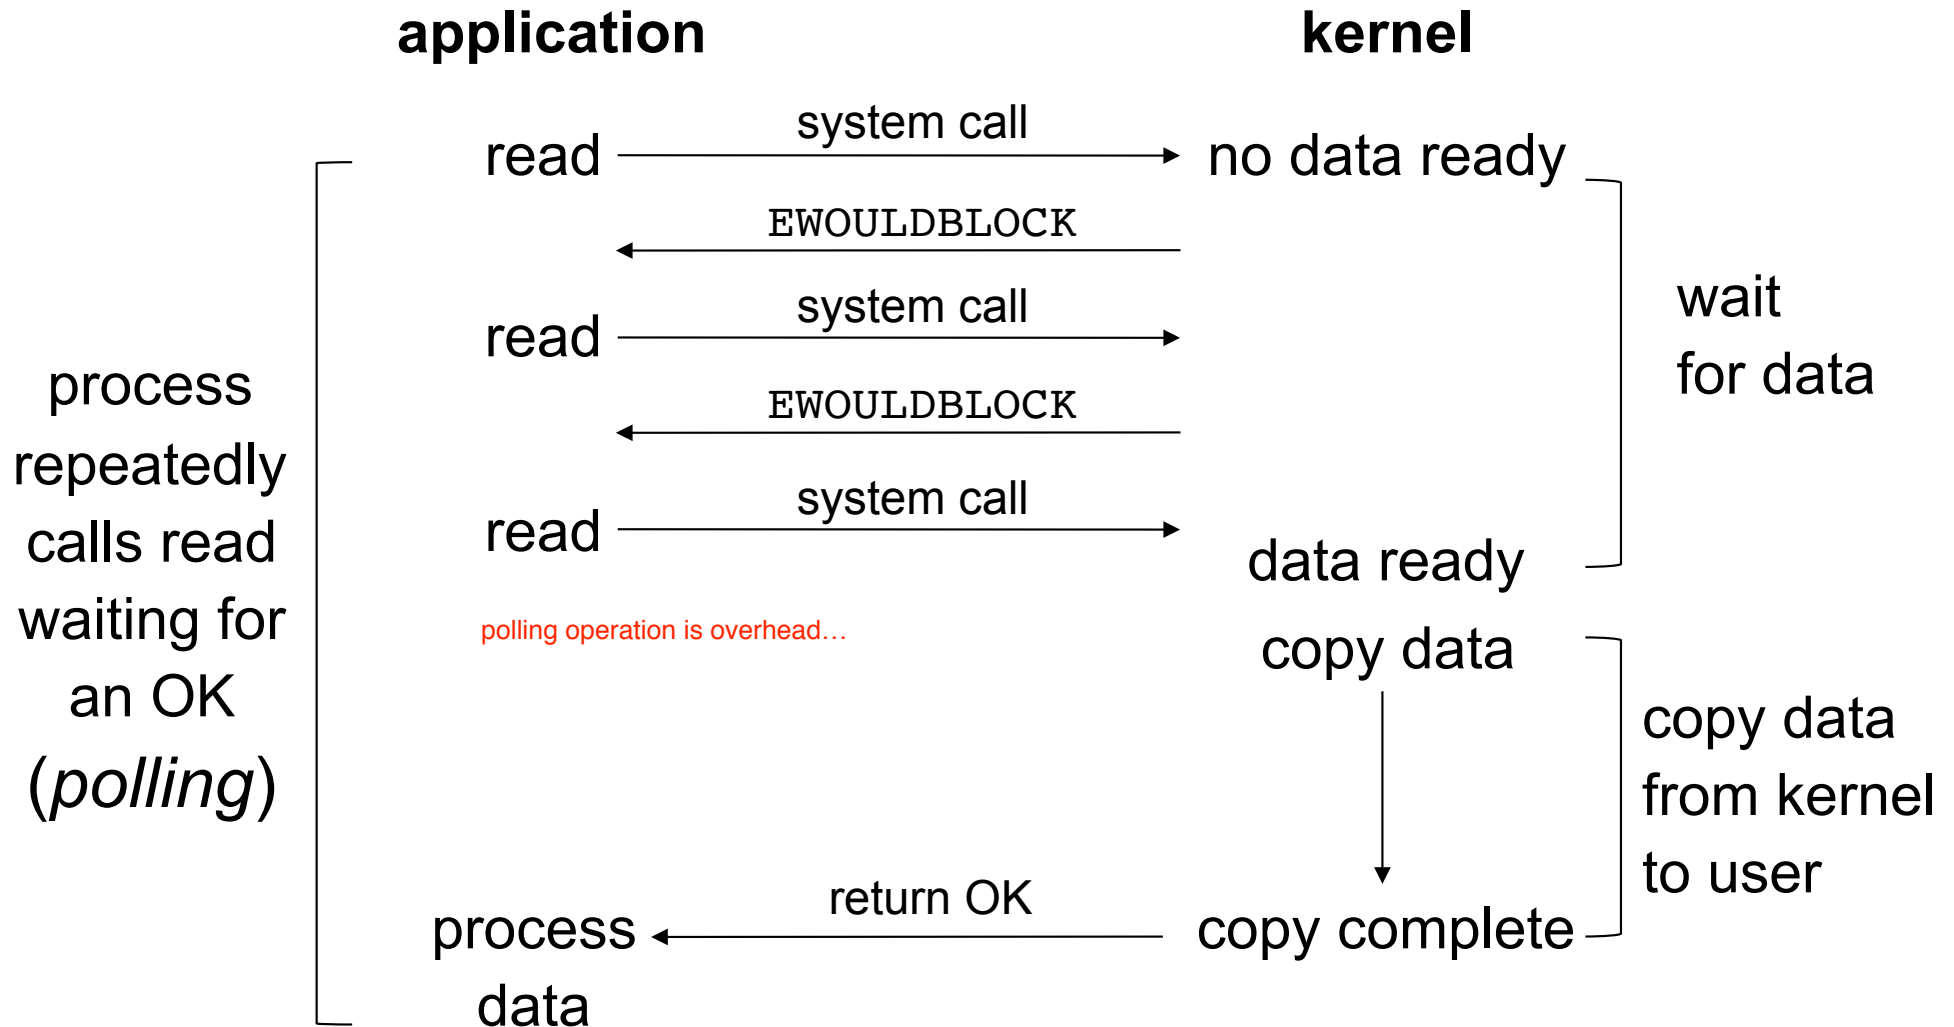
So what happens if there is only one connection?

blocked in accept call ...

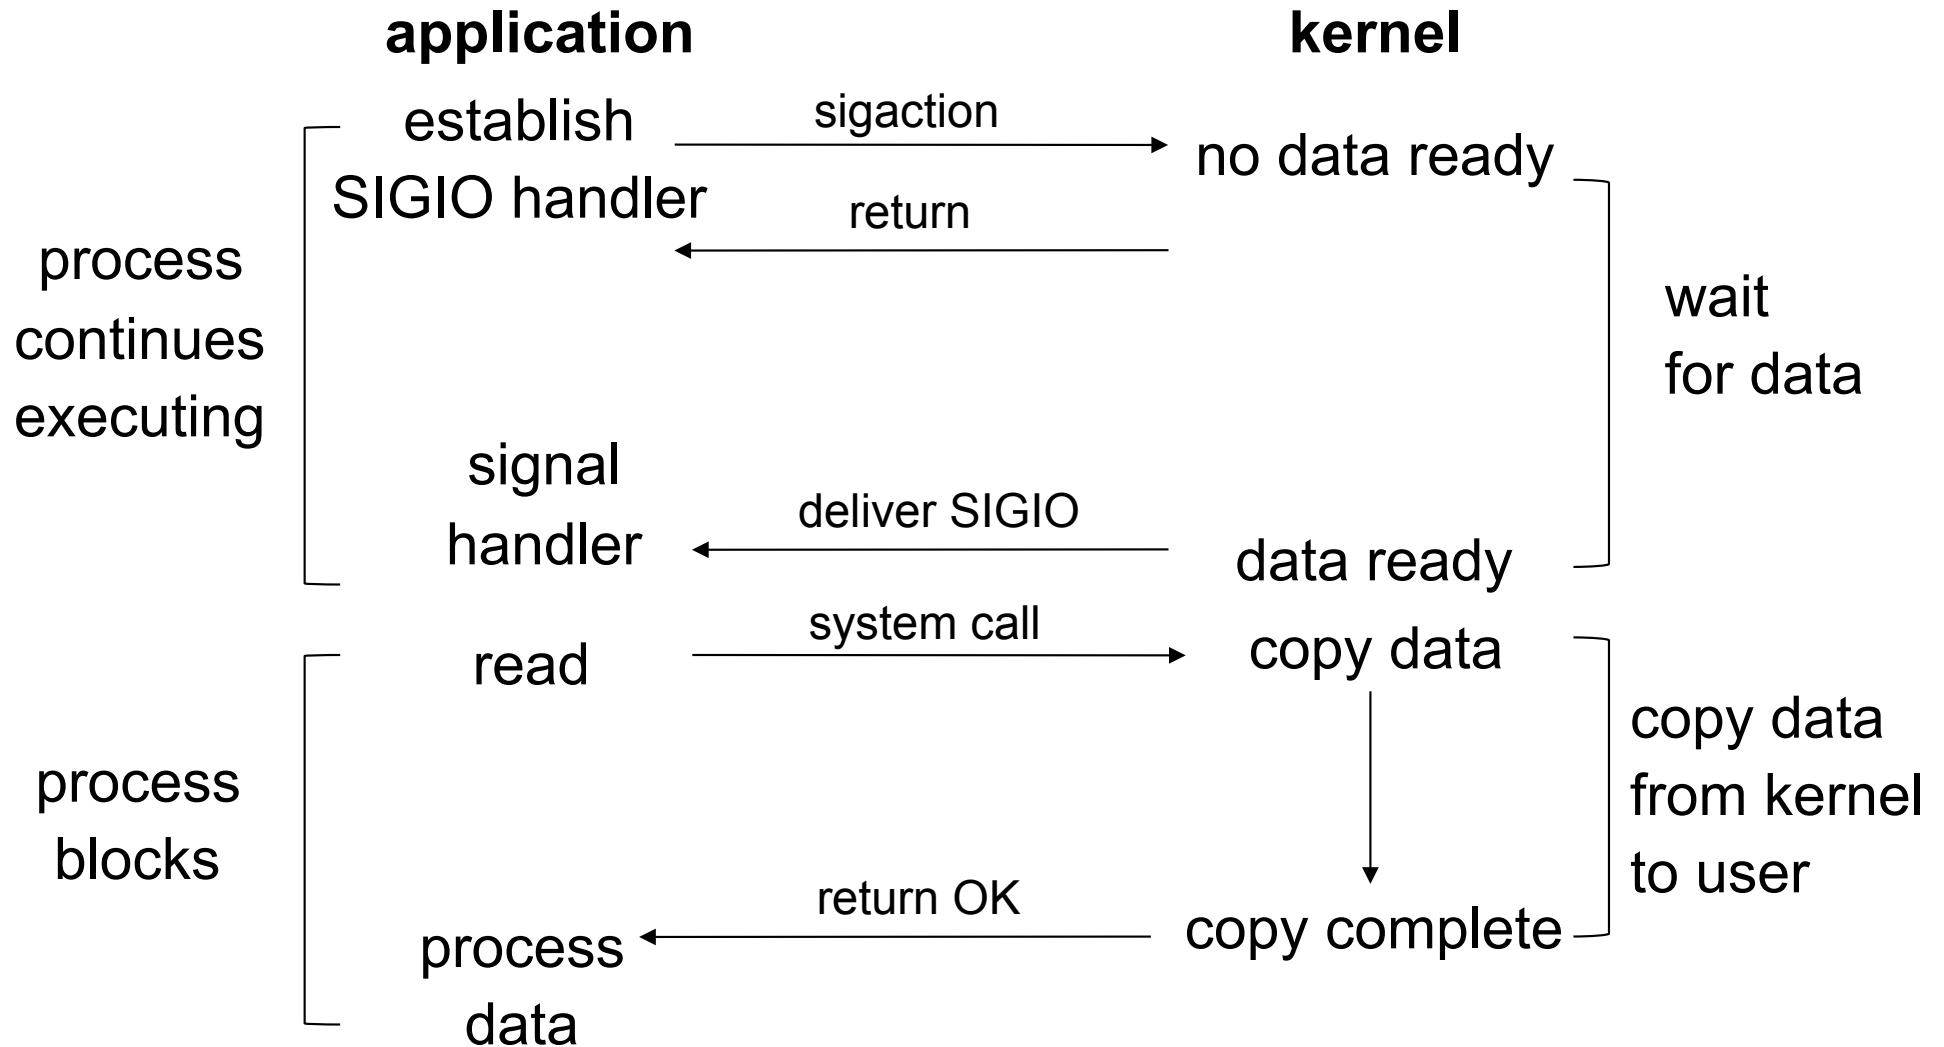
Blocking I/O Model



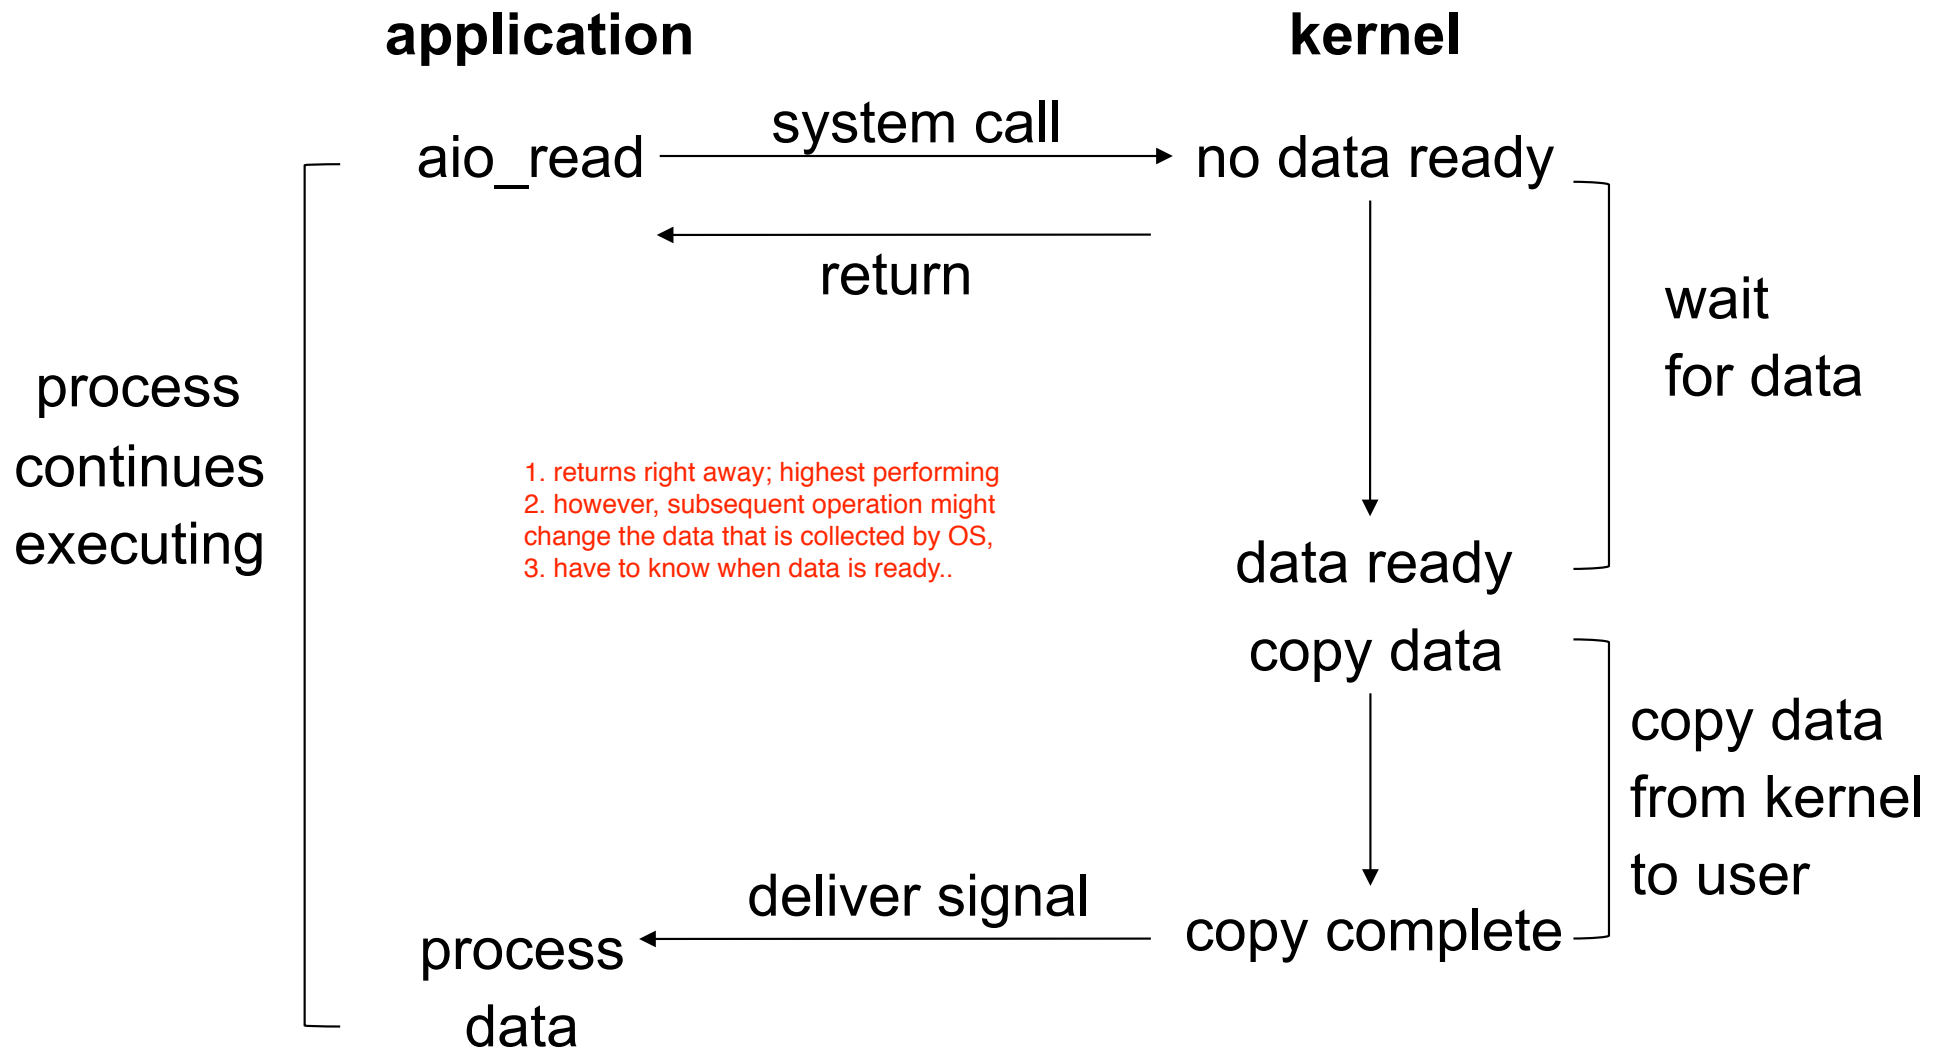
Nonblocking I/O Model



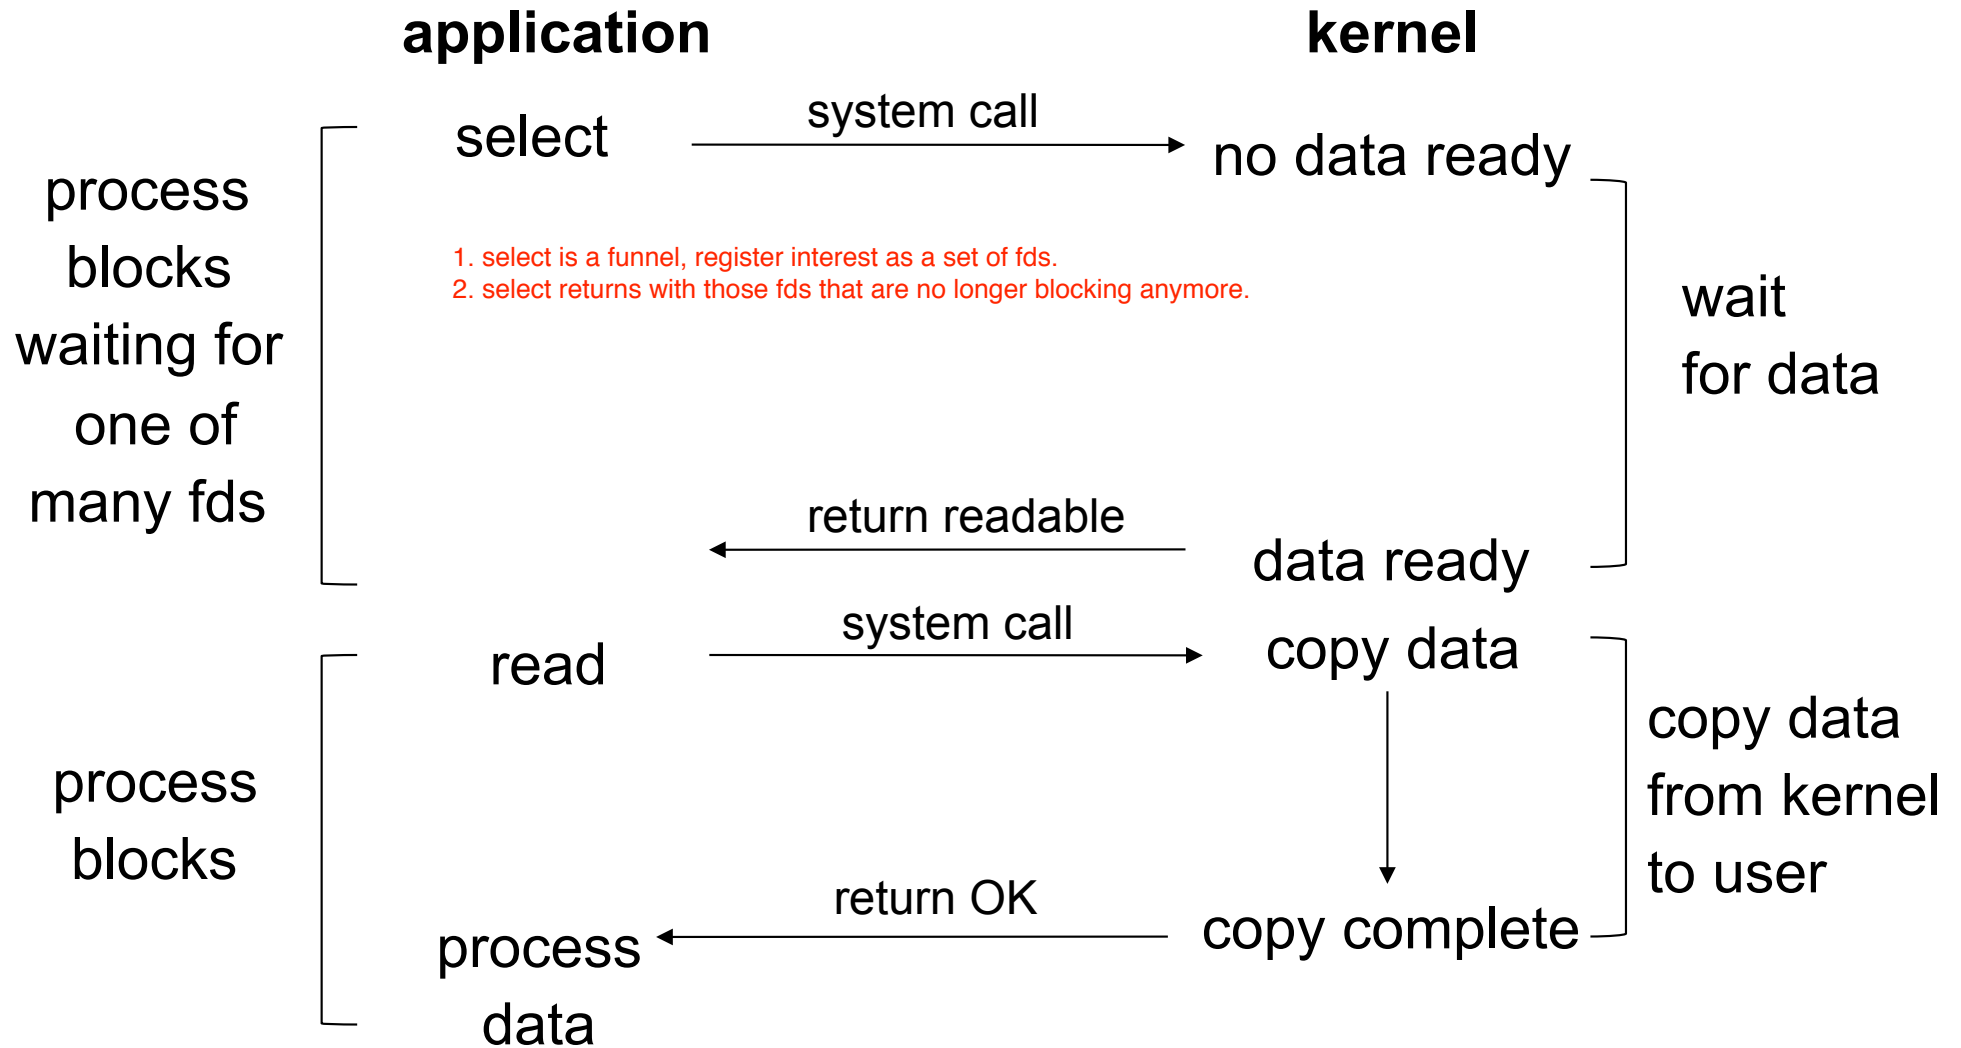
Signal Driven I/O Model



Asynchronous I/O Model



I/O Multiplexing Model



write does usually block for
period to write data.
unusual to wait indefinitely

Ch. 63.2.1

select()

```
int select(int maxfdp1,  
           fd_set *readset,  
           fd_set *writerset,  
           fd_set *exceptset,  
           const struct timeval *timeout);
```

how far into set we have to search

A call to select returns when one of the file descriptors in one of the sets is ready for I/O.

If timeout is not NULL, then select returns when a descriptor is ready or timeout time has passed.

If timeout is 0, select returns immediately after checking descriptors.

Readiness

Ready to read when

- there is data in the receive buffer to be read

- end-of-file state on file descriptor

- the socket is a listening socket and there is a connection pending

- a socket error is pending

Ready to write when

- there is space available in the write buffer

- a socket error is pending

Exception condition pending when

- TCP out-of-band data

We are typically interested in when bytes are available to be read, but sometimes we use select on write or exception sets

select timeout

- The timeout specifies how long we're willing to wait for a fd to become ready

```
struct timeval {  
    long    tv_sec;        /* seconds */  
    long    tv_usec;       /* microseconds */  
};
```

- If timeout is NULL, wait forever (or until we catch a signal)
- If timeout is zero, test and return immediately
- Otherwise wait up to specified timeout
- `select` returns when a fd ready or we timeout

Descriptor sets

Typically implemented as an array of integers where each bit corresponds to a descriptor (except in Windows).


Implementation is hidden in the `fd_set` data type
`FD_SETSIZE` is the number of descriptors in the data type

`maxfdp1` specifies the number of descriptors to test

Macros:

```
void FD_ZERO(fd_set *fdset);  
void FD_SET(int fd, fd_set *fdset);  
void FD_CLR(int fd, fd_set *fdset);  
int  FD_ISSET(int fd, fd_set *fdset);
```

Descriptor sets



select search linearly, so efficient
to tell select where to stop

| | fd0 | fd1 | fd2 | fd3 | fd4 | fd5 | fd6 | |
|--------|-----|-----|-----|-----|-----|-----|-----|--|
| allset | 0 | 0 | 0 | 1 | 1 | 0 | 1 | |

$$\text{maxfd} + 1 = 7$$

After select:

| | fd0 | fd1 | fd2 | fd3 | fd4 | fd5 | fd6 | |
|------|-----|-----|-----|-----|-----|-----|-----|--|
| rset | 0 | 0 | 0 | 1 | 0 | 0 | 0 | |

select modify set, only fds that has something to read is in the set

select example

```
fd_set rfd;
struct timeval tv;
int retval;

FD_ZERO(&rfd); /* Watch stdin (fd 0) for input */
FD_SET(STDIN_FILENO, &rfd);
tv.tv_sec = 5; /* Wait up to five seconds. */
tv.tv_usec = 0;
retval = select(1, &rfd, NULL, NULL, &tv);
if (retval == -1)
    perror("select()");
else if (retval > 0)
    printf("Data is available now.\n");
    /* FD_ISSET(0, &rfd) will be true, can use read() */
else
    printf("No data within five seconds.\n");
```

```

for( ; ; ) {
    rset = allset;
    nready = Select(maxfd+1, &rset ,NULL,NULL,NULL);
    if(FD_ISSET(listenfd, &rset)) {
        connfd = Accept(listenfd, &caddr, &crlen);
        for(i = 0; i < FD_SETSIZE; i++)
            if(client[i] < 0) {
                client[i] = connfd; break;
            }
        FD_SET(connfd, &allset);
        if(connfd > maxfd) maxfd = connfd;
    }
    for(i = 0; i <= maxi; i++) {
        if(sockfd = client[i]) < 0) continue;
        if(FD_ISSET(sockfd, &rset))
            Read(sockfd, line, MAXLINE);
    }
}

```



```

for( ; ; ) {
    rset = allset;
    nready = Select(maxfd+1, &rset ,NULL,NULL,NULL);
    if(FD_ISSET(listenfd, &rset)) {
        connfd = Accept(listenfd, &caddr, &crlen);
        for(i = 0; i < FD_SETSIZE; i++)
            if(client[i] < 0) {
                client[i] = connfd; break;
            }
        FD_SET(connfd, &allset);
        if(connfd > maxfd) maxfd = connfd;
    }
    for(i = 0; i <= maxi; i++) {
        if(sockfd = client[i]) < 0) continue;
        if(FD_ISSET(sockfd, &rset))
            Read(sockfd, line, MAXLINE);
    }
}

```

```

for( ; ; ) {
    rset = allset;
    nready = Select(maxfd+1, &rset ,NULL,NULL,NULL);
    if(FD_ISSET(listenfd, &rset)) {
        connfd = Accept(listenfd, &caddr, &crlen);
        for(i = 0; i < FD_SETSIZE; i++)
            if(client[i] < 0) {
                client[i] = connfd; break;
            }
        FD_SET(connfd, &allset);
        if(connfd > maxfd) maxfd = connfd;
    }
    for(i = 0; i <= maxi; i++) {
        if(sockfd = client[i]) < 0) continue;
        if(FD_ISSET(sockfd, &rset))
            Read(sockfd, line, MAXLINE);
    }
}

```

End of Line

- There are two characters that determine end-of-line
 - Carriage return (CR, \r, ^M)
 - Line feed (LF, \n)
- Early operating systems defined their own conventions using one or both of CR and LF.
 - Unix: LF,
 - DOS/Windows: CR LF
 - Mac classic: CR

Network Line Ending

- Transferring data between machines with different operating systems, means deciding on a common line ending.
- CR LF is the standard
- (Of course it is possible with regular expression matching to mostly ignore this issue, but still better to conform.)