


Programming with Sockets in C

- The slides by themselves will not be sufficient to learn how to write socket code.
- If you did not attend class, then you will want to review the relevant chapters in Kerrisk (ch 56, 59, 60, 61.1 63.1, 63.2(.1))
- The provided example code is also a good starting point.

TCP: Three-way handshake



TCP Server

socket()

bind()

listen()

accept()

block until connection
from client

read()

write()

close()

Connection establishment
(3-way handshake)

data transfer

end-of-file notification

TCP Client

socket()

connect()

write()

read()

close()

Connection-Oriented

Server

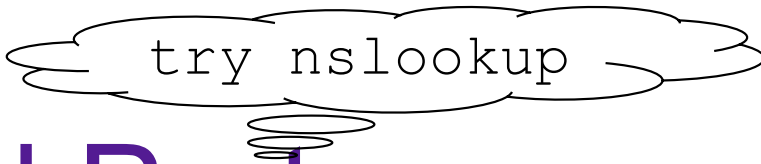
- Create a socket: `socket()`
- Assign a name to a socket: `bind()`
- Establish a queue for connections: `listen()`
- Get a connection from the queue: `accept()`

Client

- Create a socket: `socket()`
- Initiate a connection: `connect()`

Socket Types

- Two main categories of sockets
 - UNIX domain: both processes on the same machine performace better
 - INET domain: processes on different machines
- Three main types of sockets:
 - SOCK_STREAM: the one we will use for connection oriented stream
 - SOCK_DGRAM: for connectionless sockets
 - SOCK_RAW for DIY socket



try nslookup

Addresses and Ports

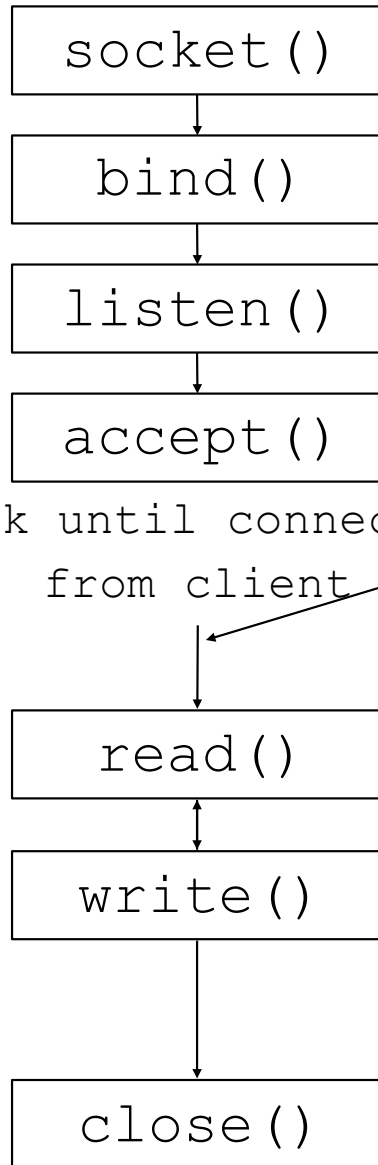
- A socket pair is the two endpoints of the connection.
- An endpoint is identified by an IP address and a port.
- IPv4 addresses are 4 8-bit numbers:
 - 128.100.31.198 = greywolf
 - 128.100.31.203 = redwolf
- Ports
 - because multiple processes can communicate with a single machine we need another identifier.

0 ~ 255


More on Ports

- Well-known ports: 0-1023
 - 80 = http
 - 21 = ftp
 - 22 = ssh
 - 25 = smtp (mail)
 - 23 = telnet
 - 194 = irc
- Registered ports: 1024-49151
 - 2709 = supermon
 - 26000 = quake
 - 3724 = world of warcraft
- Dynamic (private) ports: 49152-65535

TCP Server



TCP Client



Connection establishment
(3-way handshake)

data transfer

end-of-file notification

Server side

```
int socket(int family, int type,  
           int protocol);
```

- **family specifies protocol family:**

- PF_INET – IPv4
- PF_LOCAL – Unix domain

- **type**

- SOCK_STREAM, SOCK_DGRAM, SOCK_RAW

- **protocol**

- set to 0 except for RAW sockets

- **returns a socket descriptor**

bind to a name

```
int bind(int sockfd,  
        const struct sockaddr *servaddr,  
        socklen_t addrlen);
```

- **sockfd – returned by socket**
- **struct sockaddr_in {** cheap way of getting different socket, gives padding -> for different socket type
 short sin_family; /*PF_INET */
 u_short sin_port;
 struct in_addr sin_addr;
 char sin_zero[8]; /*filling*/
};
- **sin_addr can be set to INADDR_ANY to communicate on any network interface.**

Set up queue in kernel

creates holding area...

```
int listen(int sockfd, int backlog)
```

- after calling `listen`, a socket is ready to accept connections the queue holds a queue of request
- prepares a queue in the kernel where partially completed connections wait to be accepted.
- `backlog` is the maximum number of partially completed connections that the kernel should queue.

max length of queue

Complete the connection

```
int accept(int sockfd,  
           struct sockaddr *cliaddr,  
           socklen_t *addrlen);
```

- blocks waiting for a connection (from the queue) process is suspended, is not going to resume until conditions satisfied
- returns a new descriptor which refers to the TCP connection with the client specifically the socket descriptor of same type as `sockfd`
- `sockfd` is the listening socket socket created with `socket()` in local process
- `cliaddr` is the address of the client incoming connection
- reads and writes on the connection will use the socket returned by `accept`

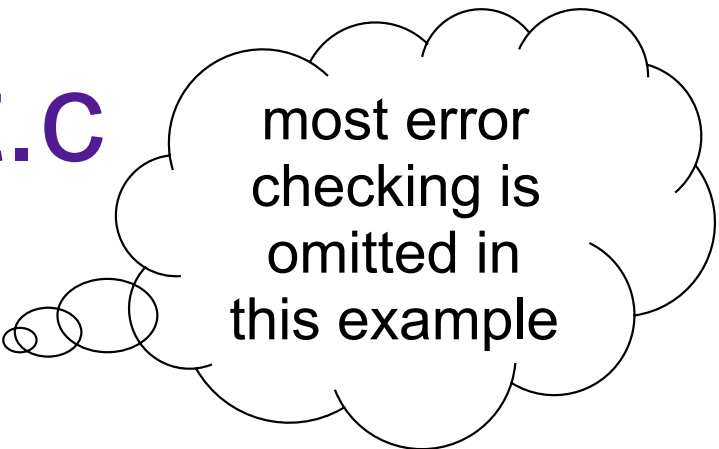
Client side

- `socket()` – same as server, to say “how” we are going to talk

```
int connect(int sockfd,  
            const struct sockaddr *servaddr,  
            socklen_t addrlen);
```

- the kernel will choose a dynamic port and source IP address.
- returns 0 on success and -1 on failure setting `errno`.
- initiates the three-way handshake.

inetclient.c



most error
checking is
omitted in
this example

```
int soc;  
struct hostent *hp;  
struct sockaddr_in peer;
```

```
peer.sin_family = AF_INET;  
peer.sin_port = htons(PORT);  
/* fill in peer address */  
hp = gethostbyname(argv[1]);  
peer.sin_addr = *((struct in_addr *)hp->h_addr);  
/* create socket */  
soc = socket(PF_INET, SOCK_STREAM, 0);  
/* request connection to server */  
if (connect(soc, (struct sockaddr *)&peer, sizeof(peer))  
    == -1) {  
    perror("client:connect"); close(soc); exit(1);  
}  
write(soc, "Hello Internet\n", 16);  
read(soc, buf, sizeof(buf));  
printf("SERVER SAID: %s\n", buf);  
close(soc);
```

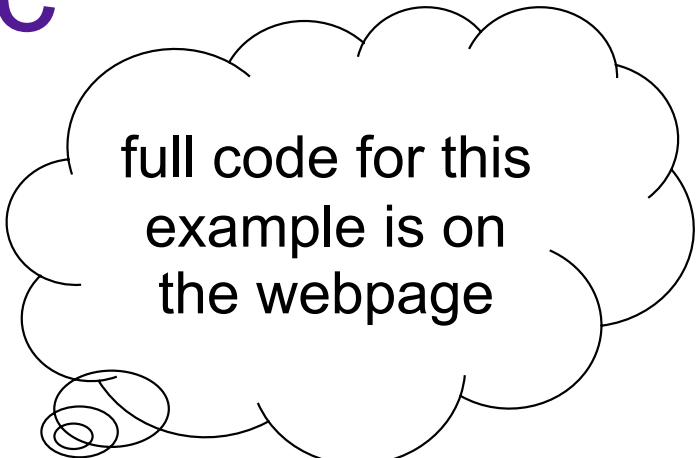
inetserver.c

```
struct sockaddr_in peer;  
struct sockaddr_in self;  
int soc, ns, k;  
int peer_len = sizeof(peer);
```

```
self.sin_family = PF_INET;  
self.sin_port = htons(PORT);  
self.sin_addr.s_addr = INADDR_ANY;  
bzero(&(self.sin_zero), 8);
```

```
peer.sin_family = PF_INET;  
/* set up listening socket soc */  
soc = socket(PF_INET, SOCK_STREAM, 0);  
if (soc < 0) {  
    perror("server:socket"); exit(1);  
}
```

```
if (bind(soc, (struct sockaddr *)&self, sizeof(self)) == -1) {  
    perror("server:bind"); close(soc); exit(1);  
}  
listen(soc, 1);  
...
```



full code for this
example is on
the webpage

inetserver.c (concluded)


```
/* ... repeated from previous slide ...
soc = socket(PF_INET, SOCK_STREAM, 0);
bind(soc, (struct sockaddr *)&self, sizeof(self))== -1){
    perror("server:bind"); close(soc); exit(1);
}
listen(soc, 1);
... and now continuing ... */

/* accept connection request */
ns = accept(soc, (struct sockaddr *)&peer, &peer_len);
if (ns < 0) {
    perror("server:accept"); close(soc); exit(1);
}
/* data transfer on connected socket ns */
k = read(ns, buf, sizeof(buf));
printf("SERVER RECEIVED: %s\n", buf);
write(ns, buf, k);


close(ns);
close(soc);
```


Byte order

- Big-endian



- Little-endian



- Intel is little-endian, and Sparc is big-endian

data may need to order properly over different byte order machines in order to interpret them properly
Need to set up a network byte order to solve this problem

Network byte order

array of char is same on different byte order machine

- To communicate between machines with unknown or different “endian-ness” we convert numbers to network byte order (big-endian) before we send them.
- There are functions provided to do this:
 - unsigned long htonl(unsigned long)
 - unsigned short htons(unsigned short)
 - unsigned long ntohl(unsigned long)
 - unsigned short ntohs(unsigned short)

host to network long/short

network to host long / short

Sending and Receiving Data

- `read` and `write` calls work on sockets, but sometimes we want more control
- `ssize_t send(int fd, const void *buf, size_t len, int flags);`
 - works like `write` if `flags==0`
 - **flags:** `MSG_OOB`, `MSG_DONTROUTE`, `MSG_DONTWAIT`
- `ssize_t recv(int fd, void *buf, size_t len, int flags);`
 - **flags:** `MSG_OOB`, `MSG_WAITALL`, `MSG_PEEK`