

Concurrency

Kerrisk, Ch 29, 47, 48, 53, 54

Concurrency

- The two key concepts driving computer systems and applications are
 - **communication**: the conveying of information from one entity to another
 - **concurrency**: the sharing of resources in the same time frame
- Concurrency can exist in a single processor as well as in a multiprocessor system
- Managing concurrency is difficult, as execution behaviour is not always reproducible.

Concurrency Example

- Program a:

```
#!/usr/bin/sh
count=1
while [ $count -le 20 ]
do
    echo -n "a"
    count=`expr $count + 1`
done
```

- Program b

```
#!/usr/bin/sh
count=1
while [ $count -le 20 ]
do
    echo -n "b"
    count=`expr $count + 1`
done
```

- When run sequentially (a ; b) output is sequential.
- When run concurrently (a & ; b &) output is interspersed and different from run to run.

output different each time

Race conditions

- A **race condition** occurs when multiple processes are trying to do something with shared data and the final outcome depends on the order in which the processes run.
- E.g., If any code after a fork depends on whether the parent or child runs first.
- A parent process can call `wait()` to wait for termination (may block)
- A child process can wait for parent to terminate by polling (wasteful) (How would you do this?)
- One standard solution is to use signals.

Example 1

Process A

```
x = get(count)
write(x + 1)
```

x = 1

write(2)

Process B

```
y = get(count)
write(y + 1)
```

y = 2

write(3)

Count

1

2

3

The value of count is what we expect.

Example 2

Process A

```
x = get(count)
write(x + 1)
```

x = 1

write(2)

Process B

```
y = get(count)
write(y + 1)
```

y = 1

write(2)

y = 2

write(3)

Not what we
wanted!

Count

1

2

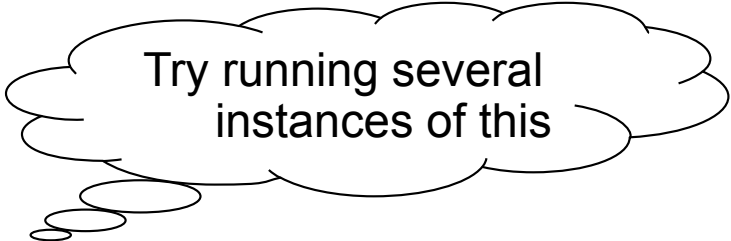
3

2

need atomic operation

Example: Race Conditions

```
#!/bin/sh
c=1
while [ $c -le 10 ]
do
    sd=`cat sharedData`
    sd=`expr $sd + 1`
    echo $sd > sharedData
    c=`expr $c + 1`
    echo d = $sd
done
#file sharedData must exist and hold
#one integer
```



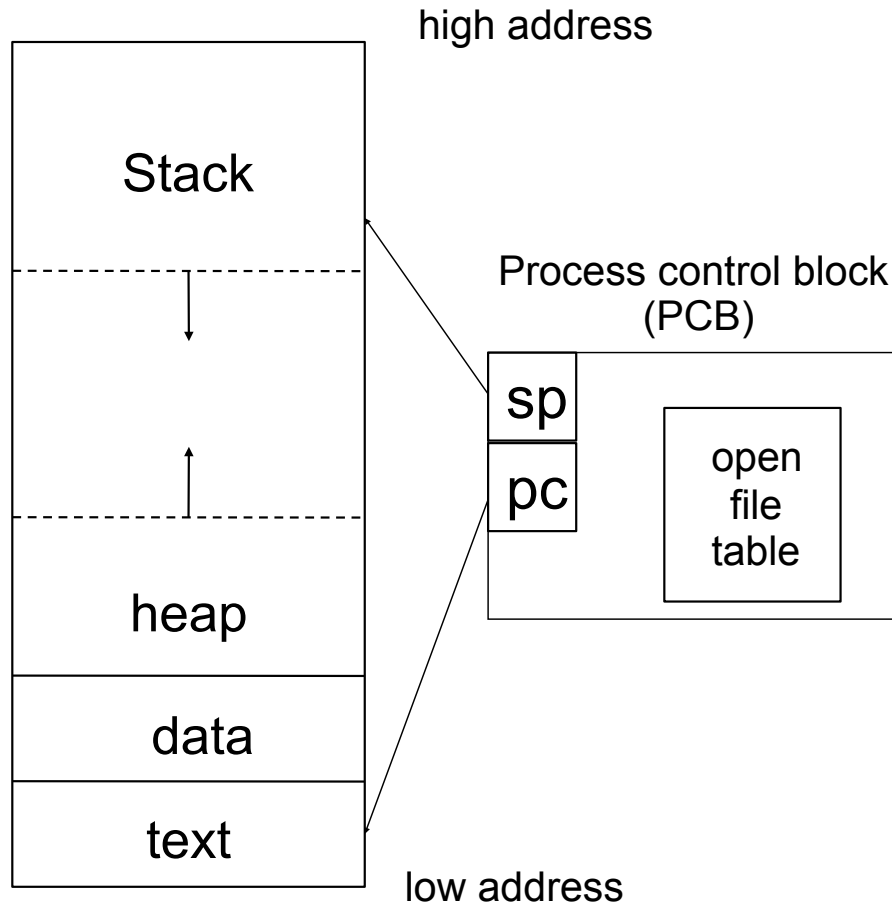
Try running several instances of this

Threads

Motivation

- Processes are expensive to create.
- It takes quite a bit of time to switch between processes
- Communication between processes must be done through an external structure
 - files, pipes, shared memory
- Synchronizing between processes is cumbersome.
- *Is there another model that will solve these problems?*

Processes



- Each process has its own
 - program counter
 - stack
 - stack pointer
 - address space
- Processes may share
 - open files
 - pipes


What is a process?

- OS abstraction for execution
- Running instance of a program
- Components of a process:
 - Address space
 - Code and data
 - Stack
 - Program Counter (PC)
 - Set of registers
 - Set of OS resources: open files, network connections...

Rethinking Processes

- What is similar in cooperating processes?
 - They all share the same code and data (address space)
 - They all share the same privileges
 - They all share the same resources (files, sockets, etc.)
- What don't they share?
 - Each has its own execution state: PC, SP, and registers
- **Key idea:** Why don't we separate the concept of a process from its execution state?
 - **Process:** address space, privileges, resources, etc.
 - **Execution state:** PC, SP, registers
- Exec state also called **thread of control**, or **thread**

Threads



- Each thread has its own
 - program counter
 - stack
 - stack pointer
- Threads share
 - address space
 - variables
 - code
 - open files

What is a Thread?

- A thread is a single control flow through a program
 - What is a “control flow”?
 - How is control flow represented?
- A program with multiple control flows is multithreaded

Control Flow



- **Control** includes all of the values that select which instructions in a program are executed.
- **Control flow**, then, is the sequence of instructions being executed.
- The hardware uses the program counter (PC) and stack to make control flow decisions.

Advantages

- Communication between threads is cheap
 - they can share variables!
- Threads are “lightweight”
 - faster to create
 - faster to switch between

Producer/Consumer Problem

- Simple example: `who | wc -l`
- Both the writing process (`who`) and the reading process (`wc`) of a pipeline execute concurrently.
- A pipe is usually implemented as an internal OS buffer.
- It is a resource that is concurrently accessed by the reader and the writer, so it must be managed carefully.

Producer/Consumer

- **consumer** should be blocked when buffer is empty
- **producer** should be blocked when buffer is full
- producer and consumer should run independently as far as buffer capacity and contents permit
- producer and consumer should never be updating the buffer at the same instant (otherwise **data integrity** cannot be guaranteed)
- producer/consumer is a harder problem if there are more than one consumer and/or more than one producer.

Pthreads

- POSIX threads (pthreads) is the most commonly used thread package on Unix/Linux

pthread_create

```
int pthread_create(pthread_t *tid,  
                  pthread_attr_t *attr,  
                  void *(*func)(void*), void *arg);
```

- `tid` uniquely identifies a thread within a process and is returned by the function
- `attr` sets attributes such as priority, initial stack size
 - can be specified as `NULL` to get defaults
- `func` - the function to call to start the thread
 - accepts one `void *` argument, returns `void *`
- `arg` is the argument to `func`
- returns 0 if successful, a positive error code if not
- does not set `errno` but returns compatible error codes
- can use `strerror()` to print error messages

pthread_join

```
int pthread_join(pthread_t tid,  
                 void **status)
```

- `tid` - the tid of the thread to wait for
 - cannot wait for any thread (as in `wait()`)
- `status`, if not NULL returns the `void *` returned by the thread when it terminates.
- a thread can terminate by
 - returning from `func`
 - the `main()` function exiting
 - `pthread_exit()`

More functions

- `void pthread_exit(void *status)`
 - a second way to exit, returns `status` explicitly
 - `status` must not point to an object local to the thread, as these disappear when the thread terminates.
- `int pthread_detach(pthread_t);`
 - if a thread is detached its termination cannot be tracked with `pthread_join()`
 - it becomes a daemon thread
- `pthread_t pthread_self(void)`
 - returns the thread ID of the thread which called it
 - often see `pthread_detach(pthread_self())`

Passing Arguments to Threads

```
pthread_t thread_ID;  int fd, result;  
fd = open("afile", "r");  
result = pthread_create(&thread_ID, NULL,  
                        myThreadFcn, (void *)&fd);  
if(result != 0)  
    printf("Error: %s\n", strerror(result));
```

- We can pass any variable (including a structure or array) to our thread function.
- It assumes the thread function knows what type it is.
- This example is **bad** if the main thread alters fd later.

Solution

- Use malloc() to create memory for the variable
 - initialize variable's value
 - pass pointer to new memory via pthread_create()
 - thread function releases memory when done.
- Example:

```
typedef struct myArg {  
    int fd;  
    char name[25];  
} MyArg;
```

```
int result;  
pthread_t thread_ID;
```


Example (cont'd)

```
MyArg *p = (MyArg *)malloc(sizeof(MyArg));
p->fd = fd; /* assumes fd is defined */
strncpy(p->name, "CSC209", 7);
result = pthread_create(&threadID, NULL,
                        myThreadFcn, (void *)p);
void *myThreadFcn(void *p) {
    MyArg *theArg = (MyArg *) p;
    write(theArg->fd, theArg->name, 7);
    close(theArg->fd);
    free(theArg);
    return NULL;
```

Thread-safe functions

- Not all functions can be called from threads
 - many use global/static variables
 - new versions of UNIX have thread-safe replacements like `strtok_r()`
- Safe:
 - `ctime_r()`, `gmtime_r()`, `localtime_r()`,
`rand_r()`, `strtok_r()`
- Not Safe:
 - `ctime()`, `gmtime()`, `localtime()`,
`rand()`, `strtok()`, `gethostxxx()`
- Could use semaphores to protect access but will generally result in poor performance.

Pthread Mutexes

```
int pthread_mutex_init(pthread_mutex_t *mp,  
    const pthread_mutexattr_t *attr);
```

```
int pthread_mutex_lock(pthread_mutex_t *mp);
```

```
int pthread_mutex_trylock(pthread_mutex_t *mp);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mp);
```

```
int pthread_mutex_destroy(pthread_mutex_t *mp);
```

- easier to use than `semget ()` and `semop ()`
- only the thread that locks a mutex can unlock it
- mutexes often declared as globals

Example

```
pthread_mutex_t myMutex;  
int status;  
  
status = pthread_mutex_init(&myMutex, NULL);  
if(status != 0)  
    printf("Error: %s \n", strerror(status));  
pthread_mutex_lock(&myMutex);  
/* critical section here */  
pthread_mutex_unlock(&myMutex);  
status = pthread_mutex_destroy(&myMutex);  
if(status != 0)  
    printf("Error: %s\n", strerror(status));
```