

Solutions for Homework Assignment #3

Answer to Question 1.

a. The scheduler \mathcal{S} can be implemented by an augmented AVL tree D . Each node u of D contains information about a thread t of \mathcal{S} . A node u contains the following fields:

- $id(u)$: the integer id of the thread.
- $status(u)$: the status of the thread, namely A , R or S .
- $ready(u)$: a boolean variable that is TRUE if and only if the subtree rooted at u (including u) has a node v such that $status(v) = R$. Note that the $ready$ field of each node u of D satisfies the identity:

$$ready(u) = ready(lchild(u)) \vee (status(u) = R) \vee ready(rchild(u)) \quad (*)$$

where $ready(NIL) = \text{FALSE}$.¹

- $lchild(u)$, $rchild(u)$ and $parent(u)$: pointers to the left child, right child, and the parent of u .
- $BF(u)$: the balance factor of u .

The id field is the *key* of the node in the AVL tree D — thus, an in-order traversal of D visits all the threads *in non-decreasing id order*.

b. The scheduler's operations are implemented as follows:

- **NEWTHREAD(t):** Given a thread $t = (i, status)$, where $status \in \{A, R, S\}$, inserts a node representing t into D ; assumes that i is different from the id of every other thread in D .

- (1) Prepare a new node u with $id(u) = i$ and $status(u) = A$ to represent thread t , and use the ordinary BST insertion algorithm to insert node u into D with i as the insertion key. After this insertion, the node u that represents t is a leaf of D .
- (2) If the **status of t is R** then traverse the path from the new leaf u to the root of the tree D to update the $ready$ fields: for each node v along this path set:

$$ready(v) := ready(lchild(v)) \vee (status(v) = R) \vee ready(rchild(v))$$

- (3) Traverse the path from the new leaf u to the root of the tree D again, performing rotations and updating the balance factors as required by the ordinary AVL insertion algorithm. In addition, update the **ready field of every node v involved in a rotation, according to the identity (*)**.

Stop the rotations and updates to the balance factors as in the ordinary AVL insertion algorithm.

Remark: The two traversals described above can be done in single pass where we first update the $ready$ field of the node under consideration and then we do the balancing/rotations part (which may also require a few additional $ready$ field updates).

Suppose that D contains n threads. Since D is an AVL tree, the height of any subtree of D is $O(\log n)$. The time required by **NEWTHREAD(t)** is that required by the ordinary AVL insertion algorithm, i.e., $O(\log n)$, plus the time required to update the $ready$ fields of the new node's ancestors (if the new node's $status$ is R), and the $ready$ fields of the nodes involved in a rotation (a few nodes for each rotation). Since there are $O(\log n)$ such nodes and each $ready$ update takes $O(1)$ time, the additional time required to update all the $ready$ fields is also $O(\log n)$. Thus, the worst-case time complexity of **NEWTHREAD(t)** is $O(\log n)$.

¹NIL denotes an empty pointer.

- **FIND(i, x):** Given a thread id i and a pointer to a node x of D , returns a pointer to the thread $t = (i, -)$ in the subtree of D rooted at x ; if no such thread exists, then it returns -1 .

This is a straightforward BST search algorithm:

```

FIND( $i, x$ )
  if  $x = \text{NIL}$  then
    return  $-1$ 
  else if  $\text{id}(x) < i$  then
    return FIND( $i, \text{lchild}(x)$ )
  else if  $\text{id}(x) = i$  then
    return  $x$ 
  else
    return FIND( $i, \text{rchild}(x)$ )
  end if

```

Note that **FIND(i)** is simply **FIND(i, r)** where r points to the root of the tree D .²

The worst-case time complexity of **FIND(i, x)** is $O(\log n)$. To see why, note that: (i) each call at a node u results in *at most one* recursive call, at the left or the right subtree of u ; and (ii) each call involves a constant amount of work. So the time to execute **FIND(i, x)** is proportional to the height of the subtree of D rooted at x , i.e., it is $O(\log n)$.

- **COMPLETED(i):** If \mathcal{S} has a thread $t = (i, -)$ then remove t from \mathcal{S} , else return -1 .

First call **Find(i)** to determine whether \mathcal{S} contains a thread $t = (i, -)$. If it does not, i.e., **Find(i)** returns -1 , then return -1 . If it does, i.e., **Find(i)** returns a pointer x to the node of D that contains $t = (i, -)$, proceed as follows:

- (1) Use the ordinary **BST deletion algorithm** to delete the node pointed to by x from D . Since D is an AVL tree, this deletion ultimately results in the removal of a leaf from the tree. Let z be the parent of that leaf. The rest of the algorithm is very similar to parts (2) and (3) of **NEWTHREAD(t)**, except that the traversals described below start from node z rather than the newly added leaf in **NEWTHREAD(t)**.
- (2) Traverse the path from z to the root of the tree D to update the *ready* fields: for each node v along this path, set $\text{ready}(v) := \text{ready}(\text{lchild}(v)) \vee (\text{status}(v) = R) \vee \text{ready}(\text{rchild}(v))$.
- (3) Traverse the path from z to the root of the tree D again, performing rotations and updating the balance factors as required by the ordinary AVL deletion algorithm. In addition, update the *ready* field of every node v involved in a rotation, according to the identity (*).

Stop the rotations and updates to the balance factors as in the ordinary AVL deletion algorithm.

Remark: The two traversals described above can be done in single pass where we first update the *ready* field of the node under consideration and then we do the balancing/rotations part (which may also require a few additional *ready* field updates).

The time required by **COMPLETED(i)** is that required by the ordinary AVL deletion algorithm, i.e., $O(\log n)$, plus the time required to update the *ready* fields of z 's ancestors, and the *ready* fields of the nodes involved in a rotation (a few nodes for each rotation). Since there are $O(\log n)$ such nodes and each *ready* update takes $O(1)$ time, the additional time required to update the *ready* fields is also $O(\log n)$. Thus, the worst-case time complexity of **COMPLETED(i)** is $O(\log n)$.

²More precisely, if the **FIND(i)** search is successful, it returns a *pointer* to a node u in D that represents some thread $t = (i, \text{status})$; given this pointer, it is trivial to return the pair $t = (i, \text{status})$ contained in u .

- **CHANGESTAT**($x, stat$): Given a pointer x to a node containing thread t in D , sets the status of t to $stat$.

To do so, we first set **status**(x) to $stat$. If $stat$ is not R , then we are done and just return. Otherwise, we update the **ready** field of every node in the path from x to the root of the tree D , according to the identity (*). Note that if the status of any node along this path does not change, this update process can stop (this obvious optimization is not shown in the basic pseudocode below).

```

CHANGESTAT( $x, stat$ )
  status( $x$ ) :=  $stat$ 
  if  $stat \neq R$  return
   $y := x$ 
  repeat
     $ready(y) = ready(lchild(y)) \vee (status(y) = R) \vee ready(rchild(y))$ 
     $y := parent(y)$ 
  until  $y = NIL$ 
  return

```

It is clear that the time to execute **CHANGESTAT**($x, stat$) is proportional to length of the path from x to the root of D . So the worst-case time complexity of **CHANGESTAT**($x, stat$) is $O(\log n)$.

- **CHANGESTATUS**($i, stat$): If \mathcal{S} has a thread $t = (i, -)$ then set the **status** of t to $stat$, else return -1 .

```

CHANGESTATUS( $i, stat$ )
   $x := \text{FIND}(i)$ 
  if  $x = -1$  then return  $-1$ 
  else return CHANGESTAT( $x, stat$ )

```

Since the worst-case time complexity of **FIND**(i) and **CHANGESTAT**($x, stat$) is $O(\log n)$, the worst-case time complexity of **CHANGESTATUS**($i, stat$) is also $O(\log n)$. pointer is fine

- **SCHEDULENEXT**(x): Given a pointer to a node x of D , returns a pointer to the thread t that has the smallest id among all the threads with **status** = R in the subtree of D rooted at x , and sets the status of t to A ; if the subtree rooted at x has no thread with **status** = R , then the procedure returns -1 .

```

SCHEDULENEXT( $x$ )
  if  $x = NIL \vee ready(x) = \text{FALSE}$  then
    return  $-1$ 
  else if  $ready(lchild(x)) = \text{TRUE}$  then
    return SCHEDULENEXT( $lchild(x)$ )
  else if  $status(x) = R$  then
    CHANGESTAT( $x, A$ )
    return  $x$ 
  else
    return SCHEDULENEXT( $rchild(x)$ )
  end if

```

Note that **SCHEDULENEXT** is simply **SCHEDULENEXT**(r) where r points to the root of the tree D .

The worst-case time complexity of **SCHEDULENEXT**(x) is $O(\log n)$. To see why, first note that **SCHEDULENEXT**(x) calls **CHANGESTAT**(x, A) at most once, and, as we have argued earlier, the time complexity of **CHANGESTAT**(x, A) is $O(\log n)$. Now consider the time required to execute **SCHEDULENEXT**(x) excluding the time to execute **CHANGESTAT**(x, A). Note that: (i) each call at a node u results in at most one recursive call, at the left or the right subtree of u ; and (ii) each call involves a constant amount of work other than the recursive call that it makes and the call to **CHANGESTAT**(x, A). So the time to execute **SCHEDULENEXT**(x) excluding the time taken by **CHANGESTAT**(x, A) is proportional to the height the subtree of D rooted at x , i.e., it is $O(\log n)$. Since the time complexity of **CHANGESTAT**(x, A) is $O(\log n)$, the total time to execute **SCHEDULENEXT**(x) is also $O(\log n)$ in the worst-case.

Answer to Question 2.

a. If the number of empty slots in T is k before inserting x , what is the probability that x is inserted into an empty slot?

Let $E = \{i_1, i_2, \dots, i_k\}$ be the set of *empty* slots of T before inserting x .

The probability that h_1 hashes x into an *empty* slot of T is:

$$\begin{aligned} \text{Prob}[h_1(x) \in E] &= \text{Prob}[h_1(x) = i_1 \vee h_1(x) = i_2 \vee \dots \vee h_1(x) = i_k] \\ &= \sum_{j=1}^{j=k} \text{Prob}[h_1(x) = i_j] \end{aligned}$$

Since T has m slots and we assume that the hash function h_1 satisfies the *Simple Uniform Hashing Assumption* (SUHA), we have $\text{Prob}[h_1(x) = i_j] = \frac{1}{m}$ for every j , $1 \leq j \leq k$. Therefore:

$$\text{Prob}[h_1(x) \in E] = \sum_{j=1}^{j=k} \frac{1}{m} = \frac{k}{m}$$

So the probability that h_1 hashes x into a *non-empty* slot of T is:

$$\text{Prob}[h_1(x) \notin E] = 1 - \frac{k}{m}$$

Since h_2 also satisfies SUHA, the probability that h_2 hashes x into a *non-empty* slot of T is also:

$$\text{Prob}[h_2(x) \notin E] = 1 - \frac{k}{m}$$

Note that x is inserted in a *non-empty* slot of T if and only if *both* h_1 and h_2 hash into a *non-empty* slot of T . Therefore:

$$\text{Prob}[x \text{ is inserted in a non-empty slot of } T] = \text{Prob}[h_1(x) \notin E \wedge h_2(x) \notin E]$$

Thus, since the hash functions h_1 and h_2 are independent:

$$\begin{aligned} \text{Prob}[x \text{ is inserted in a non-empty slot of } T] &= \text{Prob}[h_1(x) \notin E] \cdot \text{Prob}[h_2(x) \notin E] \\ &= \left(1 - \frac{k}{m}\right)^2 \end{aligned}$$

Therefore:

$$\begin{aligned} \text{Prob}[x \text{ is inserted in an empty slot of } T] &= 1 - \text{Prob}[x \text{ is inserted in a non-empty slot of } T] \\ &= 1 - \left(1 - \frac{k}{m}\right)^2 \\ &= \frac{k(2m - k)}{m^2} \end{aligned}$$

b. Suppose that $m = 4$, and $T[0]$ contains 6 elements, $T[1]$ contains 3 elements, and $T[2]$, $T[3]$ contain 9 elements each. What is the *expected* length of the chain x is inserted into, *not counting* x itself?

Since the table has 4 slots, when we insert x using the two independent hash functions h_1 and h_2 there are 16 possible hashing outcomes:

$$S = \{(i, j) \mid \text{where } 0 \leq i = h_1(x) \leq 3 \text{ and } 0 \leq j = h_2(x) \leq 3\} \quad (S \text{ is our sample space})$$

Note that:

$$\begin{aligned}
\text{Prob}[(i, j)] &= \text{Prob}[h_1(x) = i \wedge h_2(x) = j] && \text{(by definition)} \\
&= \text{Prob}[h_1(x) = i] \cdot \text{Prob}[h_2(x) = j] && \text{(because } h_1 \text{ and } h_2 \text{ are independent)} \\
&= \frac{1}{4} \cdot \frac{1}{4} = \frac{1}{16} && \text{(because } h_1 \text{ and } h_2 \text{ satisfy SUHA)}
\end{aligned}$$

Let X be the random variable denoting the length of the chain where x is inserted. That is, for each possible outcome $o = (i, j) \in S$, $X(o)$ is the length of the chain where x is inserted. We are seeking $E[X]$.

Note that:

$$\begin{aligned}
E[X] &= \sum_{o \in S} X(o) \cdot \text{Prob}[o] \\
&= \sum_{(i, j) \in S} X[(i, j)] \cdot \text{Prob}[(i, j)] \\
&= \frac{1}{16} \sum_{(i, j) \in S} X[(i, j)] \\
&= \frac{1}{16} (X[(0, 0)] + X[(0, 1)] + X[(0, 2)] + X[(0, 3)] \\
&\quad + X[(1, 0)] + X[(1, 1)] + X[(1, 2)] + X[(1, 3)] \\
&\quad + X[(2, 0)] + X[(2, 1)] + X[(2, 2)] + X[(2, 3)] \\
&\quad + X[(3, 0)] + X[(3, 1)] + X[(3, 2)] + X[(3, 3)])
\end{aligned}$$

another way:
straight calculate expected value of length.
by summing up length of slot of all 16 outcomes
then divide by 16

Recall that $T[0]$ contains 6 elements, $T[1]$ contains 3 elements, and $T[2]$, $T[3]$ contain 9 elements each. From our hashing scheme, when the hashing outcome is (i, j) , the length $X[(i, j)]$ of the chain that x enters is as follows: if $i = 1$ or $j = 1$ then $X[(i, j)] = 3$, else if $i = 0$ or $j = 0$ then $X[(i, j)] = 6$ else $X[(i, j)] = 9$. Therefore:

$$\begin{aligned}
E[X] &= \frac{1}{16} (6 + 3 + 6 + 6 \\
&\quad + 3 + 3 + 3 + 3 \\
&\quad + 6 + 3 + 9 + 9 \\
&\quad + 6 + 3 + 9 + 9) \\
&= \frac{87}{16} \\
&= 5.4375
\end{aligned}$$

Another way to compute $E[X]$ is as follows. Note that the chain where x is inserted has length 3, 6 or 9. So the only possible value v of the random variable X is 3, 6 and 9. An alternative formula for $E[X]$ is:

$$E[X] = \sum_{v \in \{3, 6, 9\}} v \cdot \text{Prob}[X = v]$$

It now suffices to compute $\text{Prob}[X = 3]$, $\text{Prob}[X = 6]$, and $\text{Prob}[X = 9]$.

Note that $X = 3$ iff the outcome of the hashing is an (i, j) with $i = 1$ or $j = 1$. Therefore:

$$\begin{aligned}
\text{Prob}[X = 3] &= \text{Prob}[\{(0, 1), (1, 0), (1, 1), (1, 2), (1, 3), (2, 1), (3, 1)\}] \\
&= \text{Prob}[(0, 1)] + \text{Prob}[(1, 0)] + \text{Prob}[(1, 1)] + \text{Prob}[(1, 2)] + \text{Prob}[(1, 3)] + \text{Prob}[(2, 1)] + \text{Prob}[(3, 1)] \\
&= 7 \cdot \frac{1}{16} = \frac{7}{16}
\end{aligned}$$

Similarly:

$$\begin{aligned}
\text{Prob}[X = 6] &= \text{Prob}[\{(0, 0), (0, 2), (0, 3), (2, 0), (3, 0)\}] \\
&= \text{Prob}[(0, 0)] + \text{Prob}[(0, 2)] + \text{Prob}[(0, 3)] + \text{Prob}[(2, 0)] + \text{Prob}[(3, 0)] \\
&= 5 \cdot \frac{1}{16} = \frac{5}{16} \\
\text{Prob}[X = 9] &= \text{Prob}[\{(2, 2), (2, 3), (3, 2), (3, 3)\}] \\
&= \text{Prob}[(2, 2)] + \text{Prob}[(2, 3)] + \text{Prob}[(3, 2)] + \text{Prob}[(3, 3)] \\
&= 4 \cdot \frac{1}{16} = \frac{4}{16}
\end{aligned}$$

Therefore:

$$\begin{aligned}
E[X] &= \sum_{v \in \{3, 6, 9\}} v \cdot \text{Prob}[X = v] \\
&= 3 \cdot \frac{7}{16} + 6 \cdot \frac{5}{16} + 9 \cdot \frac{4}{16} \\
&= \frac{87}{16}
\end{aligned}$$

Answer to Question 3.

direct addressing table is better than array can just search with $A[\text{index}]$

- (1) We will use two data structures for our algorithm: a **direct access table A with 26 indices** (corresponding to each letter in the English alphabet) and a hash table T .

We assume that the **size t of the hash table T is within a constant factor of l** ; more precisely, we assume that $t = \kappa \cdot l$, for some constant $\kappa > 0$. We will also assume our hashing function satisfies the Simple Uniform Hashing Assumption (SUHA).

- (2) The idea for our algorithm is as follows. Initially, we set every value in our direct access table A to **null**; A will store the current most frequent word starting with each letter. When we receive an input *word*, we insert $(\text{word}, 1)$ into our hash table T using *word* as our key if it does not already exist in H ; otherwise (if it does), we increment its frequency:

$$(\text{word}, \text{freq}) \rightarrow (\text{word}, \text{freq} + 1)$$

Then, let the first letter of *word* be some letter corresponding to value x (e.g. “b” corresponds to 2); if *word*’s frequency is greater than the frequency of the word at $A[x]$ (or if $A[x] == \text{null}$), then set $A[x]$ to *word*.

To perform a *query* operation, we simply read our direct access table from indices 1 to 26.

- (3) CHAINED-HASH-SEARCH(T, k):
return pointer to node with key k in list $T[h(k)]$, null if doesn’t exist

CHAINED-HASH-INSERT(T, x):
insert x at the head of list $T[h(x.\text{key})]$

INSERT(T, A, word):
// update hash table storing word frequencies
 $u := \text{CHAINED-HASH-SEARCH}(T, \text{word})$
if $u = \text{null}$:
create a new node u s.t. $u.\text{key} = \text{word}$ and $u.\text{value} = 1$
CHAINED-HASH-INSERT(T, u)
else:

```

    increment u.value by 1
    // check if there is a new most frequent word
    l = word[0] // first letter of word
    x = l - 'a' // convert to number
    current_max = CHAINED-HASH-SEARCH(T, A[x]).value
    if u.value > current_max or (u.value = current_max and word < A[x]):
        A[x] := word

```

```

QUERY(A):
    for i in 1 ... 26:
        print T[i]

```

Note: it is also possible to store $(word, freq)$ in the direct access table so that you don't have to look up the frequency of the current max. This slightly increases efficiency in exchange for a slight increase in storage.

- (4) Under the simple uniform hashing assumption (SUHA), i.e., that each distinct value is equally likely to be hashed into any one of T 's slots, the expected length of a linked list in T is α , where $\alpha = l/t$. Since t is $t = \kappa \cdot l$, α is $\Theta(1)$, and so the expected length of a linked list in T is $\Theta(1)$. Therefore the expected time for an INSERT operation is $\Theta(1)$ — a linked list in T is accessed twice; all other operations are clearly constant (in terms of the current length of the sequence). The QUERY operation runs in constant time since accessing every element in A is a constant time operation and $|A| = 26$; that is, $|A|$ is not dependant on the length of the sequence.