

# CSC321 Lecture 22: Q-Learning

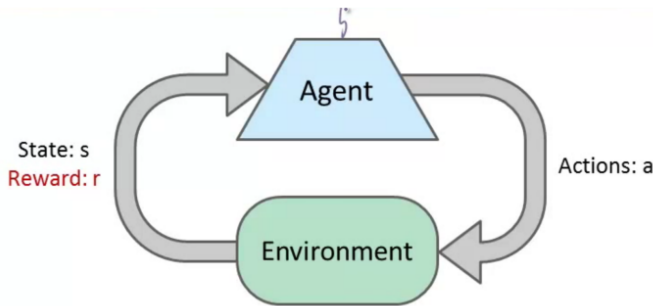
Roger Grosse

# Overview

- Second of 3 lectures on reinforcement learning
- Last time: policy gradient (e.g. REINFORCE)
  - Optimize a policy directly, don't represent anything about the environment
- Today: Q-learning
  - Learn an action-value function that predicts future returns
- Next time: AlphaGo uses both a policy network and a value network
- This lecture is review if you've taken 411
- This lecture has more new content than I'd intended. If there is an exam question about this lecture or next one, it won't be a hard question.

# Overview

- Agent interacts with an environment, which we treat as a black box
- Your RL code accesses it only through an API since it's external to the agent
  - I.e., you're not “allowed” to inspect the transition probabilities, reward distributions, etc.



# Recap: Markov Decision Processes

- The environment is represented as a **Markov decision process (MDP)**  $\mathcal{M}$ .
- Markov assumption: all relevant information is encapsulated in the current state
- Components of an MDP:
  - initial state distribution  $p(\mathbf{s}_0)$
  - transition distribution  $p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$
  - reward function  $r(\mathbf{s}_t, \mathbf{a}_t)$
- **policy**  $\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$  parameterized by  $\theta$
- Assume a **fully observable** environment, i.e.  $\mathbf{s}_t$  can be observed directly  
**real world: cant see everything, incomplete information**

# Finite and Infinite Horizon

- Last time: finite horizon MDPs
  - Fixed number of steps  $T$  per episode
  - Maximize expected return  $R = \mathbb{E}_{p(\tau)}[r(\tau)]$
- Now: more convenient to assume **infinite horizon**
  - We can't sum infinitely many rewards, so we need to discount them:  
\$100 a year from now is worth less than \$100 today
  - **Discounted return**

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \text{finite}$$

- Want to choose an action to maximize expected discounted return
- The parameter  $\gamma < 1$  is called the **discount factor**
  - small  $\gamma$  = myopic
  - large  $\gamma$  = farsighted **more weight to future rewards**

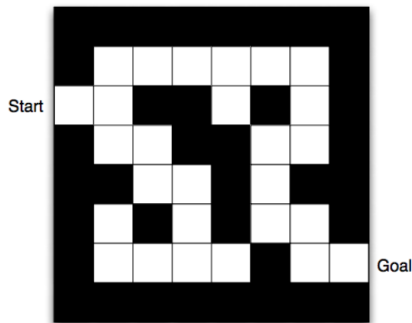
# Value Function

- **Value function**  $V^\pi(\mathbf{s})$  of a state  $\mathbf{s}$  under policy  $\pi$ : the **expected discounted return** if we start in  $\mathbf{s}$  and follow  $\pi$

$$\begin{aligned} V^\pi(\mathbf{s}) &= \mathbb{E}[G_t \mid \mathbf{s}_t = \mathbf{s}] \\ &= \mathbb{E} \left[ \sum_{i=0}^{\infty} \gamma^i r_{t+i} \mid \mathbf{s}_t = \mathbf{s} \right] \end{aligned}$$

- Computing the value function is generally impractical, but we can try to approximate (learn) it
- The benefit is credit assignment: see directly how an action affects future returns rather than wait for rollouts



# Value Function



- Rewards: -1 per time step
- Undiscounted ( $\gamma = 1$ )
- Actions: N, E, S, W 4 directions
- State: current location

# Value Function

value function: state  $s \rightarrow R$





# Action-Value Function

- Can we use a value function to choose actions?

$$\arg \max_{\mathbf{a}} r(\mathbf{s}_t, \mathbf{a}) + \gamma \mathbb{E}_{p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)} [V^{\pi}(\mathbf{s}_{t+1})]$$

can only decide which action to take,  
not what next state is in

# Action-Value Function

- Can we use a value function to choose actions?

$$\arg \max_{\mathbf{a}} r(\mathbf{s}_t, \mathbf{a}) + \gamma \mathbb{E}_{p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)} [V^{\pi}(\mathbf{s}_{t+1})]$$

- Problem: this requires taking the expectation with respect to the environment's dynamics, which we don't have direct access to!
- Instead learn an **action-value function**, or **Q-function**: expected returns if you take action  $\mathbf{a}$  and then follow your policy

$$Q^{\pi}(\mathbf{s}, \mathbf{a}) = \mathbb{E}[G_t | \mathbf{s}_t = \mathbf{s}, \mathbf{a}_t = \mathbf{a}]$$

- Relationship:

$$V^{\pi}(\mathbf{s}) = \sum_{\mathbf{a}} \pi(\mathbf{a} | \mathbf{s}) Q^{\pi}(\mathbf{s}, \mathbf{a})$$

- Optimal action:

$$\arg \max_{\mathbf{a}} Q^{\pi}(\mathbf{s}, \mathbf{a})$$

# Bellman Equation

- The **Bellman Equation** is a recursive formula for the action-value function:

$$Q^{\pi}(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{p(\mathbf{s}' | \mathbf{s}, \mathbf{a}) \pi(\mathbf{a}' | \mathbf{s}')} [Q^{\pi}(\mathbf{s}', \mathbf{a}')] ]$$

- There are various Bellman equations, and most RL algorithms are based on repeatedly applying one of them.

# Optimal Bellman Equation

- The **optimal policy**  $\pi^*$  is the one that maximizes the expected discounted return, and the **optimal action-value function**  $Q^*$  is the action-value function for  $\pi^*$ .
- The **Optimal Bellman Equation** gives a recursive formula for  $Q^*$ :

$$Q^*(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{p(\mathbf{s}' | \mathbf{s}, \mathbf{a})} \left[ \max_{\mathbf{a}'} Q^*(\mathbf{s}_{t+1}, \mathbf{a}') \mid \mathbf{s}_t = \mathbf{s}, \mathbf{a}_t = \mathbf{a} \right]$$

- This system of equations characterizes the optimal action-value function. So maybe we can **approximate  $Q^*$**  by trying to solve the optimal Bellman equation!

# Q-Learning

algorithm for learning (approx) Q

- Let  $Q$  be an action-value function which hopefully approximates  $Q^*$ .
- The **Bellman error** is the update to our expected return when we observe the next state  $\mathbf{s}'$ .

$$\underbrace{r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \max_{\mathbf{a}} Q(\mathbf{s}_{t+1}, \mathbf{a}) - Q(\mathbf{s}_t, \mathbf{a}_t)}_{\text{inside } \mathbb{E} \text{ in RHS of Bellman eqn}}$$

- The Bellman equation says the Bellman error is 0 in expectation
- **Q-learning** is an algorithm that repeatedly adjusts  $Q$  to minimize the Bellman error
- Each time we **sample** consecutive states and actions  $(\mathbf{s}_t, \mathbf{a}_t, \mathbf{s}_{t+1})$ :

$$Q(\mathbf{s}_t, \mathbf{a}_t) \leftarrow Q(\mathbf{s}_t, \mathbf{a}_t) - \alpha \underbrace{\left[ r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \max_{\mathbf{a}} Q(\mathbf{s}_{t+1}, \mathbf{a}) - Q(\mathbf{s}_t, \mathbf{a}_t) \right]}_{\text{Bellman error}}$$

at each step

# Exploration-Exploitation Tradeoff

- Notice: Q-learning only learns about the states and actions it visits.
- **Exploration-exploitation tradeoff**: the agent should sometimes **pick suboptimal actions** in order to visit new states and actions.
- Simple solution:  **$\epsilon$ -greedy policy**
  - With probability  $1 - \epsilon$ , choose the optimal action according to  $Q$
  - With probability  $\epsilon$ , choose a random action
- Believe it or not,  $\epsilon$ -greedy is still used today! doctor do random thing 1% of time...

# Exploration-Exploitation Tradeoff

- You can't use an epsilon-greedy strategy with policy gradient because it's an **on-policy algorithm**: the agent can only learn about the policy it's actually following.
- Q-learning is an **off-policy** algorithm: the agent can learn  $Q$  regardless of whether it's actually following the optimal policy
- Hence, Q-learning is typically done with an  $\epsilon$ -greedy policy, or some other policy that encourages exploration.

# Q-Learning

```
Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
  Initialize  $S$ 
  Repeat (for each step of episode):
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
    Take action  $A$ , observe  $R, S'$ 
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$ 
     $S \leftarrow S'$ ;
  until  $S$  is terminal
```



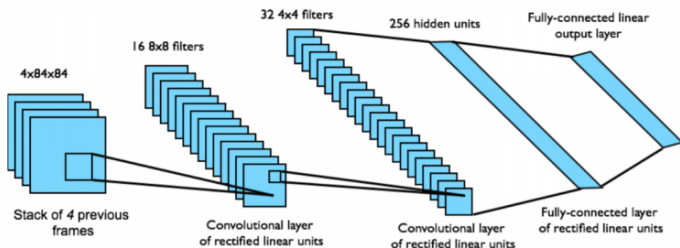
# Function Approximation

- So far, we've been assuming a **tabular representation** of  $Q$ : one entry for every state/action pair.
- This is impractical to store for all but the simplest problems, and doesn't share structure between related states.
- Solution: approximate  $Q$  using a parameterized function, e.g.
  - linear function approximation:  $Q(\mathbf{s}, \mathbf{a}) = \mathbf{w}^\top \psi(\mathbf{s}, \mathbf{a})$
  - compute  $Q$  with a neural net
- Update  $Q$  using backprop:

$$t \leftarrow r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \max_{\mathbf{a}} Q(\mathbf{s}_{t+1}, \mathbf{a})$$
$$\theta \leftarrow \theta + \alpha (t - Q(\mathbf{s}, \mathbf{a})) \frac{\partial Q}{\partial \theta}$$

# Function Approximation

- Approximating  $Q$  with a neural net is a decades-old idea, but DeepMind got it to work really well on Atari games in 2013 (“deep Q-learning”)
- They used a very small network by today’s standards



- Main technical innovation: store experience into a **replay buffer**, and perform Q-learning using stored experience
  - Gains sample efficiency by separating environment interaction from optimization — don’t need new experience for every SGD update!

- Mnih et al., *Nature* 2015. Human-level control through deep reinforcement learning
- Network was given raw pixels as observations
- Same architecture shared between all games
- Assume fully observable environment, even though that's not the case
- After about a day of training on a particular game, often beat “human-level” performance (number of points within 5 minutes of play)
  - Did very well on reactive games, poorly on ones that require planning (e.g. Montezuma's Revenge)
- <https://www.youtube.com/watch?v=V1eYniJ0Rnk>
- <https://www.youtube.com/watch?v=4MlZncshy1Q>

# Wireheading

- If rats have a lever that causes an electrode to stimulate certain “reward centers” in their brain, they’ll keep pressing the lever at the expense of sleep, food, etc.
- RL algorithms show this “wireheading” behavior if the reward function isn’t designed carefully
- <https://blog.openai.com/faulty-reward-functions/>

# Policy Gradient vs. Q-Learning

- Policy gradient and Q-learning use two very different choices of representation: **policies and value functions**
- Advantage of both methods: don't need to model the environment
- Pros/cons of policy gradient
  - Pro: unbiased estimate of gradient of expected return
  - Pro: can handle a large space of actions (since you only need to sample one)
  - Con: high variance updates (implies poor sample efficiency)
  - Con: doesn't do credit assignment  
doesnt tell which action in a sequence is more important
- Pros/cons of Q-learning
  - Pro: lower variance updates, more sample efficient
  - Pro: does credit assignment
  - Con: biased updates since Q function is approximate (drinks its own Kool-Aid)
  - Con: hard to handle many actions (since you need to take the max)

# Actor-Critic (optional)

Actor-critic methods combine the best of both worlds

- Fit both a policy network (the “actor”) and a value network (the “critic”)
- Repeatedly update the value network to estimate  $V^\pi$
- Unroll for only a few steps, then compute the REINFORCE policy update using the expected returns estimated by the value network
- The two networks adapt to each other, much like GAN training
- Modern version: Asynchronous Advantage Actor-Critic (A3C)