

CSC367 Parallel computing

Lecture 9: Parallel Architectures and Parallel Algorithm Design Continued!


Next up ...

- Shared memory architecture
- Parallel programming models: shared memory
- Pthreads: Synchronization, Races, Locks
- OpenMP
- Cache coherency

Next up ...

- Shared memory architecture
- Parallel programming models: shared memory
- Pthreads: Synchronization, Races, Locks
- OpenMP
- Cache coherency

Shared Memory Architecture



Chip Multiprocessor (CMP)

Next up ...

- Shared memory architecture
- Parallel programming models: shared memory
- Pthreads: Synchronization, Races, Locks
- OpenMP
- Cache coherency

Parallel Programming Models

- **Programming model** is made up of the languages and libraries that create an abstract view of the machine: Pthreads!


The programming model enables us to identify

- Control
 - How is parallelism **created**?
 - What **orderings** exist between operations?
- Data:
 - What data is **private** vs. **shared**?
 - How is logically shared data accessed or **communicated**?
- Synchronization
 - What operations can be used to coordinate parallelism?
 - What are the **atomic** (indivisible) operations?

Programming Model: Shared Memory

Program is a collection of threads of control, can be created mid-execution.


Thread



Programming Model: Shared Memory

Program is a collection of threads of control, can be created mid-execution.

Thread




Programming Model: Shared Memory

Program is a collection of threads of control, can be created mid-execution.

Each thread has a set of **private variables**, e.g., local stack variables.

Thread



Programming Model: Shared Memory


Program is a collection of threads of control, can be created mid-execution.

Each thread has a set of **private variables**, e.g., local stack variables.

Also a set of **shared variables**, e.g., static variables.

- Threads communicate **implicitly** by writing and reading shared variables.

Thread



Slide Source: Demmel

Next up ...

- Shared memory architecture
- Parallel programming models: shared memory
- Pthreads: Synchronization, Races, Locks
- OpenMP
- Cache coherency

Overview of POSIX Threads

- POSIX: *Portable Operating System Interface*
 - Interface to Operating System utilities
- PThreads: The POSIX threading interface
 - System calls to create and synchronize threads
 - Should be relatively uniform across UNIX-like OS platforms
- PThreads contain support for
 - Creating parallelism
 - Synchronizing
 - No explicit support for communication, because shared memory is implicit; a pointer to shared data is passed to a thread

Forking Posix Threads

Signature:

```
int pthread_create(pthread_t *, const pthread_attr_t *, void * (*)(void *), void *);
```

Example call:

```
errcode = pthread_create(&thread_id; &thread_attribute; &thread_fun; &fun_arg);
```

- `thread_id` is the thread id or handle (used to halt, etc.)
- `thread_attribute` various attributes
 - Standard default values obtained by passing a NULL pointer
 - Sample attributes: minimum stack size, priority
- `thread_fun` the function to be run (takes and returns void*)
- `fun_arg` an argument can be passed to `thread_fun` when it starts
- `errorcode` will be set nonzero if the create operation fails

“Simple” Threading Example

```
void* SayHello(void *foo) {  
    printf( "Hello, world!\n" );  
    return NULL;  
}
```

Compile using gcc -lpthread

```
int main() {  
    pthread_t threads[16];  
    int tn;  
    for(tn=0; tn<16; tn++) {  
        pthread_create(&threads[tn], NULL, SayHello, NULL);  
    }  
    for(tn=0; tn<16 ; tn++) {  
        pthread_join(threads[tn], NULL);  
    }  
    return 0;  
}
```

Synchronization

- Threads interact in a multiprogrammed system
 - To share resources (such as shared data)
 - To coordinate their execution
- Arbitrary interleaving of thread executions can have unexpected consequences
 - We need a way to restrict the possible interleavings of executions
 - Scheduling is invisible to the application => cannot know when we lose control of the CPU and another thread/process runs
- **Synchronization** is the mechanism that gives us this control

Motivating Example

```
EggRun(fridge *f) {  
    int eggs_left = f->egg_count;  
    if(eggs_left == 0) {  
        eggs_left = buy_carton();  
        f->egg_count += eggs_left;  
    }  
}
```

```
EggRun(fridge *f) {  
    int eggs_left = f->egg_count;  
    if(eggs_left == 0) {  
        eggs_left = buy_carton();  
        f->egg_count += eggs_left;  
    }  
}
```

- Separate threads, which may run concurrently; `eggs_left` is local to each thread while the `f->egg_count` is shared
- Assume fridge has no eggs initially
- Think about potential schedules for these two threads

Interleaved Schedules

- The execution of the two threads can be interleaved:

T1:

```
int eggs_left = f->egg_count;  
if(eggs_left == 0) {
```

```
    eggs_left = buy_carton();  
    f->egg_count += eggs_left;  
}
```

T2:

```
int eggs_left = f->egg_count;  
if(eggs_left == 0) {
```

```
    eggs_left = buy_carton();  
    f->egg_count += eggs_left;  
}
```

time



Interleaved Schedules

- The execution of the two threads can be interleaved:

T1:

```
int eggs_left = f->egg_count;  
if(eggs_left == 0) {
```

```
    eggs_left = buy_carton();  
    f->egg_count += eggs_left;  
}
```

T2:

```
int eggs_left = f->egg_count;  
if(eggs_left == 0) {
```

```
    eggs_left = buy_carton();  
    f->egg_count += eggs_left;  
}
```

time



We end up buying **two** cartons of eggs

Interleaved Schedules

- The execution of the two threads can be interleaved:

T1:

```
int eggs_left = f->egg_count;
if(eggs_left == 0) {
    eggs_left = buy_carton();
    f->egg_count += eggs_left;
}
```

T2:

```
int eggs_left = f->egg_count;
if(eggs_left == 0) {
    eggs_left = buy_carton();
    f->egg_count += eggs_left;
}
```

time



Interleaved Schedules

- The execution of the two threads can be interleaved:

T1:

```
int eggs_left = f->egg_count;  
if(eggs_left == 0) {  
    eggs_left = buy_carton();  
    f->egg_count += eggs_left;  
}
```

T2:

```
int eggs_left = f->egg_count;  
if(eggs_left == 0) {  
    eggs_left = buy_carton();  
    f->egg_count += eggs_left;  
}
```

time



We end up buying **one** carton of eggs

Race conditions and synchronization

- What happens when 2 or more concurrent threads manipulate a *shared resource* (e.g., a piece of data) without any synchronization?
 - The outcome depends on the order in which accesses take place!
 - This is called a *race condition* `f->egg_count += (unsynchronized write to shared memory)`
 - We need to ensure that only one thread at a time can manipulate the shared resource
 - So that we can reason about correct program behavior
- => We need *synchronization*

How do we handle this?

- How about whoever gets to check first, locks the fridge and takes the sole key, for the duration of the entire grocery run?
 - Nobody else can unlock the shared resource until the key owner unlocks it

Mutual Exclusion

- Given:
 - A set of n threads, T_0, T_1, \dots, T_{n-1}
 - A set of resources shared between threads
 - A segment of code which accesses the shared resources, called the *critical section, CS*
- We want to ensure that:
 - Only one thread at a time can execute in the critical section
 - All other threads are forced to wait on entry
 - When a thread leaves the CS, another can enter

Mutex locks

- Typically associated to a resource, to ensure one access at a time, to that resource
- Ensure mutual exclusion to a critical section
- For Mutexes, a thread go to sleep when they see the lock is busy.

Spinlock Implementation

- There are two operations on locks: *acquire()* and *release()*

```
boolean lock;  
  
void acquire(boolean *lock) {  
    while(test_and_set(lock)) ;  
}  
  
void release(boolean *lock) {  
    *lock = false;  
}
```

When false, we know
that we've acquired it

To release, simply
turn it to false.

- This is a *spinlock*
 - Uses *busy waiting* - thread continually executes *while* loop in *acquire()* , consumes CPU cycles

Spinlocks vs Mutex

- Spinlocks are built on machine instructions and because *busy waiting* they will waste CPU cycles.
- Mutexes are usually the better choice because a thread goes to sleep if the lock is busy allowing another thread to execute on that core/CPU.
- Putting threads to sleep is expensive so for very short tasks it might not be beneficial to use a mutex.
- However, most modern systems will allow a mutex to spinlock for a very short amount of time, if this seems beneficial.

Next up ...

- Using locks for synchronization
- Common mistakes, potential correctness problems
- Coarse-grained vs. fine-grained locking
- Deadlocks

POSIX mutex API

- Pthreads library has builtin mutexes
 - You've seen these in the labs already
- Basic API:
 - `pthread_mutex_t mutex;`
 - `pthread_mutex_init(pthread_mutex_t *mutex);`
 - `pthread_mutex_lock(pthread_mutex_t *mutex);`
 - `pthread_mutex_unlock(pthread_mutex_t *mutex);`

Potential correctness problems

- Both reads and writes to shared data must be locked, if a concurrent write is possible

```
typedef struct {  
    int egg_count;  
    double milk_qty;  
    pthread_mutex_t lock;  
} fridge;
```

```
EatEggOrDieTrying(fridge *f) {  
    pthread_mutex_lock(f->lock);  
    if(f->egg_count > 0) {  
        f->egg_count --;  
    }  
    else {  
        printf("Plan B: cereal\n");  
    }  
    pthread_mutex_unlock(f->lock);  
}
```

```
EggRun(fridge *f) {  
    pthread_mutex_lock(f->lock);  
    int eggs_left = f->egg_count;  
    if(eggs_left == 0) {  
        eggs_left = buy_carton();  
        f->egg_count += eggs_left;  
    }  
    pthread_mutex_unlock(f->lock);  
  
    printf("Eggs refilled: %d remaining!", f->egg_count);  
}
```

Problem!

No lock around printf in the
yellow box so possible bogus
output of
Eggs refilled:
0 remaining!

not as expected i guess

Potential correctness problems

- Careful about losing track of a lock without unlocking
 - e.g., what happens here:

```
bool CanEatEggs(fridge *f) {  
    pthread_mutex_lock(f->lock);  
    int eggs_left = f->egg_count;  
    if(eggs_left == 0){  
        printf("Oh no!\n");  
        return false;    lock not released !  
    }  
    printf("Yummy, eggs!\n");  
    pthread_mutex_unlock(f->lock);  
    return true;  
}
```

- If a thread never releases a lock, all other waiting threads are stuck
 - Such concurrency bugs are called **deadlocks**! (more on this later...)

Locking – coarse vs fine-grained

- Locking large sections of code may not be efficient => limits concurrency
- What if T1 wants to do a MilkRun, while T2 does an EggRun?
 - Locking the fridge for the EggRun won't allow a MilkRun to happen
- Solution: fine-grained locking
 - Use smaller locks, lock **only what is needed**...
- Advantage: reduces unnecessary waiting/blocking, more parallelism

Example

- Separate locks => can run in parallel, higher degree of concurrency

```
typedef struct {  
    int egg_count;  
    pthread_mutex_t egg_lock;  
  
    double milk_qty;  
    pthread_mutex_t milk_lock;  
} fridge;
```

```
MilkRun(fridge *f) {  
    pthread_mutex_lock(f->milk_lock);  
    double milk_left = f->milk_qty;  
    if(milk_left == 0) {  
        milk_left = buy_carton();  
        f->milk_qty += milk_left;  
    }  
    pthread_mutex_unlock(f->milk_lock);  
}
```

```
EggRun(fridge *f) {  
    pthread_mutex_lock(f->egg_lock);  
    int eggs_left = f->egg_count;  
    if(eggs_left == 0) {  
        eggs_left = buy_carton();  
        f->egg_count += eggs_left;  
    }  
    pthread_mutex_unlock(f->egg_lock);  
}
```


Careful with fine-grained locking

- Morning routine includes eating eggs and drinking milk
- Must have both eggs and milk to eat breakfast, otherwise breakfast is ruined
- MorningRoutine thread executes concurrently with other threads that perform other breakfast routines => Locks are used (see code)
- Problem? Two threads might try to eat the same egg!

```
MorningRoutine(fridge *f, int e, double m) {
    pthread_mutex_lock(f->egg_lock);
    int eggs_left = f->egg_count;
    if(eggs_left > 0) {
        pthread_mutex_unlock(f->egg_lock);

        pthread_mutex_lock(f->milk_lock);
        double milk_left = f->milk_qty;
        if(milk_left > 0) {
            pthread_mutex_unlock(f->milk_lock);

            pthread_mutex_lock(f->egg_lock);
            f->egg_count -= e;
            pthread_mutex_unlock(f->egg_lock);

            pthread_mutex_lock(f->milk_lock);
            f->milk_qty -= m;
            pthread_mutex_unlock(f->milk_lock);
        } else {
            printf("Breakfast is ruined\n");
            pthread_mutex_unlock(f->milk_lock);
        }
    } else {
        printf("Breakfast is ruined\n");
        pthread_mutex_unlock(f->egg_lock);
    }
}
```

Fine-grained locking and atomicity


- Must be aware of the program semantics to correctly use fine-grained locking and guarantee atomicity where race conditions are possible
- Solutions:
 - Restructure the code if possible
 - Lock larger sections to guarantee atomicity
- Let's fix the example...

Fine-grained locking and atomicity

```
MorningRoutine(fridge *f, int e, double m) {  
    pthread_mutex_lock(f->egg_lock);  
    int eggs_left = f->egg_count;  
    if(eggs_left > 0) {  
  
        pthread_mutex_lock(f->milk_lock);  
        double milk_left = f->milk_qty;  
        if(milk_left > 0) {  
  
            f->egg_count -= e;  
            pthread_mutex_unlock(f->egg_lock);  
  
            f->milk_qty -= m;  
            pthread_mutex_unlock(f->milk_lock);  
        } else {  
            printf("Breakfast is ruined\n");  
            pthread_mutex_unlock(f->milk_lock);  
        }  
    } else {  
        printf("Breakfast is ruined\n");  
        pthread_mutex_unlock(f->egg_lock);  
    }  
}
```

a bit more coarser parallelism

The egg_lock might
never get unlocked if a
thread makes it to this
line



Other problems with fine-grained locking

- If we have two morning routines, eating egg then milk, eating milk then egg

```
MorningRoutine_1(fridge *f, int e, double m) {
    pthread_mutex_lock(f->egg_lock);
    int eggs_left = f->egg_count;
    if(eggs_left > 0) {

        pthread_mutex_lock(f->milk_lock);
        double milk_left = f->milk_qty;
        if(milk_left > 0) {

            f->egg_count -= e;
            pthread_mutex_unlock(f->egg_lock);

            f->milk_qty -= m;
            pthread_mutex_unlock(f->milk_lock);

        } else {
            printf("Breakfast is ruined\n");
            pthread_mutex_unlock(f->milk_lock);
            pthread_mutex_unlock(f->egg_lock);
        }

    } else {
        printf("Breakfast is ruined\n");
        pthread_mutex_unlock(f->egg_lock);
    }
}
```

```
MorningRoutine_2(fridge *f, int e, double m) {
    pthread_mutex_lock(f->milk_lock);
    double milk_left = f->milk_qty;
    if(milk_left > 0) {

        pthread_mutex_lock(f->egg_lock);
        int eggs_left = f->egg_count;
        if(eggs_left > 0) {

            f->milk_qty -= m;
            pthread_mutex_unlock(f->milk_lock);

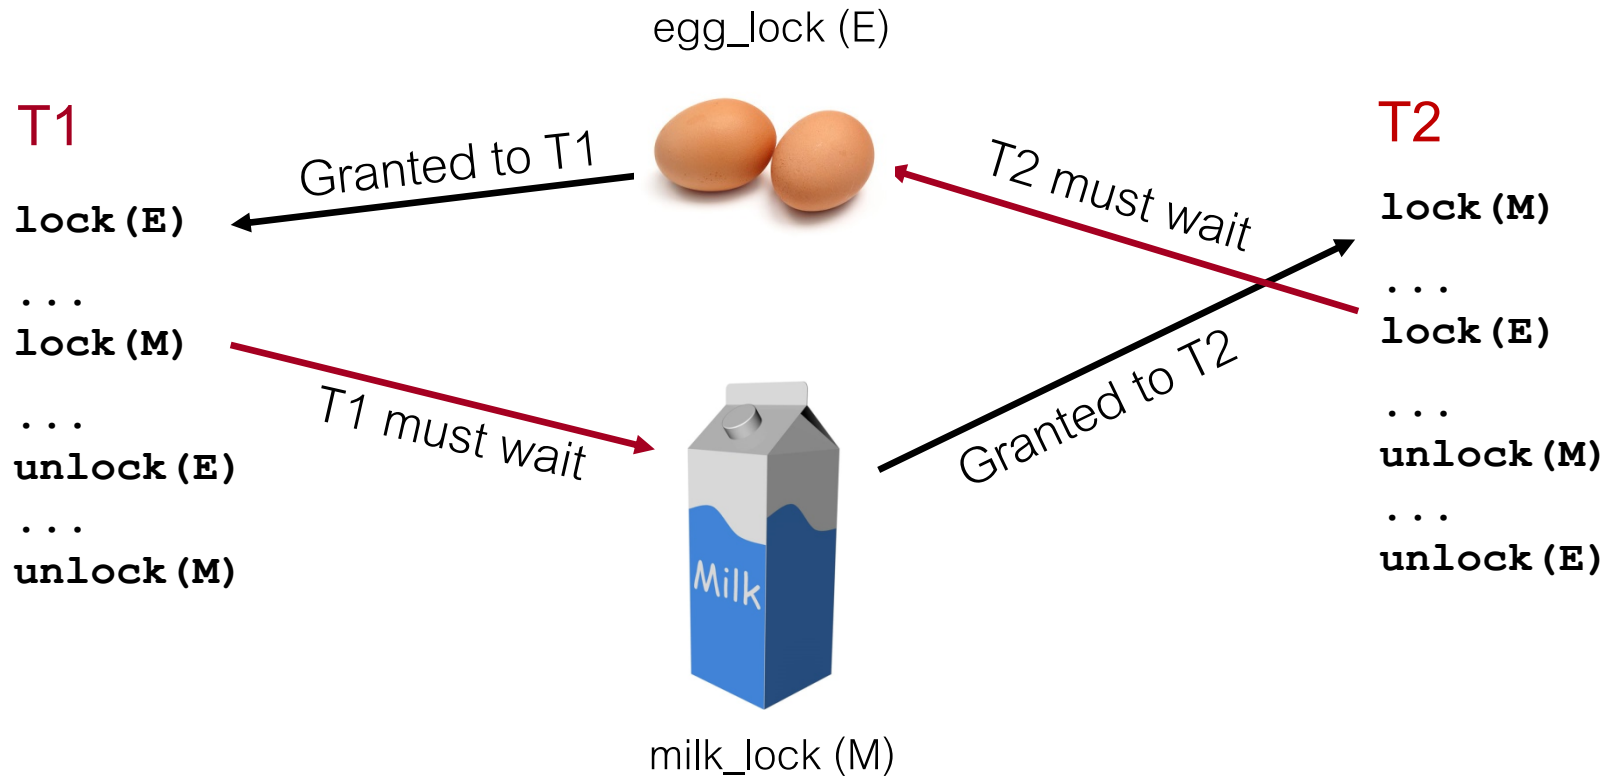
            f->egg_count -= e;
            pthread_mutex_unlock(f->egg_lock);

        } else {
            printf("Breakfast is ruined\n");
            pthread_mutex_unlock(f->milk_lock);
            pthread_mutex_unlock(f->egg_lock);
        }

    } else {
        printf("Breakfast is ruined\n");
        pthread_mutex_unlock(f->milk_lock);
    }
}
```

Deadlocks

- The **mutual** blocking of a set of threads (or processes)
- Each process/thread in the set is blocked, waiting for a lock which can only be unlocked by another process/thread in the set



- Simplest way to **break the deadlock**: always acquire locks in the same order!
 - Must enforce the **same ordering** in every piece of code where we acquire more than 1 lock

Next up...

- Overheads of locking
- Barrier construct

Overheads of locking

- When threads access the same locks => **lock contention!**
- If lots of threads are contending for the same lock, impacts performance
- Example: simple access to a shared counter by a handful of threads

```
void* do_work(void* arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&mutex);  
        counter++;  
        pthread_mutex_unlock(&mutex);  
    }  
    return NULL;  
}
```

- Even when no lock contention, acquiring a lock has overheads
 - Try with 1 thread and lots of iterations, with and without the mutex

"Localize" your computations

- Idea: Compute as much as possible locally, use synchronization scarcely
 - Mind you, do not break mutual exclusion when needed for correctness!
- Example:

```
void* do_work(void* arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&mutex);  
        counter++;  
        pthread_mutex_unlock(&mutex);  
    }  
    return NULL;  
}
```

- Idea: use local counter, update the global shared counter much more rarely...

Other synchronization primitives

- Semaphores
- Condition variables
- We'll focus only on the latter - powerful semantics

Barriers

- Threads that reach the barrier stop until all threads have reached it as well
 - If execution has stages, barrier ensures that data needed from a previous stage is ready
- POSIX has built-in barrier implementation
 - `int pthread_barrier_init(pthread_barrier_t *barrier, const pthread_barrier_attr_t *attr, unsigned count);`
 - `int pthread_barrier_wait(pthread_barrier_t *barrier);`
 - `int pthread_barrier_destroy(pthread_barrier_t* barrier);`
 - Check pthread documentation for more details...