

Operating Systems

Sina Meraji
U of T





Announcement

- Check discussion board for announcements
- A1 is posted

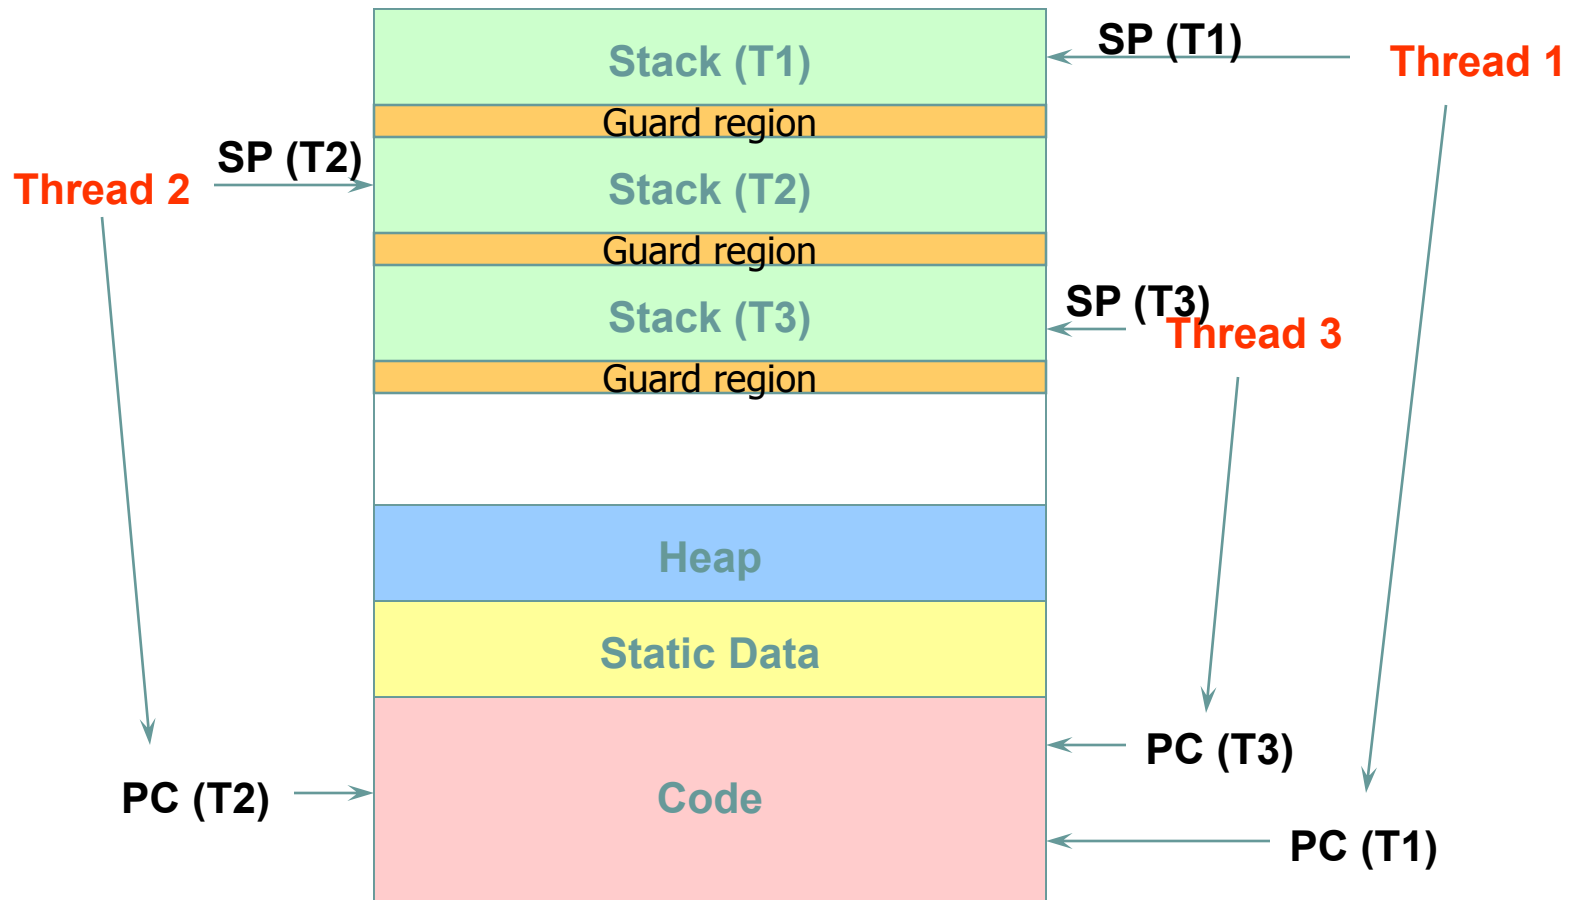
Recap:

Process Creation: Unix




- In Unix, processes are created using `fork()`
`int fork()`
- `fork()`
 - Creates a new address space
 - Initializes the address space with a **copy** of the entire contents of the address space of the parent
 - Initializes the kernel resources to point to the resources used by parent (e.g., open files)

Recap: Threads



TODAY:

- System Calls
- Intro to Synchronization



Bootstrapping



- Hardware stores small program in non-volatile memory
 - BIOS – Basic Input Output System
 - Knows how to access simple hardware devices
 - Disk, keyboard, display
- When power is first supplied, this program executes
- What does it do?
 - Checks that RAM, keyboard, and basic devices are installed and functioning correctly
 - Scans buses to detect attached devices and configures new ones
 - Determines boot device (tries list of devices in order)
 - Reads first sector from boot device and executes it (bootloader)
 - Bootloader reads partition table, finds active partition, reads secondary bootloader
active partition where OS resides
 - Secondary bootloader reads OS into memory and executes it

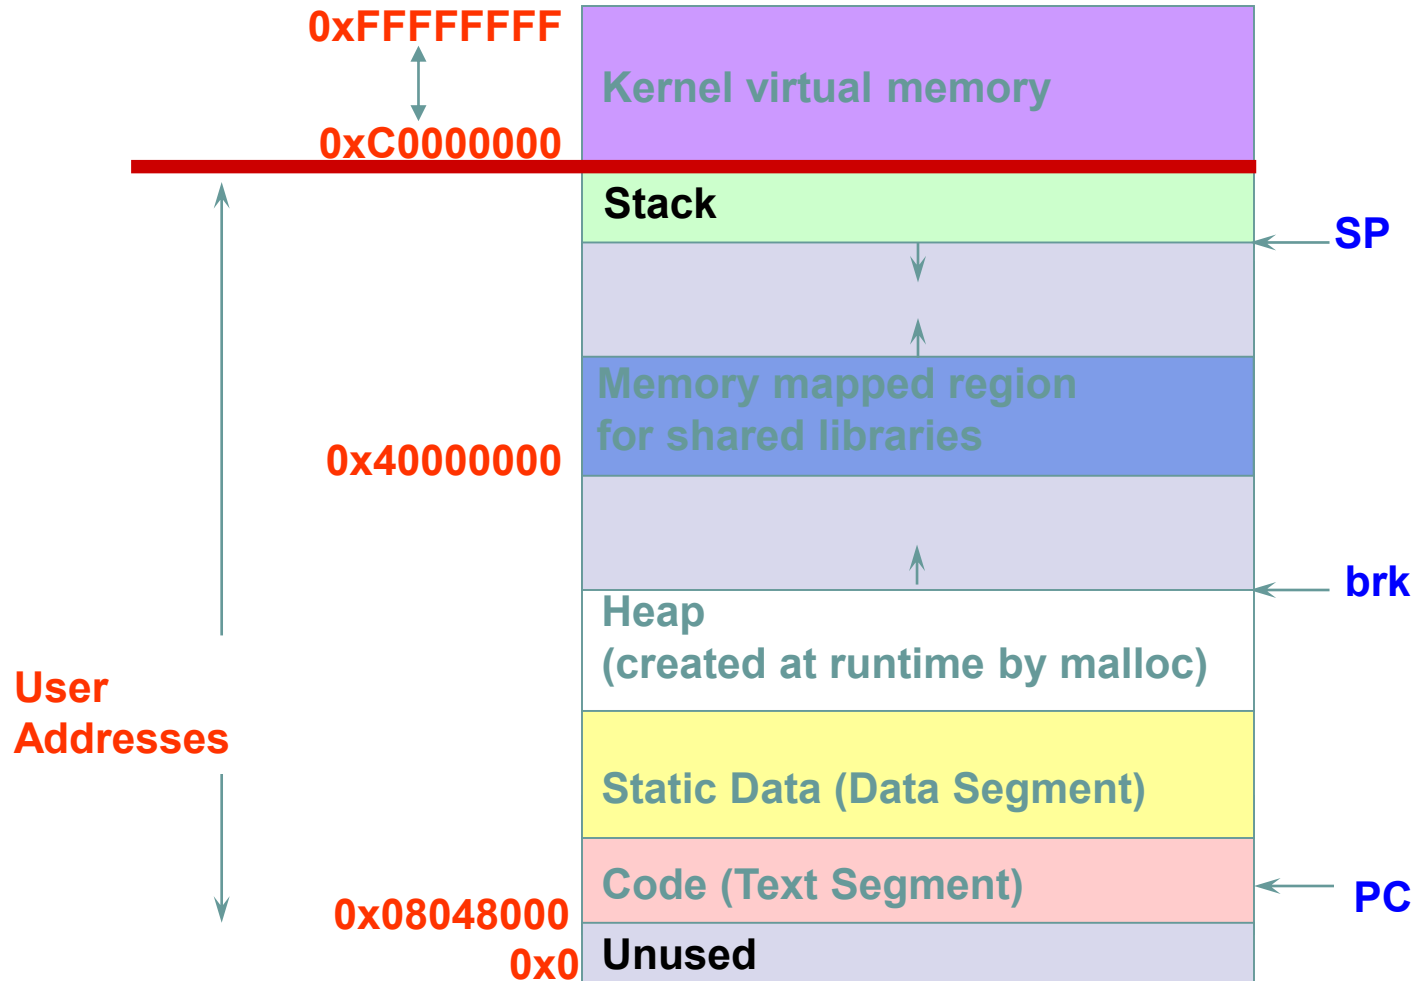
creates process for OS, brings OS from hard disk to memory
since OS is in a process. OS has its own pcb, heap, stack, ...

Operating System Startup

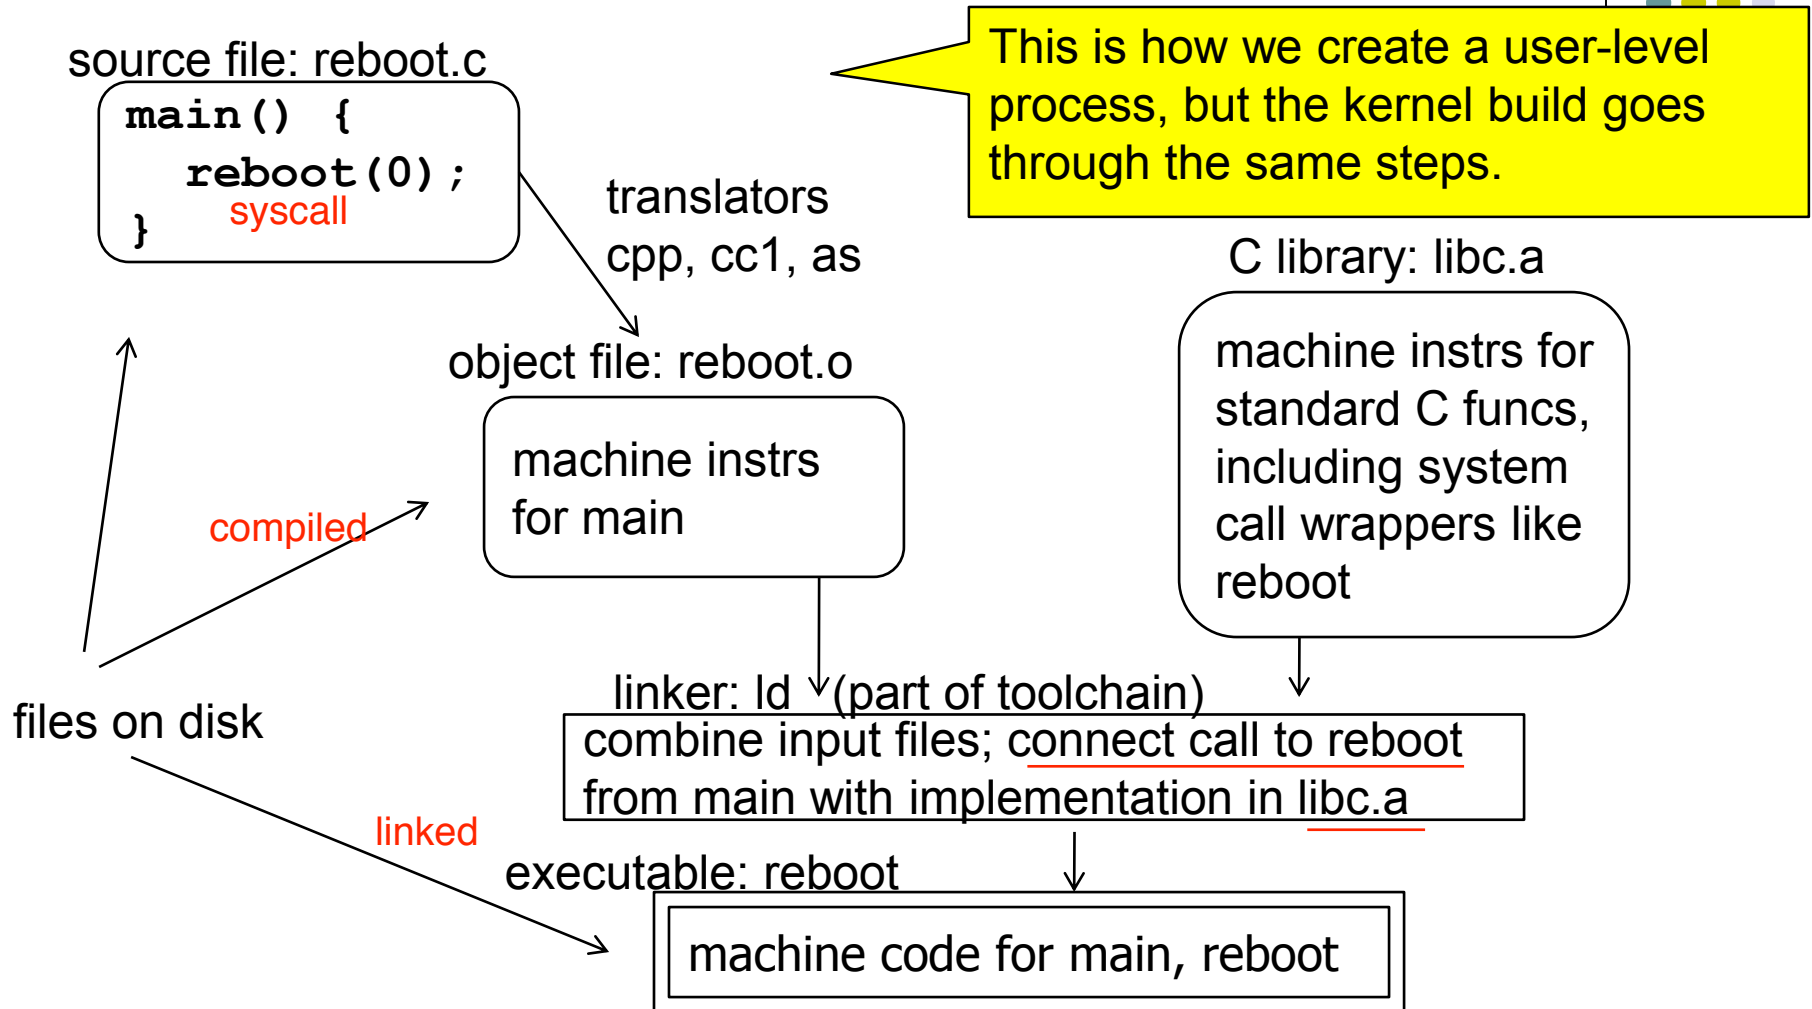


- Machine starts in system mode, so kernel code can execute immediately
- OS initialization:
 - Initialize internal data structures
 - Machine dependent operations are typically done first
 - Create first process
 - Switch mode to user and start running first process when login occurs
 - Wait for something to happen
 - OS is entirely driven by external events

Memory Layout (Linux, x86)



From Program to Process... 1




Unix Shells

syscall

+ execution mode will change from user to kernel mode

+ hence permission matters here



```
while (1) {  
    char *cmd = read_command();  
    int child_pid = fork();  
    if (child_pid == 0) {  
        exec(cmd); //cmd=executable name(reboot)  
    } else {  
        wait(child_pid);  
    }  
}
```

will likely give permission error



Process Creation: Unix (2)

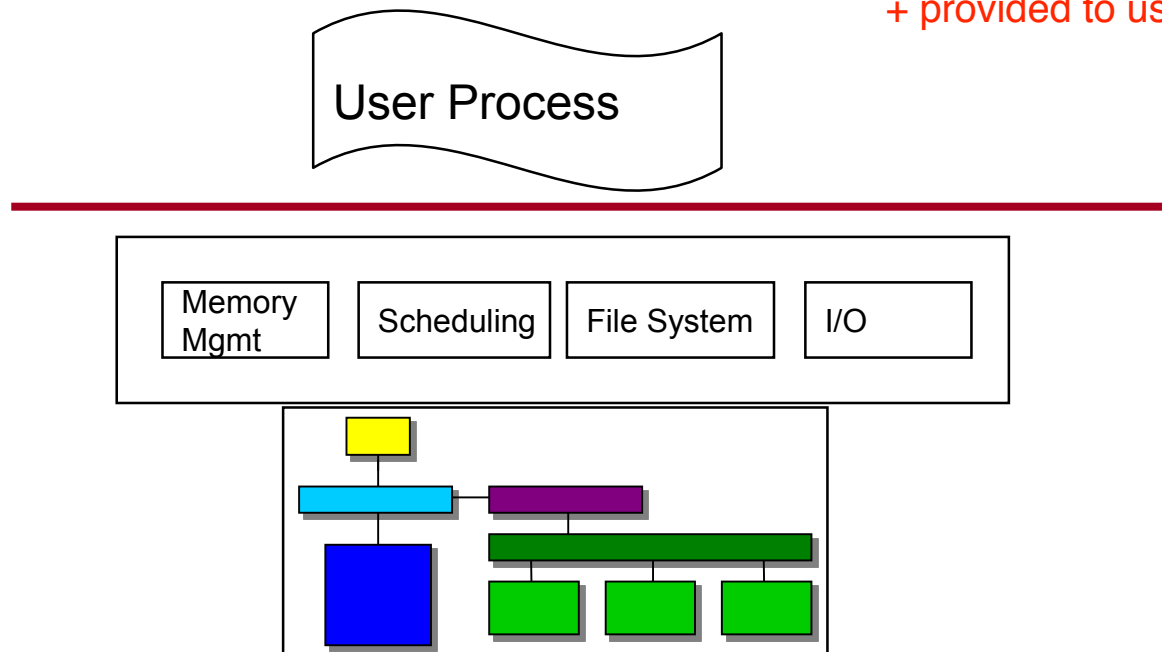
- Wait a sec ... How do we actually start a new program?
`int exec(char *prog, char *argv[])`
- `exec()`
 - Stops the current process only the CODE portion replaced
 - Loads the program “prog” into the process’ address space
 - Initializes hardware context and args for the new program
 - Places the PCB onto the ready queue
 - Note: It **does not** create a new process



Requesting OS Services

- Operating System and user programs are isolated from each other
- But OS provides service to user programs...
- So, how do they communicate?

communicate
1. syscall
+ provided to user






Boundary Crossings

- Getting to kernel mode
 - Boot time (not really a crossing, starts in kernel)
 - Explicit system call – request for service by application
 - Hardware interrupt
 - Software trap or exception i.e. division by 0, invalid mem access
 - Hardware has table of “Interrupt service routines”
table contains map from syscall name to function
- Kernel to user
 - Jumps to next application instruction
return from successful syscall

a software routine that hardware invokes in response to an interrupt. ISRs examine an interrupt and determine how to handle it.

System Calls for Process Management




Process management

Call	Description
<code>pid = fork()</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

Some of the major system calls.

System Calls for File Management




File management

Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing, or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &buf)</code>	Get a file's status information

Some of the major system calls.


System Calls



Read(fd, buffer, nbytes).

trap to kernel
+ Uses a hashtable which
maps from syscall name to
address of function

System Call Interface



- User program calls C library function with arguments
- C library function arranges to pass arguments to OS, including a system call identifier
- Executes special instruction to trap to system mode
 - Interrupt/trap vector transfers control to a system call handling routine
- Syscall handler figures out which system call is needed and calls a routine for that operation
- How does this differ from a normal C language function call? Why is it done this way?
 - Extra level of indirection through system call handler, rather than direct control flow to called function
 - Hardware support is needed to enforce separation of userspace and kernel

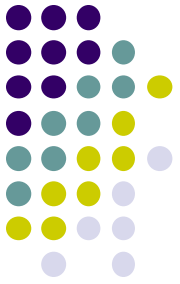
System Call Operation




- Kernel must verify arguments that it is passed
 - Why?
- A fixed number of arguments can be passed in registers
 - Often pass the address of a user buffer containing data (e.g., for write())
 - Kernel must copy data from user space into its own buffers
 - Result of system call is returned in register

hence data not passed directly but
address to which data will be stored

Intro to Synchronization



Cooperating Processes



- A process is *independent* if it cannot affect or be affected by the other processes executing in the system
- No data sharing \Rightarrow process is independent
- A process is *cooperating* if it is not independent
- Cooperating processes must be able to communicate with each other and to synchronize their actions



Interprocess Communication

- Cooperating processes need to exchange information, using either
 - Shared memory (e.g. `fork()`)
 - Message passing
- Message passing models
 - `Send(P, msg)` – send msg to process P
 - `Receive(Q, msg)` – receive msg from process Q



Motivating Example

- Suppose we write functions to handle withdrawals and deposits to a bank account:

update and return the balance to shared account on > 1 processes/threads

```
Withdraw(acct, amt) {  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

```
Deposit(account, amount) {  
    balance = get_balance(acct);  
    balance = balance + amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

shared: each account may be accessed by more than one instance⁵

- Idea: Create separate threads for each action, which may run at the bank's central server
- What's wrong with this implementation?
 - Think about potential schedules for these two threads



Motivating Example

- Suppose we write functions to handle withdrawals and deposits to a bank account:

```
Withdraw(acct, amt) {  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

```
Deposit(account, amount) {  
    balance = get_balance(acct);  
    balance = balance + amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

- Suppose you share this account with someone and the balance is \$1000
- You each go to separate ATM machines - you withdraw \$100 and your S.O. deposits \$100

Interleaved Schedules



- The problem is that the execution of the two processes can be interleaved:

Schedule A

```
balance = get_balance(acct);  
balance = balance - amt;
```

```
balance = get_balance(acct);  
balance = balance + amt;  
put_balance(acct, balance);
```


Context
switch

```
put_balance(acct, balance);
```

900

- What is the account balance now?
- Is the bank happy with our implementation?
 - Are you?

Interleaved Schedules



- The problem is that the execution of the two processes can be interleaved:

Schedule A

```
balance = get_balance(acct);  
balance = balance - amt;
```

```
balance = get_balance(acct);  
balance = balance + amt;  
put_balance(acct, balance);
```

```
put_balance(acct, balance);
```

900

Context
switch

Schedule B

```
balance = get_balance(acct);  
balance = balance - amt;
```

```
balance = get_balance(acct);  
balance = balance + amt;
```

```
put_balance(acct, balance);
```

```
put_balance(acct, balance);
```

1100

- What is the account balance now?
- Is the bank happy with our implementation?
 - Are you?



What Went Wrong

- Two concurrent threads manipulated a shared resource (the account) without any synchronization
 - Outcome depends on the order in which accesses take place
 - This is called a *race condition* result depends on execution order
- We need to ensure that only one thread at a time can manipulate the shared resource
 - So that we can reason about program behavior
 - We need *synchronization*



Example continued ...

- Could the same problem occur with a simple shared variable:
 - T_1 and T_2 share variable X
 - T_1 increments X ($X := X+1$)
 - T_2 decrements X ($X := X-1$)
 - At the machine level, we have:

changing value of a variable take multiple steps.


```
T1:   LOAD X
      INCR
      STORE X
```

```
T2:   LOAD X
      DECR
      STORE X
```

- Same problem of interleaving can occur!

atomic operation: execute or cannot execute at all,
i.e. CPU guarantees to finish execution

Mutual Exclusion




- Given:
 - A set of n threads, T_0, T_1, \dots, T_n
 - A set of resources shared between threads
 - A segment of code which accesses the shared resources, called the *critical section, CS*

```
Withdraw(acct, amt) {  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

CS

- We want to ensure that:
 - Only one thread at a time can execute in the critical section
 - All other threads are forced to wait on entry
 - When a thread leaves the CS, another can enter

Aside: What program data is shared between threads?




- Local variables are not shared (*private*)
 - Each thread has its own stack
 - Local vars are allocated on this private stack
- Global variables and static objects are *shared*
 - Stored in the static data segment, accessible by any thread
- Dynamic objects and other heap objs are *shared*
 - Allocated from heap with malloc/free or new/delete



The Critical Section Problem

- Design a protocol that threads can use to cooperate




- Each thread must request permission to enter its CS, in its *entry* section
- CS may be followed by an *exit* section
- Remaining code is the *remainder* section

Critical Section Requirements (1)



- Design a protocol that threads can use to cooperate




1) Mutual Exclusion

- If one thread is in the CS, then no other is




Critical Section Requirements (2)

- Design a protocol that threads can use to cooperate




2) Progress

- If no thread is in the CS, and some threads want to enter CS, it should be able to enter in definite time



Critical Section Requirements (3)

- Design a protocol that threads can use to cooperate




3) Bounded waiting (no starvation)

- If some thread T is waiting on the CS, then there is a limit on the number of times other threads can enter CS before this thread is granted access hence waiting time is approx. uniform.



Critical Section Requirements (4)

- Design a protocol that threads can use to cooperate



4) Performance

- The overhead of entering and exiting the CS is small with respect to the work being done within it




Critical Section Requirements

- 1) Mutual Exclusion
 - If one thread is in the CS, then no other is
- 2) Progress
 - If no thread is in the CS, and some threads want to enter CS, it should be able to enter in definite time
- 3) Bounded waiting (no starvation)
 - If some thread T is waiting on the CS, then there is a limit on the number of times other threads can enter CS before this thread is granted access
- Performance
 - The overhead of entering and exiting the CS is small with respect to the work being done within it

Some Assumptions & Notation



- Assume no special hardware instructions, no restrictions on the # of processors (for now)
- Assume that basic machine language instructions (LOAD, STORE, etc.) are atomic:
 - If two such instructions are executed concurrently, the result is equivalent to their sequential execution in some unknown order
- If only two threads, we number them T_0 and T_1
 - Use T_i to refer to one thread, T_j for the other ($j=1-i$) when the exact numbering doesn't matter
to reflect that order is unknown
- Let's look at one solution...



2-Thread Solutions: 1st Try

- Let the threads share an integer variable *turn* initialized to 0 (or 1)
- If $turn=i$, thread T_i is allowed into its CS

```
My_work(id_t id) { /* id_t can be 0 or 1 */  
    ...  
    while (turn != id) ; /* entry section */  
    /* critical section, access protected resource */  
    turn = 1 - id; /* exit section */  
    ... /* remainder section */  
}
```


- ✓ Only one thread at a time can be in its CS
- ✗ Progress is not satisfied strictly alternating. if T_i has a lot of task while T_j has one. Progress stagnate after one alteration.
 - Requires strict alternation of threads in their CS: if $turn=0$, T_1 may not enter, even if T_0 is in the code section



2-Thread Solutions: 2nd Try

- First attempt does not have enough info about state of each process. It only remembers which process is allowed to enter its CS
i.e. the value of Turn
- Replace turn with a shared flag for each thread
 - **boolean** `flag[2] = {false, false}`
 - Each thread may update its own flag, and read the other thread's flag
 - If `flag[i]` is true, T_i is ready to enter its CS

A Closer Look at 2nd Attempt



```
My_work(id_t id) { /* id can be 0 or 1 */
    ...           check the other thread
    while (flag[1-id]) ; /* entry section */
    flag[id] = true; /* indicate entering CS */
    /* critical section, access protected resource */
    flag[id] = false; /* exit section */
    ...           /* remainder section */
}
```

while loops indefinitely if any other thread is true;
and enter CS if other threads has flag=false.

- Mutual exclusion is not guaranteed
 - Each thread executes *while* statement, finds *flag* set to false at the same time. Context switching right after while check
 - Each thread sets own *flag* to *true* and enters CS
- Can't fix this by changing order of testing and setting *flag* variables (leads to *deadlock*)

if flag both set to true. check other's flag, true, hence both threads stay in while loop

2-Thread Solutions: 3rd Try



- Combine key ideas of first two attempts for a correct solution
- The threads share the variables *turn* and *flag* (where *flag* is an array, as before)

```
Enter_region(id_t id) {      /* id can be 0 or 1 */  
  
    flag[id] = true;        /* indicate entering CS */  
    turn = id;              turn <- intend to enter  
    while (turn == id && flag[other] == true);  
                                loops if other is in CS...  
}
```

```
Leave_region(id_t id) {      /* id can be 0 or 1 */  
  
    flag[id] = false;  
                                turn not set on exit, since turn reflects other threads' intention  
}
```


2-Thread Solutions: 3rd Try



- Imagine two threads *i* and *j* execute `Enter_region()` at the same time:

Thread *i*

Thread *j*

```
flag[i] = true;  
turn = i;  
while(turn==i && flag[j]==true);
```

```
flag[j] = true;  
turn = j;  
while(turn==j && flag[i]==true);
```

- Basic idea: if both try to enter at the same time, *turn* will be set to both 0 and 1 at roughly the same time. Only one assignments will last. The final value of *turn* decides who gets to go first. since turn holds 1 value at any time
- This is the basis of *Peterson's Algorithm*

Peterson's Solution



```
#define FALSE 0
#define TRUE 1
#define N      2                /* number of processes */

int turn;                       /* whose turn is it? */
int interested[N];              /* all values initially 0 (FALSE) */

void enter_region(int process); /* process is 0 or 1 */
{
    int other;                  /* number of the other process */

    other = 1 - process;        /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;             /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process) /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

Peterson's solution for achieving mutual exclusion.



Higher-level Abstractions for CS's

- Locks
 - Very primitive, minimal semantics
- Semaphores
 - Basic, easy to understand, hard to program with
- Monitors
 - High-level, ideally has language support (Java)
- Messages
 - Simple model for communication & synchronization
 - Direct application to distributed systems



Synchronization Hardware

- To build these higher-level abstractions, it is useful to have some help from the hardware
- On a uniprocessor:
 - Disable interrupts before entering critical section
 - Prevents context switches
 - Doesn't work on multiprocessor
- Need some special atomic instructions

Atomic Instructions:

Test-and-Set Lock (TSL)



- Test-and-set uses a *lock* variable
 - Lock == 0 => nobody is using the lock
 - Lock == 1 => lock is in use
 - In order to acquire lock, must change it's value from 0=>1

provided on single processor machine for implementing locks at hardware level

```
boolean test_and_set(boolean *lock)
{
    boolean old = *lock;
    *lock = True;
    return old;
}
```

Test if lock is available,

if lock is not available, returns false -> implies no function called before current
if lock is available, returns true -> implies no other function called before current
as sideeffect, lock always set to true

- Hardware executes this atomically!



Atomic Instructions: Test-and-Set

- The semantics of test-and-set are:
 - Record the old value of the variable
 - Set the variable to some non-zero value
 - Return the old value

```
boolean test_and_set(boolean *lock)
{
    boolean old = *lock;
    *lock = True;
    return old;
}
```

- *lock* is always True on exit from test-and-set
 - Either it was *True* (locked) already, and nothing changed
 - or it was *False* (available), but the caller now holds it
- Return value is either *True* if it was locked already, or *False* if it was previously available

A Lock Implementation



- There are two operations on locks: *acquire()* and *release()*

```
boolean lock;  
  
void acquire(boolean *lock) {  
    while (test_and_set(lock)) ;  
} lock taken -> returns true -> while loop blocks  
lock not taken -> returns false -> while loop finishes  
  
void release(boolean *lock) {  
    *lock = false;  
}
```

- This is a *spinlock*
 - Uses *busy waiting* - thread continually executes *while* loop in *acquire()* , consumes CPU cycles

Using Locks



Function Definitions

```
Withdraw(acct, amt) {  
  
    acquire(lock) ;  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    release(lock) ;  
    return balance;  
}
```

```
Deposit(account, amount) {  
  
    acquire(lock) ;  
    balance = get_balance(acct);  
    balance = balance + amt;  
    put_balance(acct, balance);  
    release(lock) ;  
    return balance;  
}
```

Possible schedule


```
acquire(lock) ;  
balance = get_balance(acct);  
balance = balance - amt;
```

```
acquire(lock) ;
```

```
put_balance(acct, balance);  
release(lock) ;
```

```
balance = get_balance(acct);  
balance = balance + amt;  
put_balance(acct, balance);  
release(lock) ;
```


Next Week



- More on Synchronization

Announcement



- Check course website regularly
- Attend Tutorials