

Synchronization Primitives

Synchronization Mechanisms

- Locks
 - Very primitive constructs with minimal semantics
- Semaphores
 - A generalization of locks
 - Easy to understand, hard to program with
- Condition Variables
 - Constructs used in implementing *monitors* (more on this later...)

Locks

- Synchronization mechanisms with 2 operations: acquire(), and release()
- In simplest terms: an object associated with a particular critical section that you need to “own” if you wish to execute in that region
- Simple semantics to provide mutual exclusion:
 acquire(lock);
 //CRITICAL SECTION
 release(lock);
- Downsides: spinlock has performance issues
 - Can cause deadlock if not careful i.e. not releasing lock properly
 - Cannot allow multiple concurrent accesses to a resource

POSIX Locks

- POSIX locks are called mutexes (since locks provide mutual exclusion...)

- A few calls associated with POSIX mutexes:

pthread_mutex_init (mutex, attr)

- Initialize a mutex

pthread_mutex_destroy (mutex)

- Destroy a mutex

pthread_mutex_lock (mutex) blocked if lock not available

- Acquire the lock

pthread_mutex_trylock(mutex) return value indicates if lock is available, up to programmer to decide what to do

- Try to acquire the lock (more on this later...)

pthread_mutex_unlock (mutex)

- Release the lock

Initializing & Destroying POSIX Mutexes

- POSIX mutexes can be created statically or dynamically
 - Statically, using PTHREAD_MUTEX_INITIALIZER

```
pthread_mutex_t mx = PTHREAD_MUTEX_INITIALIZER;
```

 - Will initialize the mutex with default attributes
 - Only use for static mutexes; no error checking is performed
 - Dynamically, using the pthread_mutex_init call

```
int pthread_mutex_init(pthread_mutex_t * mutex, const  
pthread_mutexattr_t * attr);
```

 - mutex: the mutex to be initialized
 - attr: structure whose contents are used at mutex creation to determine the mutex's attributes
 - Same idea as pthread_attr_t attributes for threads
- Destroy using pthread_mutex_destroy

```
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

 - mutex: the mutex to be destroyed
 - Make sure it's unlocked! (destroying a locked mutex leads to undefined behaviour...)

Acquiring and Releasing POSIX Locks

- Acquire

int pthread_mutex_lock(pthread_mutex_t *mutex);

- mutex: the mutex to lock (acquire)
- If mutex is already locked by another thread, the call will block until the mutex is unlocked

int pthread_mutex_trylock(pthread_mutex_t *mutex);

- mutex: the mutex to TRY to lock (acquire)
- If mutex is already locked by another thread, the call will return a “busy” error code (EBUSY)

- Release

int pthread_mutex_unlock(pthread_mutex_t *mutex);

- mutex: the mutex to unlock (release)

Banking Example

- Bank account balance maintained in one variable `int balance`
- Transactions: deposit or withdraw some amount from the account (+/- balance)
- Unprotected, concurrenented accesses to your balance could create race conditions

Banking Example

- Thread 1 withdraws 100
- Thread 2 withdraws 100

```
int new_balance = balance –  
amount;
```

```
balance = new_balance;
```

```
int new_balance = balance –  
amount;
```

```
balance = new_balance;
```

- End with balance – 100 instead of balance – 200
- Bank error in your favour? Could be the other way around!
- Idea: put a lock around the code that modifies balance so only a single thread accesses it at any given time

Banking Example

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS200
int balance=0;
pthread_mutex_t bal_mutex;

int main (int argc, char *argv[]){
    pthread_t thread[NUM_THREADS];
    int rc;
    long t;
    void *status;

    pthread_mutex_init(&bal_mutex, NULL);
    for(t=0; t<NUM_THREADS; t+=2) {
        rc = pthread_create(&thread[t], NULL, deposit, (void *)10);
        if (rc) {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
        rc = pthread_create(&thread[t+1], NULL, widthdraw, (void *)10);
        if (rc) {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    (...)
}
```

Banking Example

(...)

```
for(t=0; t<NUM_THREADS; t++) {  
    rc = pthread_join(thread[t], &status);  
    if (rc) {  
        printf("ERROR; return code from pthread_join() is %d\n", rc);  
        exit(-1);  
    }  
}  
  
printf("Final Balance is %d.\n", balance);  
pthread_exit(NULL);  
}
```

Banking Example - Transactions

```
void *deposit(void *amt){
```

```
    pthread_mutex_lock(&bal_mutex);
```

```
    //CRITICAL SECTION
```

```
    int amount = (int)amt;
```

```
    int new_balance = balance +  
    amount;
```

```
    balance = new_balance;
```

```
    pthread_mutex_unlock(&bal_mutex)  
    ;
```

```
    pthread_exit((void *)0);
```

```
}
```

```
void *withdraw(void *amt){
```

```
    pthread_mutex_lock(&bal_mutex);
```

```
    //CRITICAL SECTION
```

```
    int amount = (int)amt;
```

```
    int new_balance = balance -  
    amount;
```

```
    balance = new_balance;
```

```
    pthread_mutex_unlock(&bal_mutex)  
    ;
```

```
    pthread_exit((void *)0);
```

```
}
```

Semaphore

- Synchronization mechanism that generalizes locks to more than just “acquired” and “free” (or “released”)
- A semaphore provides you with:
 - An integer count accessed through 2 atomic operations
 - Wait - aka: down, decrement, P (for proberen)
 - taking a spot
 - Block until semaphore is free, then decrement the variable
 - Signal - aka: up, post, increment, V (for verhogen)
 - giving up a spot
 - Increment the variable and unblock a waiting thread (if there are any)
- A mutex was just a binary semaphore (remember pthread_mutex_lock blocked if another thread was holding the lock)
 - mutex = binary semaphore !
- A queue of waiting threads

POSIX Semaphores

- Declared in semaphore.h
- A few calls associated with POSIX semaphores:
 - sem_init
 - Initialize the semaphore
 - sem_wait
 - Wait on the semaphore (decrement value)
 - sem_post
 - Signal (post) on the semaphore (increment value)
 - sem_getvalue
 - Get the current value of the semaphore
 - sem_destroy
 - Destroy the semaphore

Initializing & Destroying POSIX Semaphores

- Initialize semaphores using `sem_init`

`int sem_init(sem_t *sem, int pshared, unsigned int value);`

- `sem`: the semaphore to initialize
- `pshared`: non-zero to share between processes defaults to be shared between threads
- `value`: initial count value of the semaphore the integer count

- Destroy semaphores using `sem_destroy`

`int sem_destroy(sem_t *sem);`

- `sem`: semaphore to destroy
- Semaphore must have been created using `sem_init`
- Destroying a semaphore that has threads blocked on it is undefined.

every threads has to be unblocked

Decrementing & Incrementing POSIX Semaphores

- Decrement semaphores using `sem_wait`

`int sem_wait(sem_t *sem);`

- `sem`: the semaphore to decrement (wait on)

- Increment semaphores using `sem_post`

`int sem_post(sem_t *sem);`

- `sem`: semaphore to increment

- Let's look at an example of a very simple server simulation...

Server Example

```
(...)
#define NUM_THREADS200
#define NUM_RESOURCES10sem_t resource_sem; //Semaphore declaration

int main (int argc, char *argv[])
{ pthread_t thread[NUM_THREADS];
  int rc;
  int i;
  void *status;
  sem_init(&resource_sem, 0, NUM_RESOURCES); //Resource Semaphore

  for(i=0; i<NUM_THREADS; i++) {

    rc = pthread_create(&thread[i], NULL, handle_connection, (void *)i);
    if (rc) {
      printf("ERROR; return code from pthread_create() is %d\n", rc);
      exit(-1);
    }
  }
  (...)
  for(i=0; i<NUM_THREADS; i++) {
    rc = pthread_join(thread[i], &status);
    if (rc) {
      printf("ERROR; return code from pthread_join() is %d\n", rc);
      exit(-1);
    }
  }
  return 0;
} //End of main
```


Server Example – Connection Handler

```
void *handle_connection(void *client){  
    printf ("Handler for client %d created!\n", (int)client);  
  
    sem_wait(&resource_sem);  
  
    //DO WORK TO HANDLE CONNECTION HERE  
    sleep(1);  
    printf ("Done servicing client %d\n", (int) client);  
  
    sem_post(&resource_sem);  
  
    pthread_exit((void *)0);  
}
```

Condition Variables

- Another useful synchronization construct used in implementing monitors - only a single process execute inside the monitor
- Locks control thread access to data; condition variables allow threads to synchronize based on the value of the data.
- Alternative to condition variables is to constantly poll the variable (from the critical section)
 - BAD!
 - Ties up a lot of CPU resources
 - Could potentially lead to synchronization problems
- Monitors support suspending execution within the monitor
 - wait() (suspend the invoking process and release the lock)
 - signal() (resume exactly one suspended process)
 - broadcast() (resumes all suspended processes)
 - If no process is suspended, signal/broadcast has no effect (in contrast to semaphores, where signal always changes state of the semaphore)

POSIX Condition Variables

- POSIX condition variables: pthread_cond_t

- A few calls associated with POSIX CVs:

```
int pthread_cond_init(pthread_cond_t *cond, pthread_condattr_t *attr);
```

- Initialize a condition variable

```
int pthread_cond_destroy(pthread_cond_t *cond);
```

- Destroy a condition variable

```
int pthread_cond_wait (pthread_cond_t *cond, pthread_mutex_t *mutex);
```

- Wait on a condition variable

```
int pthread_cond_signal(pthread_cond_t *cond);
```

- Wake up one thread waiting on this condition variable

```
int pthread_cond_broadcast(pthread_cond_t *cond);
```

- Wake up all threads waiting on this condition variable

Using Condition Variables (from LLNL tutorial)

Main Thread

- Declare and initialize global data/variables which require synchronization (such as "count")
- Declare and initialize a condition variable object
- Declare and initialize an associated mutex
- Create threads A and B to do work

Thread A

- Do work up to the point where a certain condition must occur (such as "count" must reach a specified value)
- Lock associated mutex and check value of a global variable
- Call pthread_cond_wait() to perform a blocking wait for signal from Thread-B. Note that a call to pthread_cond_wait() automatically and atomically unlocks the associated mutex variable so that it can be used by Thread-B.
- When signalled, wake up. Mutex is automatically and atomically locked.
- Explicitly unlock mutex
- Continue

Thread B

- Do work
- Lock associated mutex
- Change the value of the global variable that Thread-A is waiting upon.
- Check value of the global Thread-A wait variable. If it fulfills the desired condition, signal Thread-A.
- Unlock mutex.
- Continue

Main Thread: Join / Continue

Monitors

- Locks
 - Provide mutual exclusion
 - 2 operations: acquire() and release()
- Semaphores
 - Generalize locks with an integer count variable and a thread queue
 - 2 operations: wait() and signal()
 - If the integer count is negative, threads wait in a queue until another thread signals the semaphore
- Monitors
 - An abstraction that encapsulates shared data and operations on it in such a way that only a single process at a time may be executing “in” the monitor

More on Monitors

- Programmer defines the scope of the monitor
 - ie: which data is “monitored”
- Local data can be accessed only by the monitor’s procedures (not by any external procedures)
- Before any monitor procedure may be invoked, mutual exclusion must be guaranteed
 - There is often a lock associated with each monitored object
- Other processes that attempt to enter the monitor are blocked. They must first acquire the lock before becoming active in the monitor

Complications With Monitors

- Complication
 - A process may need to wait for something to happen
 - Input from another thread might be necessary for example
 - The other thread may require access to the monitor to produce that event
- Solution?
 - Monitors support suspending execution within the monitor
 - wait() (suspend the invoking process and release the lock)
 - signal() (resume exactly one suspended process)
 - broadcast() (resumes all suspended processes)
 - If no process is suspended, signal/broadcast has no effect (in contrast to semaphores, where signal always changes state of the semaphore)

Monitor signal() ; who goes first?

- Suppose P executes a signal operation that would wake up a suspended process Q
 - Either process can continue execution, but both cannot simultaneously be active in the monitor
- Who goes first?
 - Hoare monitors: waiter first
 - signal() immediately switches from the caller to a waiting thread
 - Condition that the waiter was blocked on is guaranteed to hold when the waiter resumes
 - Mesa monitors: signaler first
 - signal() places a waiter on the ready queue, but signaler continues inside the monitor
 - Condition that the waiter was blocked on is not guaranteed to hold when the waiter resumes (must check again...)

Hoare vs. Mesa Monitors

- Hoare monitor wait

```
if(...){  
    wait(cv, lock);  
}
```

- Mesa monitor wait

```
while(...){  
    wait(cv, lock);  
}
```

- Tradeoffs

- Hoare monitors are easier to reason with, but hard to implement
- Mesa monitors are easier to implement, and support additional operations like broadcast()

Monitor Example - Bounded Buffers

- We have a buffer of limited size N
 - Producers add to the buffer if it is not full
 - Consumers remove from the buffer if it is not empty
- Want to control buffer as a monitor
 - Buffer can only be accessed by methods that are “part of” the monitor, that only give one producer or consumer access to the buffer at a time

i.e. read/write buffer are procedures inside of monitor
- Need 2 functions
 - `add_to_buffer()`
 - `remove_from_buffer()`
- Need
 - One lock
 - Two conditions
 - One for producers to wait
 - One for consumers to wait

Monitor Example - Bounded Buffers

```
#define N 100

typedef struct buf_s {
    int data[N];
    int inpos; /* producer inserts here */
    int outpos; /* consumer removes from here */
    int numelements; /* # items in buffer */
    struct lock *mylock; /* access to monitor */
    struct cv *notFull; /* for producers to wait */
    struct cv *notEmpty; /* for consumers to wait */
} buf_t;

buf_t buffer;


void add_to_buff(int value);

int remove_from_buff();
```

Monitor Example - Bounded Buffers

```
void add_to_buf(int value) {  
    lock_acquire(buffer.mylock);  
    while (nelements == N) {  
        /* buffer is full, wait */  
        cv_wait(buffer.notFull, buffer.mylock);  
    }  
    buf.data[inpos] = value;  
    inpos = (inpos + 1) % N;  
    nelements++;  
    cv_signal(buffer.notEmpty, buffer.mylock);  
    lock_release(buffer.mylock);  
}
```

mesa monitor



What kind of
monitor is this?

Monitor Example - Bounded Buffers

```
int remove_from_buf() {  
    int val;  
    lock_acquire(buffer.mylock);  
    while (nelements == 0) {  
        /* buffer is empty, wait */  
        cv_wait(buffer.notEmpty, buffer.mylock);  
    }  
    val = buf.data[outpos];  
    outpos = (outpos + 1) % N;  
    nelements--;  
    cv_signal(buffer.notFull, buffer.mylock);  
    lock_release(buffer.mylock);  
}
```