

STA 414/2104:

Lecture 8

5 March 2018:
Continuous Latent Variable Models,
Neural networks

Delivered by Mark Ebden
With thanks to Russ Salakhutdinov,
Jimmy Ba and others

Announcements

- Midterm regrading requests:
 - You should receive a reply later this week
 - Today (5 March) is the request deadline
- Next week's topics:
 - Graphical Models
 - Modelling Sequential Data

Outline

- Continuous latent variable models
 - Background
 - PCA
- Neural networks
 - Introduction
 - Autoencoders
 - Learning neural networks

Reminder from last week: latent variables

“Latent variables are entities that we invent to explain patterns we see in observable variables – for instance, doctors have invented *diseases* to explain commonalities of symptoms seen in patients.”

– Radford Neal

The values of these latent variables are inferred from those of observable variables.

Continuous Latent Variable Models

- Often there are some **unknown underlying causes** of the data
- So far we have looked at models with discrete latent variables, such as the mixture of Gaussians
- Sometimes, it is more appropriate to think in terms of **continuous factors** which control the data we observe
- **Motivation**: for many datasets, data points lie close to a manifold of **much lower dimensionality** compared to that of the original data space
- Training continuous latent variable models is often called **dimensionality reduction**, since there are typically **fewer latent dimensions**
- Examples: Principal Component Analysis, Factor Analysis, Independent Component Analysis

Intrinsic Latent Dimensions


- What are the **intrinsic latent dimensions** in these two datasets?



- How can we find the latent dimensions from this high-dimensional data?

Intrinsic Latent Dimensions


- In this dataset, there are only **3 degrees of freedom of variability**, corresponding to vertical- and horizontal translations, and the rotations



- Each image undergoes a random displacement and rotation within some larger image field
- The resulting images have $100 \times 100 = 10,000$ pixels

Generative View

- Each data example was generated by first selecting a point from a **distribution in the latent space**, then **generating a point from the conditional distribution** in the input space
- **Simple example of a latent variable model**: Assume a Gaussian distribution for both the latent and observed variables
- This can lead to a probabilistic formulation of **Principal Component Analysis** and **Factor Analysis**
- We will look at standard PCA, then briefly note its probabilistic formation.
- Advantages of the **probabilistic formulation**: use of EM for parameter estimation, mixture of PCAs, Bayesian PCA, etc



Principal Component Analysis

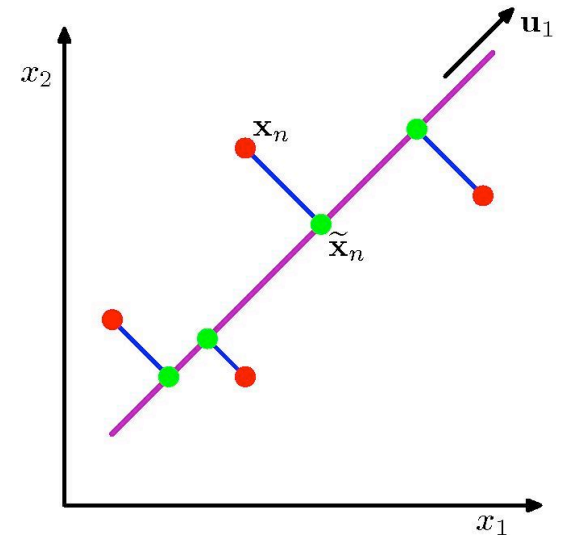
- Used for data compression, visualization, feature extraction, dimensionality reduction

- The goal is to:

- Find the M principal components underlying the D -dimensional data: select the top M **eigenvectors** of \mathbf{S} (data covariance matrix):

$$\{\mathbf{u}_1, \dots, \mathbf{u}_M\}.$$

- **Project** each input vector \mathbf{x} into this subspace, e.g. $z_{n1} = \mathbf{x}_n^T \mathbf{u}_1$.



- The full projection into M dimensions takes the form:

$$\begin{bmatrix} \mathbf{u}_1^T \\ \vdots \\ \mathbf{u}_M^T \end{bmatrix} [\mathbf{x}_1 \cdots \mathbf{x}_N] = [\mathbf{z}_1 \cdots \mathbf{z}_N]$$

Two equivalent views / derivations:

- PCA **maximizes the variance** of the projected data (the scatter of the green points)
- PCA **minimizes the error** of the projected data (the mean of the squared blue lines)


Maximum Variance Formulation

- Consider a dataset $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, $\mathbf{x}_n \in \mathbb{R}^D$. Our goal is to **project data** onto a space having dimensionality $M < D$
- Consider the projection into $M=1$ dimensional space
- Define **the direction of this space** using a D -dimensional unit vector \mathbf{u}_1 , so that $\mathbf{u}_1^T \mathbf{u}_1 = 1$.
- **Objective:** maximize the variance of the projected data with respect to \mathbf{u}_1

$$\frac{1}{N} \sum_{n=1}^N \{\mathbf{u}_1^T \mathbf{x}_n - \mathbf{u}_1^T \bar{\mathbf{x}}\}^2 = \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1$$

where sample **mean** and **data covariance** is given by:

$$\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$
$$\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T$$



Maximum Variance Formulation

- Maximize the variance of the projected data:

$$\frac{1}{N} \sum_{n=1}^N \{ \mathbf{u}_1^T \mathbf{x}_n - \mathbf{u}_1^T \bar{\mathbf{x}} \}^2 = \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1$$

- Must constrain $\|\mathbf{u}_1\|$, and we choose $\|\mathbf{u}_1\| = 1$.
Using a Lagrange multiplier, maximize:

$$\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 + \lambda(1 - \mathbf{u}_1^T \mathbf{u}_1)$$

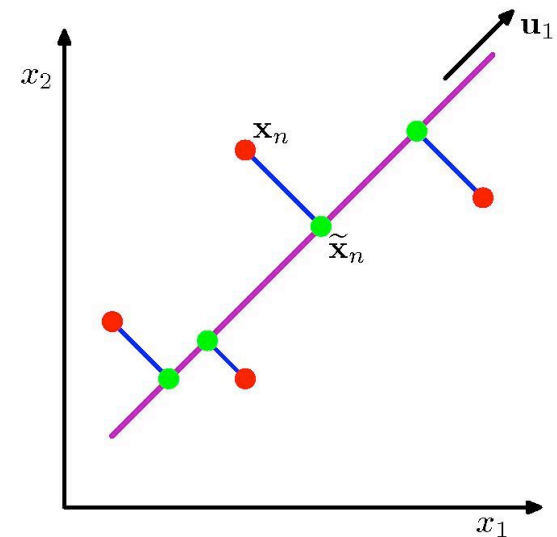
- Setting the derivative with respect to \mathbf{u}_1 to zero:

$$\mathbf{S} \mathbf{u}_1 = \lambda_1 \mathbf{u}_1.$$

- Hence \mathbf{u}_1 must be an eigenvector of \mathbf{S}
- The maximum variance of the projected data is given by:

$$\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 = \lambda_1.$$

- Optimal \mathbf{u}_1 is the principal component (eigenvector with maximal eigenvalue)




Minimum Error Formulation

- Introduce a **complete orthonormal set** of D -dimensional basis vectors: $\{\mathbf{u}_1, \dots, \mathbf{u}_D\}$:

$$\mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}.$$

- Without loss of generality, we can write:

$$\mathbf{x}_n = \sum_{i=1}^D \alpha_{ni} \mathbf{u}_i, \quad \alpha_{ni} = \mathbf{x}_n^T \mathbf{u}_i.$$


Rotation of the coordinate system to a new system defined by \mathbf{u}_i .

- Our goal is to represent data points by the **projection into an M -dimensional subspace** (plus some distortion)
- Represent the M -dimensional linear subspace by the **first M basis vectors**:

$$\tilde{\mathbf{x}}_n = \sum_{i=1}^M z_{ni} \mathbf{u}_i + \sum_{i=M+1}^D b_i \mathbf{u}_i.$$

first M basis are principle components

Minimum Error Formulation

- Represent the M -dimensional linear subspace by the first M basis vectors:

$$\tilde{\mathbf{x}}_n = \sum_{i=1}^M z_{ni} \mathbf{u}_i + \sum_{i=M+1}^D b_i \mathbf{u}_i.$$

where z_{ni} depend on the particular data point and b_i are constants

- Objective:** minimize distortion with respect to \mathbf{u}_i , z_{ni} , and b_i


$$J = \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \tilde{\mathbf{x}}_n\|^2.$$

- Minimizing with respect to z_{nj} , b_j :

$$z_{nj} = \mathbf{x}_n^T \mathbf{u}_j$$

$$b_j = \bar{\mathbf{x}}^T \mathbf{u}_j$$
- Hence, the objective reduces to:

$$J = \frac{1}{N} \sum_{n=1}^N \sum_{i=M+1}^D (\mathbf{x}_n^T \mathbf{u}_i - \bar{\mathbf{x}}^T \mathbf{u}_i)^2 = \sum_{i=M+1}^D \mathbf{u}_i^T \mathbf{S} \mathbf{u}_i.$$



Minimum Error Formulation

- Minimizing distortion with respect to \mathbf{u}_i is a constrained minimization problem:


$$J = \frac{1}{N} \sum_{n=1}^N \|\mathbf{x}_n - \tilde{\mathbf{x}}_n\|^2 = \sum_{i=M+1}^D \mathbf{u}_i^T \mathbf{S} \mathbf{u}_i.$$

- You are not responsible for showing that the *general* solution is to choose \mathbf{u}_i to be **eigenvectors of the covariance matrix**: $\mathbf{S} \mathbf{u}_i = \lambda_i \mathbf{u}_i$.

- The distortion is then given by: $J = \sum_{i=M+1}^D \lambda_i$.

- The objective is minimized when the remaining $D - M$ components are the **eigenvectors** of \mathbf{S} with **lowest eigenvalues** → same result

- Exercise:** show that for $D=2$, $M=1$, \mathbf{u}_2 is the eigenvector of \mathbf{S} corresponding to the second-largest eigenvalue



- We will later see a generalization of PCA: deep autoencoders

Applications of PCA

- Run PCA on 2429 19x19 grayscale images (CBCL database)



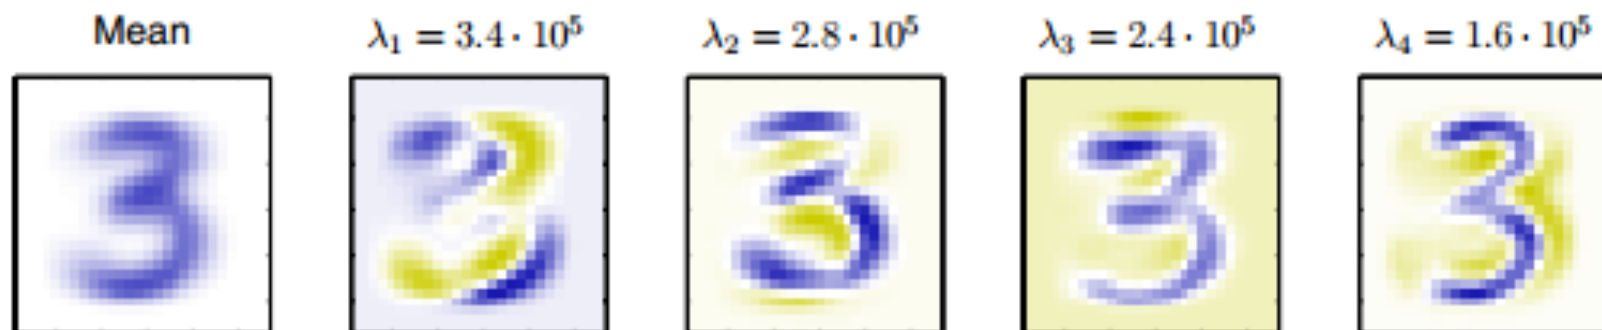
- **Data compression**: We can get good reconstructions with only 3 components
- **Pre-processing**: We can apply a **standard classifier to latent representation** – PCA with 3 components obtains 79% accuracy on face/non-face discrimination in test data, vs. 76.8% for a mixture of Gaussians with 84 components
- **Data visualization**: by projecting the data onto the first two principal components

Learned Basis

- Run PCA on 2429 19x19 grayscale images (CBCL database)



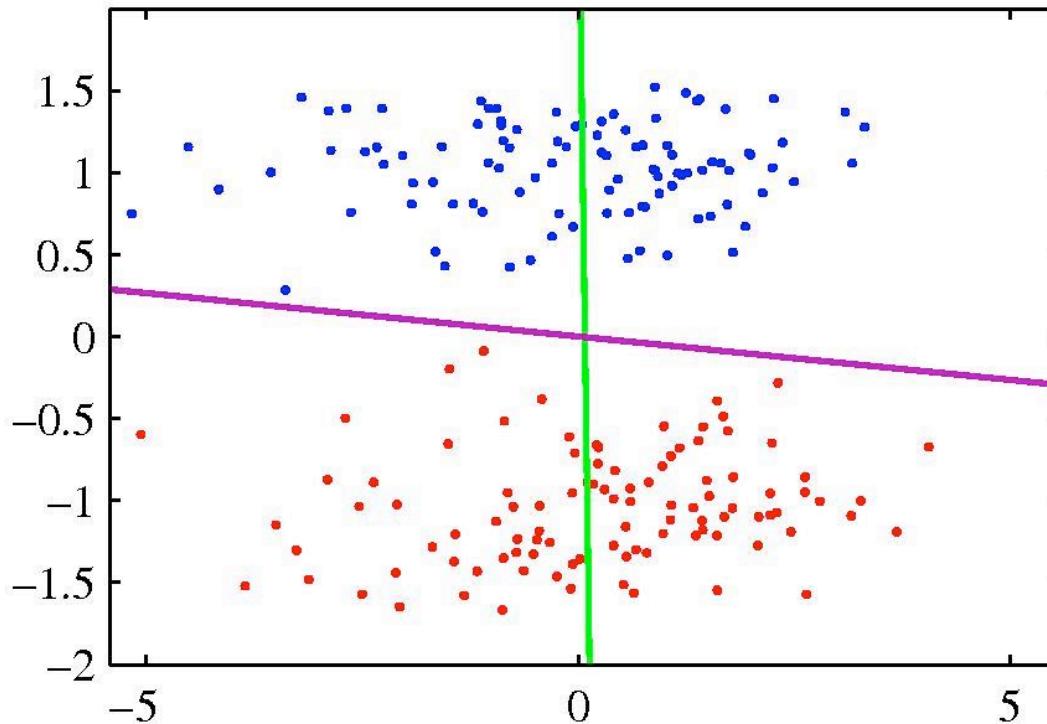
Eigenvectors for 3's



The mean vector \bar{x} along with the first four PCA eigenvectors u_1, \dots, u_4 for the off-line digits data set, together with the corresponding eigenvalues.

PCA vs. Fisher's LDA

- A comparison of PCA with Fisher's LDA for linear dimensionality reduction



- PCA chooses the **direction of maximum variance** (magenta curve) leading to strong class overlap (unsupervised)
- LDA takes into account the **class labels** (supervised), leading to a projection into the green curve

PCA for High-Dimensional Data

- In some applications of PCA, the number of data points is smaller than the dimensionality of the data space, i.e. $N < D$
- So far, we've needed to find the eigenvectors of the $D \times D$ data covariance matrix \mathbf{S} , which scales as $\mathcal{O}(D^3)$
- Direct application of PCA may be computationally infeasible
- **Solution:** Let \mathbf{X} be the $N \times D$ centred data matrix. The corresponding eigenvector equation becomes:

$$\frac{1}{N} \mathbf{X}^T \mathbf{X} \mathbf{u}_i = \lambda_i \mathbf{u}_i.$$

- Pre-multiply by \mathbf{X} :

$$\frac{1}{N} \mathbf{X} \mathbf{X}^T (\mathbf{X} \mathbf{u}_i) = \lambda_i (\mathbf{X} \mathbf{u}_i).$$

PCA for High-Dimensional Data

- Define $\mathbf{v}_i = \mathbf{X}\mathbf{u}_i$, and hence we have:

$$\frac{1}{N} \mathbf{X} \mathbf{X}^T \mathbf{v}_i = \lambda_i \mathbf{v}_i.$$

- This is an **eigenvector equation** for the $N \times N$ matrix
- It has the same $N - 1$ eigenvalues as the original data covariance matrix \mathbf{S} (which itself has **an additional $D - N + 1$ zero eigenvalues**).
- Computational cost scales as $\mathcal{O}(N^3)$ rather than $\mathcal{O}(D^3)$
- To determine eigenvectors, we pre-multiply by \mathbf{X}^T :

$$\left(\frac{1}{N} \mathbf{X}^T \mathbf{X} \right) (\mathbf{X}^T \mathbf{v}_i) = \lambda_i \mathbf{X}^T \mathbf{v}_i.$$

- Hence $\mathbf{X}^T \mathbf{v}_i$ is **an eigenvector of \mathbf{S} with eigenvalue λ_i**
- These eigenvectors may not be normalized. **Exercise:** How do we ensure that they are normalized? (Answer: Bishop exercise 12.3)


Probabilistic PCA

- A probabilistic, generative view of data which is not covered in STA414
- Advantages of probabilistic PCA (PPCA):
 - We can **derive an EM algorithm** for PCA, which in some settings is more computationally efficient
 - PPCA allows us to deal with **missing values** in the data set
 - We can formulate a **mixture of PPCAs** in a principled way
 - PPCA forms the basis for **Bayesian PCA**, in which the dimensionality of the principal subspace can be determined from the data
 - The **existence of a likelihood function** allows direct comparisons with other probabilistic density models
 - And more

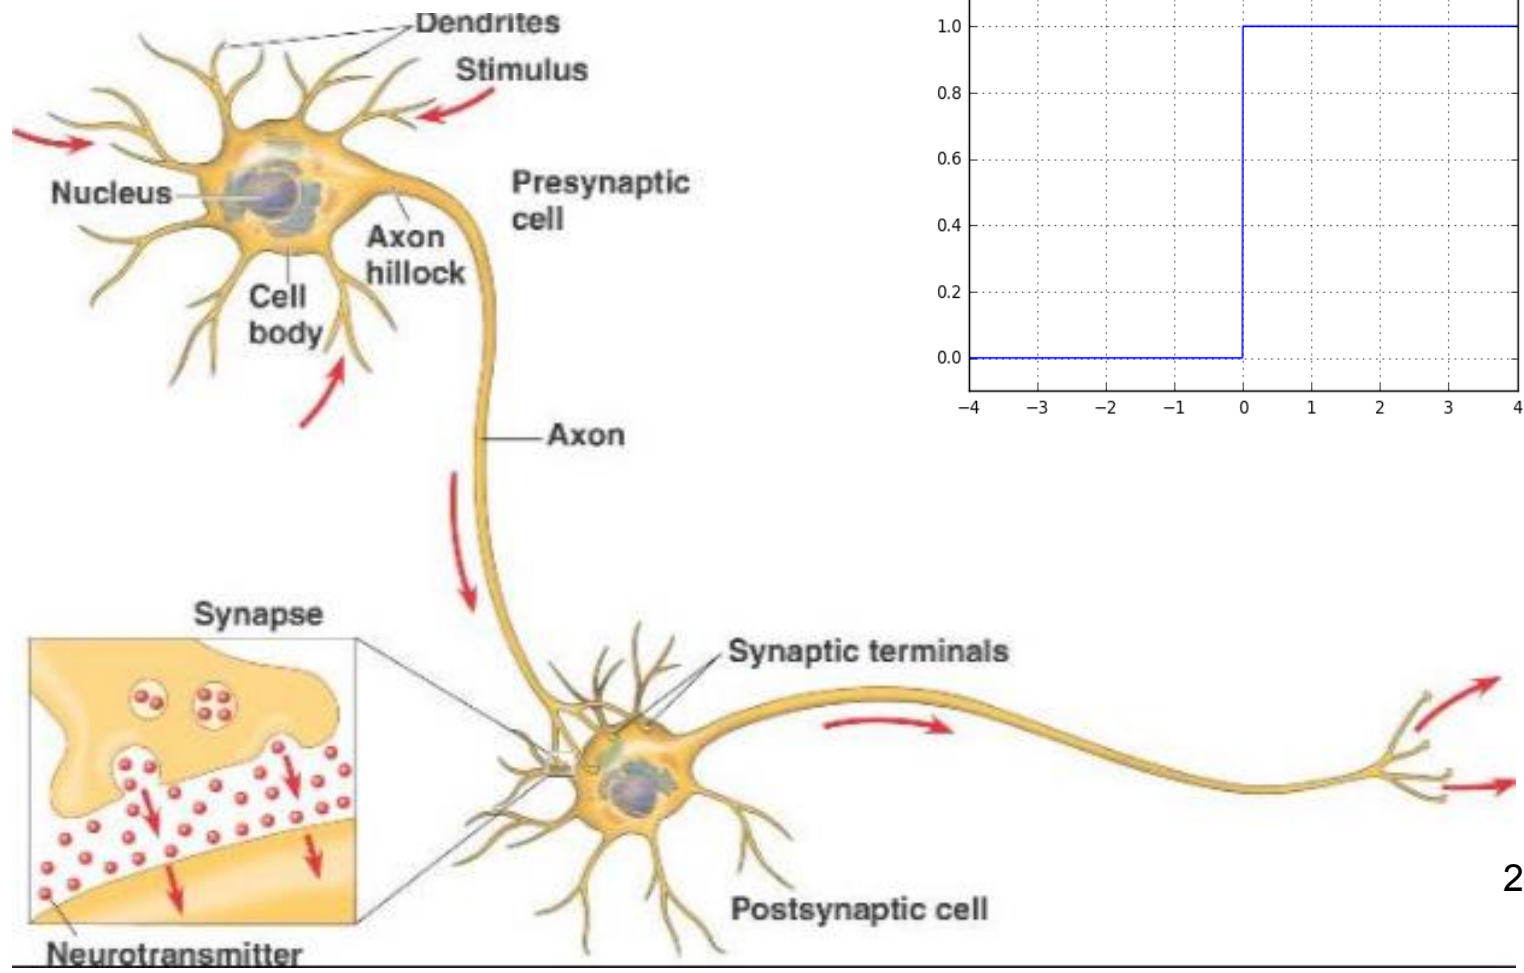
Outline

- Continuous latent variable models
 - Background
 - PCA
- Neural networks
 - Introduction
 - Autoencoders
 - Learning neural networks

One neuron (biological)

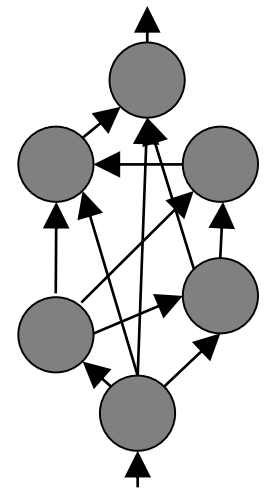


Two neurons




Neural networks

- Neural networks are flexible computation models that consist of many smaller computational modules called neurons or **hidden units**:
 - Neural networks are real-valued
 - It is very modular and some special modules are designed for reusability and abstraction
 - All continuous functions can be represented as neural networks
 - All the learnt knowledge of a neural network is stored in its **weight connections**; it is also called “connectionism” (a name popular before the first AI Winter)



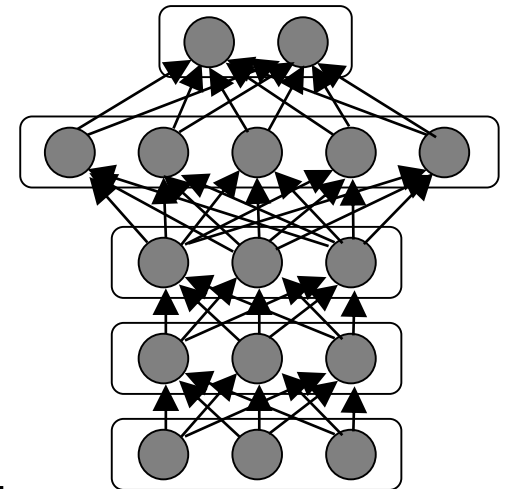
Neural networks

- One very useful abstraction is the concept of a “layer”:
 - A hidden layer is a group of hidden units that have connections one layer above and one layer below
 - There is no connection among the hidden units within a layer
 - This abstraction is computationally efficient because all the hidden units within a layer can be computed in parallel



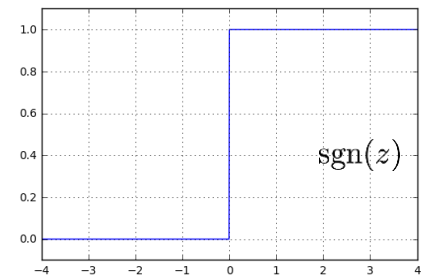
Neural networks

- Deep learning typically refers to a neural network with more than three hidden layers
 - Deep neural networks can mathematically represent any continuous function given enough layers, but they also require additional tricks to learn useful representations for any tasks
 - They work really well in supervised learning, given enough data
 - A deep neural network is like a complex system in biology: we understand a lot about what the simple module does, but it quickly becomes hard to understand what the system does, i.e. a “black box”



An artificial neuron

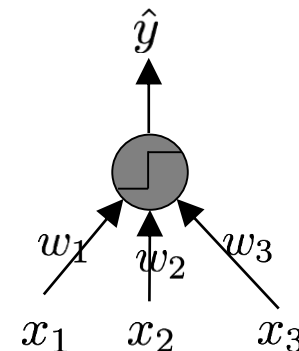
- An artificial neuron is a simple computation unit that receives inputs from other simple computation units:
 - The effect of each input on the final output of the neuron is controlled by a weight
 - The weights can be positive or negative values for encoding +ve or -ve contributions from the inputs
 - A weighted sum of the inputs was first proposed by McCulloch-Pitts (1943)



$$\text{sgn}(z) = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases}$$

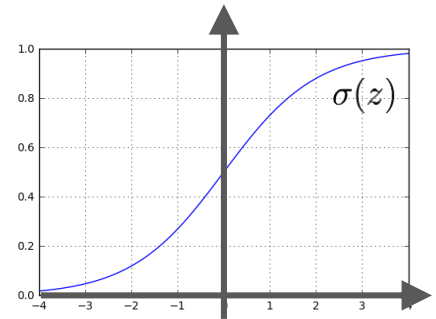
$$\frac{\partial \text{sgn}}{\partial z} = 0, \forall z \neq 0$$

$$\hat{y} = \text{sgn}\left(\sum_n w_n x_n + b\right)$$



Some simple neurons: sigmoid neurons


- An alternative to a hard step function is a soft, differentiable function; desirable when using a gradient-descent algorithm (e.g. Week 3) to learn our model
 - Sigmoid neurons can be thought of as soft-thresholding units
 - Logistic regression models are simply neural networks with a single logistic neuron
 - With high enough w values, the sigmoid can behave like the *sgn* function (slides 47-48)



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

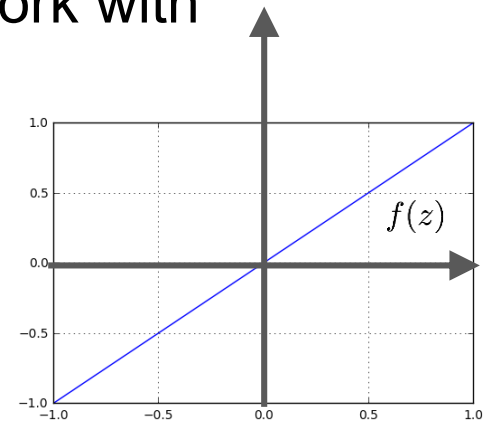
$$\frac{\partial \sigma}{\partial z} = \sigma(z)(1 - \sigma(z))$$

$$\hat{y} = \sigma\left(\sum_n w_n x_n + b\right)$$



Some simple neurons: linear neurons

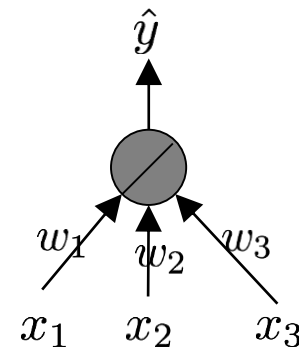
- A Linear neuron directly outputs the weighted sum of the inputs:
 - Linear regression is the simplest neural network with a single linear neuron
 - It has a constant partial derivative which is great for gradient descent
 - However, stacking layers of linear neurons does not increase the representational power of a model. Nonlinearity is important for building richer and more flexible models



$$f(z) = z$$

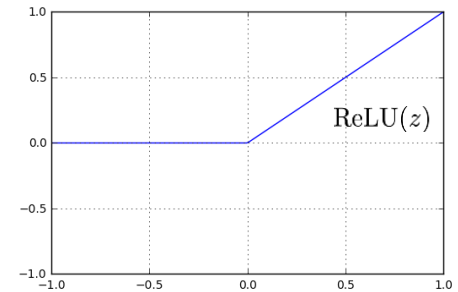
$$\frac{\partial f}{\partial z} = 1$$

$$\hat{y} = \sum_n w_n x_n + b$$




Some simple neurons: rectified linear units (ReLU)

- Linear neurons can be modified to exhibit nonlinear behaviour:
 - The non-positive values are forced to be zero
 - ReLU neurons still have a constant gradient if the weighted sum of the inputs is positive
 - It is mathematically non-differentiable at zero, but we ignore that and use gradient descent anyways. It works well (numerically, we tend not to get exactly zero summed inputs)



$$\text{ReLU}(z) = \max(0, z) \quad \hat{y} = \text{ReLU}\left(\sum_n w_n x_n + b\right)$$

$$\frac{\partial \text{ReLU}}{\partial z} = \begin{cases} 1 & z > 0 \\ 0 & z < 0 \end{cases}$$



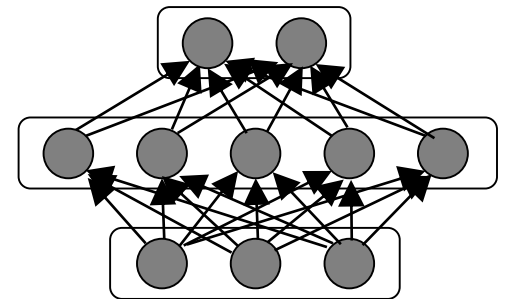
Context of neural networks

Two common ways to solve a problem:

1. Hire people to hard-code a program
2. Gather a huge dataset and learn the program from the data

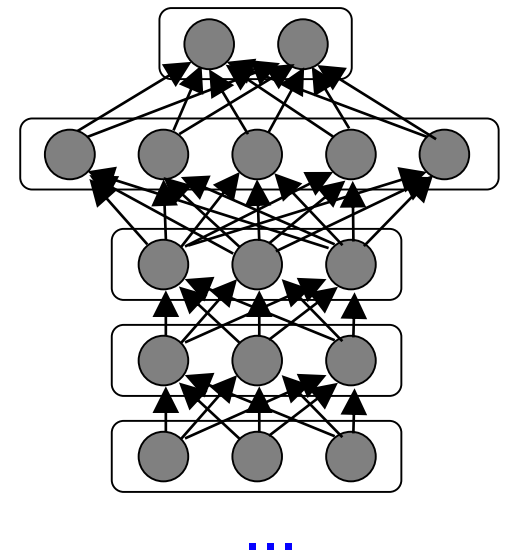
Deep neural networks avoid time-consuming feature-engineering by hand, and as the datasets grow larger they can discover better and better features without human intervention.

Neural networks can also be understood as a form of *adaptive basis function model* in which the model learns layers of basis functions. The activation function used for a neuron is similar to the nonlinear basis functions.




Hyper-parameters

- Choices:
 - How many hidden units to use in each hidden layer?
 - How many layers in total?
 - Which hidden activation function?
- Good answer: decide these hyper-parameters using a **validation set**
- *One* common, practical answer:
 - Around 500-2000 hidden units
 - 2-3 layers
 - Choosing ReLU often leads to fast convergence



Autoencoders

- Neural networks can be used for **nonlinear dimensionality reduction**
- This is achieved by having the **same number of outputs as inputs**. These models are called **autoencoders**
- Consider a multilayer perceptron that has D inputs, D outputs, and M hidden units such that $M < D$
- It is useful if we can squeeze the information **through some kind of bottleneck**
- If we use a **linear network**, this is very similar to Principal Component Analysis



Autoencoders and PCA


- Given an input \mathbf{x} , its corresponding reconstruction is given by:

$$y_k(\mathbf{x}, \mathbf{w}) = \sum_{j=1}^M w_{kj}^{(2)} \sigma \left(\sum_{i=1}^D w_{ji}^{(1)} x_i \right), \quad k = 1, \dots, D.$$

- We can determine the network parameters \mathbf{w} by minimizing the reconstruction error:

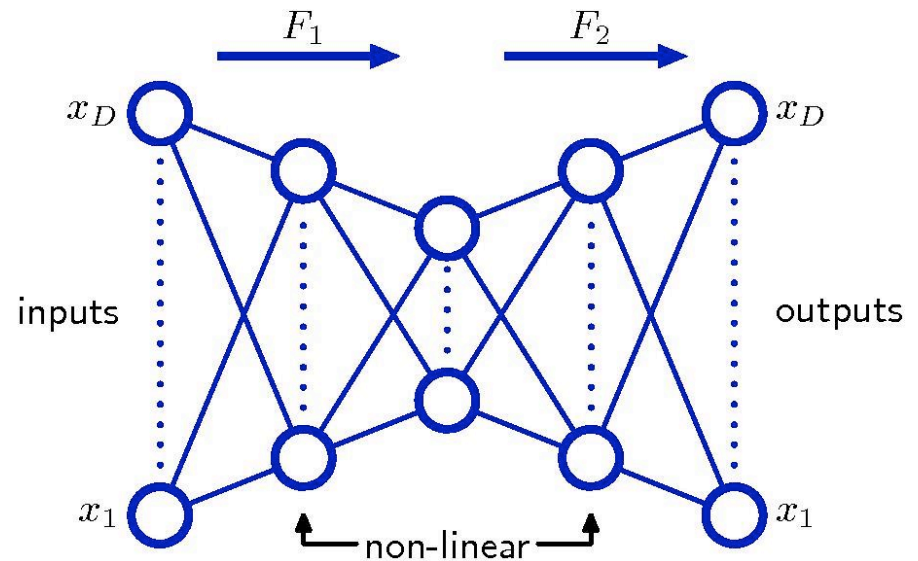
$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \|y(\mathbf{x}_n, \mathbf{w}) - \mathbf{x}_n\|^2.$$

- If the hidden and output layers are linear, we will learn hidden units that are a linear function of the data and minimize the squared error
- The M hidden units will span the same space as the first m principal components. The weight vectors may not be orthogonal



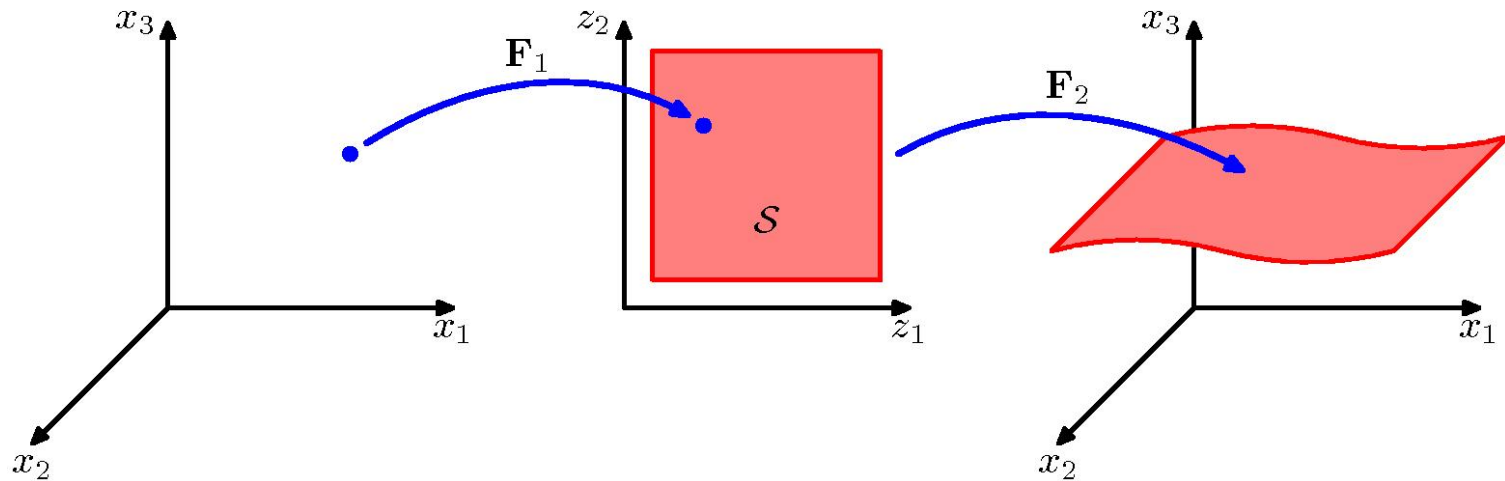
Deep Autoencoders

- We can put **extra nonlinear hidden layers** between the input and the bottleneck and between the bottleneck and the output
- This gives a **nonlinear generalization** of PCA
- It is good for nonlinear dimensionality reduction
- The network can be trained by the **minimization of the reconstruction error function**
- Much harder to train



Geometrical Interpretation

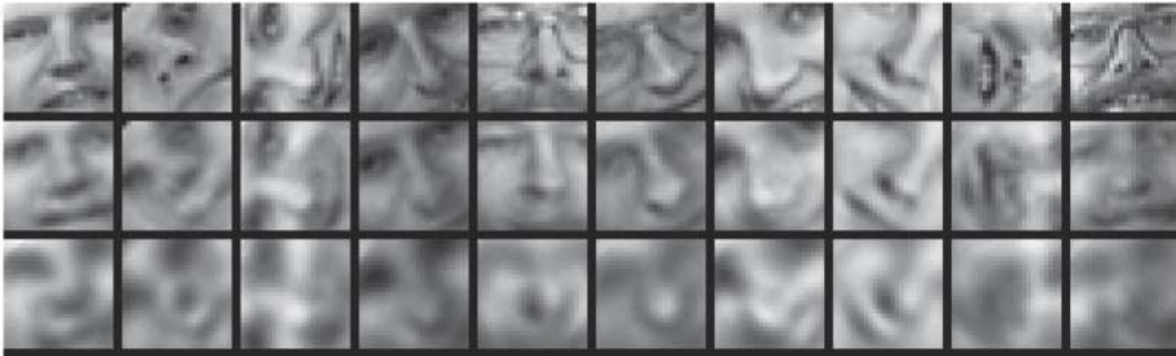
- Geometrical interpretation of the mappings performed by the network with 2 hidden layers, for the case of $D=3$ and $M=2$ units in the middle layer



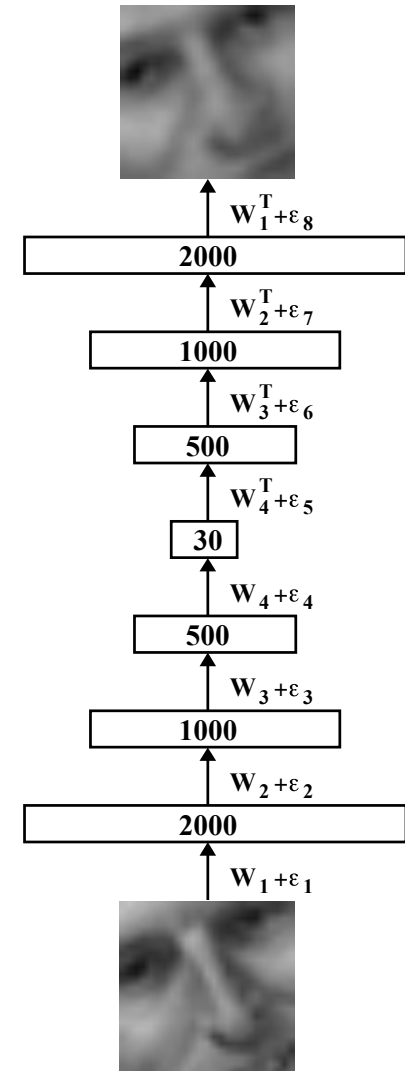
- The mapping F_1 defines a nonlinear projection of points in the original D -space into the M -dimensional subspace
- The mapping F_2 maps from an M -dimensional space into D -dimensional space

Deep Autoencoders

- We can consider deep autoencoders
- There is an efficient way to learn these deep autoencoders



- By row: Real data, Deep autoencoder with a bottleneck of 30 linear units, and 30-d PCA



Deep Autoencoders

- We can consider deep autoencoders
- Similar model for MNIST handwritten digits:



Real data

30-d deep autoencoder

30-d logistic PCA

30-d PCA


- The Deep autoencoder produces much better reconstructions

Class Structure of the Data

- Do the 30-D codes found by the deep autoencoder preserve the class structure of the data?
- Take the 30-D activity patterns in the code layer and display them in 2-D using **nonlinear multi-dimensional scaling** (UNI-SNE)
- Will the learning find the natural classes?

Class Structure of the Data

- Do the 30-D codes found by the deep autoencoder preserve the class structure of the data?




Outline

- Continuous latent variable models
 - Background
 - PCA
- Neural networks
 - Introduction
 - Autoencoders
 - Learning neural networks

Notation for neural networks

- The model now consists of many artificial neurons wired together into a large network. For clarity, we will use the following notation for our algorithms:
 - The output of a neuron or the hidden activation is denoted as h
 - Scalar weight connections are indexed by the two neurons it connects
 - The input to the network is denoted x
 - The output of the network is denoted as \hat{y}
 - The element-wise hidden activation function or the activation function or nonlinearity, denoted as $\phi(\cdot)$, is the nonlinear transformation for the weighted sum of the inputs of a neuron, e.g. sigmoid, ReLU...
 - The weighted sum of a neuron's inputs is denoted as z

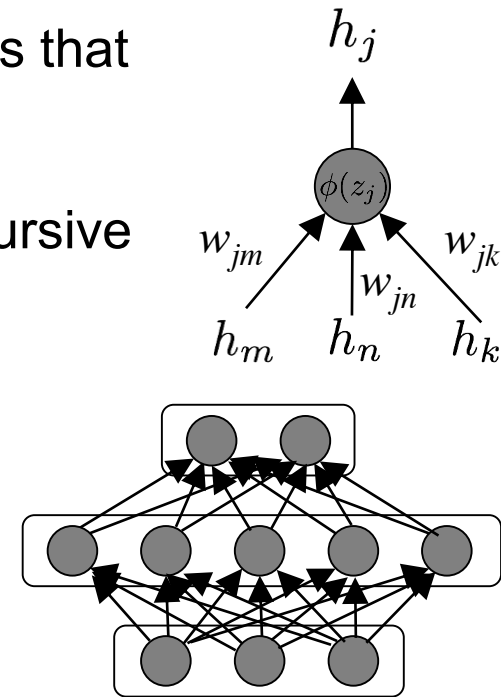


Forward propagation

- Forward propagation computes all the hidden activations h and the output of the neural network \hat{y}

- $h_j = \phi(z_j) = \phi(\sum_n w_{jn} h_n + b_j)$

- This requires computing all the hidden activations that are the inputs to the current hidden units.
 - The forward propagation can be written as a recursive algorithm
 - The naive recursive algorithm is bad because there are a lot of redundant computations. We would like to cache the appropriate intermediate values and reuse them



Back-propagation


- Back-propagation (Rumelhart, Hinton and Williams, 1986) is a dynamic programming method to reuse previous computations when computing the gradient of some variable using the chain rule from calculus

- $h_j = \phi(z_j) = \phi(\sum_n w_{jn} h_n + b_j)$

- In its simplest form: $\frac{\partial \mathcal{L}}{\partial w_{jn}} = \frac{\partial \mathcal{L}}{\partial h_j} \frac{\partial h_j}{\partial z_j} \frac{\partial z_j}{\partial w_{jn}}$

- $\frac{\partial \mathcal{L}}{\partial h_j}$ can be further expanded until the output of the neural network

- The key observation here is that the gradient of a connection is a product of the input and the partial derivative of the weighted sum of that neuron



$$\frac{\partial \mathcal{L}}{\partial w_{jn}} = \frac{\partial \mathcal{L}}{\partial z_j} h_n$$


Appendix

- Slides 47-48: an example of a simple neural network
- Slides 49-54: backpropagation, in which w_{ij} becomes w_{ji}

Example 1: representing digital circuits with neural networks

- Let us look at a simple example of a soft OR gate simulated by a neural network:
 - Use a single sigmoid neuron with two inputs and a bias unit.


	$x_2=0$	$x_2=1$
$x_1=0$	0	1
$x_1=1$	1	1



Example 1: representing digital circuits with neural networks

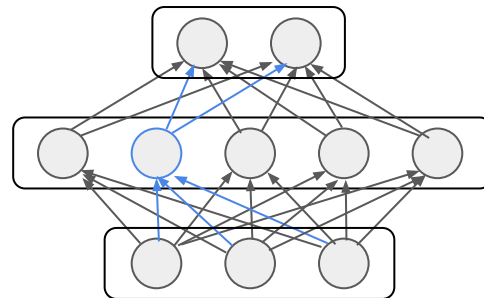
- Let us look at a simple example of a soft OR gate simulated by a neural network:
 - Use a single sigmoid neuron with two inputs and a bias unit.
 - One possible solution is to use the bias as a threshold while setting w_1 and w_2 to be large positive values. When either of the inputs is non-zero, the sigmoid neuron will be turned on and the output will be 1.

	$x_2=0$	$x_2=1$
$x_1=0$	0	1
$x_1=1$	1	1



Learning fully connected multi-layer neural networks

- There are many choices when “crafting” the architecture of a neural network. The fully connected multi-layer NN is the most general multi-layer NN:
 - Each neuron has its incoming weights connected to *all* the neurons from the previous layer and its outgoing weights connected to *all* the neurons in the next layer.
- Fully connected network is the go-to architecture choice if we do not have any additional information about the dataset.
 - After choosing the network architecture, there are a few more engineering choices: #hidden units, #layers, the type of activation function.
 - The output units type: linear, logistic or softmax are determined by output tasks, i.e. regression or classification task




Learning fully connected multi-layer neural networks

- Consider a fully connected neural network with 3 hidden layers:
 - The input to the neural network is an N -dimensional vector \mathbf{x} . There are H_1 , H_2 , and H_3 hidden units in the three hidden layers. We use superscript to index the layers.
 - There are four weight matrices among the hidden layers, e.g. $W^{(2)} \in \mathbb{R}^{H_2 \times H_1}$, $b^{(2)} \in \mathbb{R}^{H_2}$
 - The j th row of the weight matrix $W^{(2)}$ is denoted as $W_j^{(2)} \in \mathbb{R}^{H_1}$
- The hidden activation of the j th hidden unit $h_j^{(2)}$ in the second hidden layer is the weighted sum of the first hidden layer:

$$h_j^{(2)} = \phi(z_j^{(2)}) = \phi\left(\sum_i w_{ij}^{(2)} h_i^{(1)} + b_j^{(2)}\right) = \phi\left(W_j^{(2)T} \mathbf{h}^{(1)} + b_j^{(2)}\right)$$

- We can use vector notation to express the hidden vector:

$$\mathbf{h}^{(2)} = \begin{bmatrix} h_1^{(2)} \\ \vdots \\ h_{H_2}^{(2)} \end{bmatrix} = \begin{bmatrix} \phi(z_1^{(2)}) \\ \vdots \\ \phi(z_{H_2}^{(2)}) \end{bmatrix} = \phi\left(\begin{bmatrix} W_1^{(2)T} \\ \vdots \\ W_{H_2}^{(2)T} \end{bmatrix} \mathbf{h}^{(1)} + b^{(2)}\right) = \phi\left(W^{(2)} \mathbf{h}^{(1)} + b^{(2)}\right)$$



Learning fully connected multi-layer neural networks

- For a single data point, we can write the the hidden activations of the fully connected neural network as a recursive computation using the vector notation:

$$\mathbf{z}^{(1)} = W^{(1)}\mathbf{x} + b^{(1)}, \quad \mathbf{h}^{(1)} = \phi\left(\mathbf{z}^{(1)}\right)$$


$$\mathbf{z}^{(2)} = W^{(2)}\mathbf{h}^{(1)} + b^{(2)}, \quad \mathbf{h}^{(2)} = \phi\left(\mathbf{z}^{(2)}\right)$$

$$\mathbf{z}^{(3)} = W^{(3)}\mathbf{h}^{(2)} + b^{(3)}, \quad \mathbf{h}^{(3)} = \phi\left(\mathbf{z}^{(3)}\right)$$

$$\mathbf{z}^{(4)} = W^{(4)}\mathbf{h}^{(3)} + b^{(4)}, \quad \hat{y} = f\left(\mathbf{z}^{(4)}\right)$$

- $f()$ is the output activation function
- The output of the network is then used to compute the loss function on the training data

forward
propagation



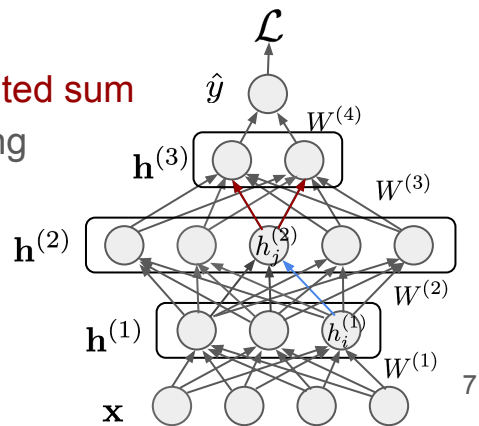
Learning fully connected multi-layer neural networks

- Learning neural networks using stochastic gradient descent requires the gradient of the weight matrices from each hidden layer.
 - Let us consider the gradient of the loss for a single training example. The gradient w.r.t. the incoming weights $w_{ij}^{(2)}$ of the j th hidden unit in the second layer is the *product* of the hidden activation from layer 1 *and* the partial derivative w.r.t. z_j . Remember: $h_j^{(2)} = \phi(z_j^{(2)}) = \phi\left(\sum_i w_{ij}^{(2)} h_i^{(1)} + b_j^{(2)}\right)$

$$\frac{\partial \mathcal{L}}{\partial w_{ij}^{(2)}} = \frac{\partial \mathcal{L}}{\partial z_j^{(2)}} \frac{\partial z_j^{(2)}}{\partial w_{ij}^{(2)}} = \frac{\partial \mathcal{L}}{\partial z_j^{(2)}} h_i^{(1)}$$

- The partial derivative w.r.t. z_j in the second hidden layer is the **weighted sum of the partial derivatives** from the third layer, weighted by the outgoing weights of the j th hidden units:

$$\frac{\partial \mathcal{L}}{\partial z_j^{(2)}} = \frac{\partial \mathcal{L}}{\partial h_j^{(2)}} \frac{\partial h_j^{(2)}}{\partial z_j^{(2)}} = \left(\sum_i \frac{\partial \mathcal{L}}{\partial z_i^{(3)}} \frac{\partial z_i^{(3)}}{\partial h_j^{(2)}} \right) \frac{\partial h_j^{(2)}}{\partial z_j^{(2)}} = \left(\sum_i \frac{\partial \mathcal{L}}{\partial z_i^{(3)}} w_{ji}^{(3)} \right) \frac{\partial h_j^{(2)}}{\partial z_j^{(2)}}$$



Learning fully connected multi-layer neural networks

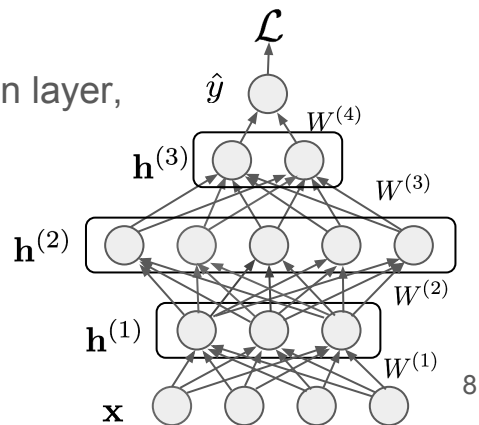
- Similar to the hidden-activation computation (slide 10), the weighted sum of the partial derivatives can be rewritten using vector notation:

$$\frac{\partial \mathcal{L}}{\partial z_j^{(2)}} = \left(\sum_i \frac{\partial \mathcal{L}}{\partial z_i^{(3)}} w_{ji}^{(3)} \right) \frac{\partial h_j^{(2)}}{\partial z_j^{(2)}} = \left(\mathcal{W}_{j \cdot}^{(3)T} \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(3)}} \right) \frac{\partial h_j^{(2)}}{\partial z_j^{(2)}}$$

$$\frac{\partial \mathbf{h}^{(2)}}{\partial \mathbf{z}^{(2)}} = \begin{bmatrix} \frac{\partial h_1^{(2)}}{\partial z_1^{(2)}} & 0 & \dots & 0 \\ \vdots & \frac{\partial h_j^{(2)}}{\partial z_j^{(2)}} & \vdots & \vdots \\ 0 & \dots & \frac{\partial h_{H_2}^{(2)}}{\partial z_{H_2}^{(2)}} \end{bmatrix} = \text{diag} \left\{ \begin{bmatrix} \frac{\partial h_1^{(2)}}{\partial z_1^{(2)}} \\ \vdots \\ \frac{\partial h_{H_2}^{(2)}}{\partial z_{H_2}^{(2)}} \end{bmatrix} \right\}$$

- Here, $\mathcal{W}_{j \cdot}^{(3)}$ is the j th column of the weight matrix $W^{(3)}$
- To express the partial derivatives w.r.t. \mathbf{z} for the entire second hidden layer, we can use a matrix-vector product:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(2)}} = \begin{bmatrix} \frac{\partial \mathcal{L}}{\partial z_1^{(2)}} \\ \vdots \\ \frac{\partial \mathcal{L}}{\partial z_{H_2}^{(2)}} \end{bmatrix} = \frac{\partial \mathbf{h}^{(2)}}{\partial \mathbf{z}^{(2)}} \left(\begin{bmatrix} \mathcal{W}_{1 \cdot}^{(3)T} \\ \vdots \\ \mathcal{W}_{H_2 \cdot}^{(3)T} \end{bmatrix} \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(3)}} \right) = \frac{\partial \mathbf{h}^{(2)}}{\partial \mathbf{z}^{(2)}} \left(W^{(3)T} \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(3)}} \right)$$



Learning fully connected multi-layer neural networks

- For a single training datum, computing the gradient w.r.t. the weight matrices is also a recursive procedure:

- Remember: $\mathbf{z}^{(4)} = W^{(4)}\mathbf{h}^{(3)} + b^{(4)}$, $\hat{y} = f(\mathbf{z}^{(4)})$
- Back-propagation is similar to running the neural network backwards using the transpose of the weight matrices

$$\frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(4)}} = \frac{\partial \hat{y}}{\partial \mathbf{z}^{(4)}} \frac{\partial \mathcal{L}}{\partial \hat{y}}, \quad \frac{\partial \mathcal{L}}{\partial W^{(4)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(4)}} \mathbf{h}^{(3)T}$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(3)}} = \frac{\partial \mathbf{h}^{(3)}}{\partial \mathbf{z}^{(3)}} \left(W^{(4)T} \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(4)}} \right), \quad \frac{\partial \mathcal{L}}{\partial W^{(3)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(3)}} \mathbf{h}^{(2)T} \quad \text{back-propagation}$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(2)}} = \frac{\partial \mathbf{h}^{(2)}}{\partial \mathbf{z}^{(2)}} \left(W^{(3)T} \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(3)}} \right), \quad \frac{\partial \mathcal{L}}{\partial W^{(2)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(2)}} \mathbf{h}^{(1)T}$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(1)}} = \frac{\partial \mathbf{h}^{(1)}}{\partial \mathbf{z}^{(1)}} \left(W^{(2)T} \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(2)}} \right), \quad \frac{\partial \mathcal{L}}{\partial W^{(1)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(1)}} \mathbf{x}^T$$

What about the expression for the bias units?

